

Docker and Kubernetes





Somkiat Puisungnoen

Somkiat Puisungnoen

Update Info 1

View Activity Log 10+

...
Timeline About Friends 3,138 Photos More

When did you work at Opendream?
... 22 Pending Items

Post Photo/Video Live Video Life Event

What's on your mind?

Public Post

Intro
Software Craftsmanship

Software Practitioner at สยามชัมนาภิกิจ พ.ศ. 2556

Agile Practitioner and Technical at SPRINT3r

Somkiat Puisungnoen 15 mins · Bangkok · ⚙️

Java and Bigdata





Page

Messages

Notifications 3

Insights

Publishing Tools

Settings

Help ▾



somkiat.cc

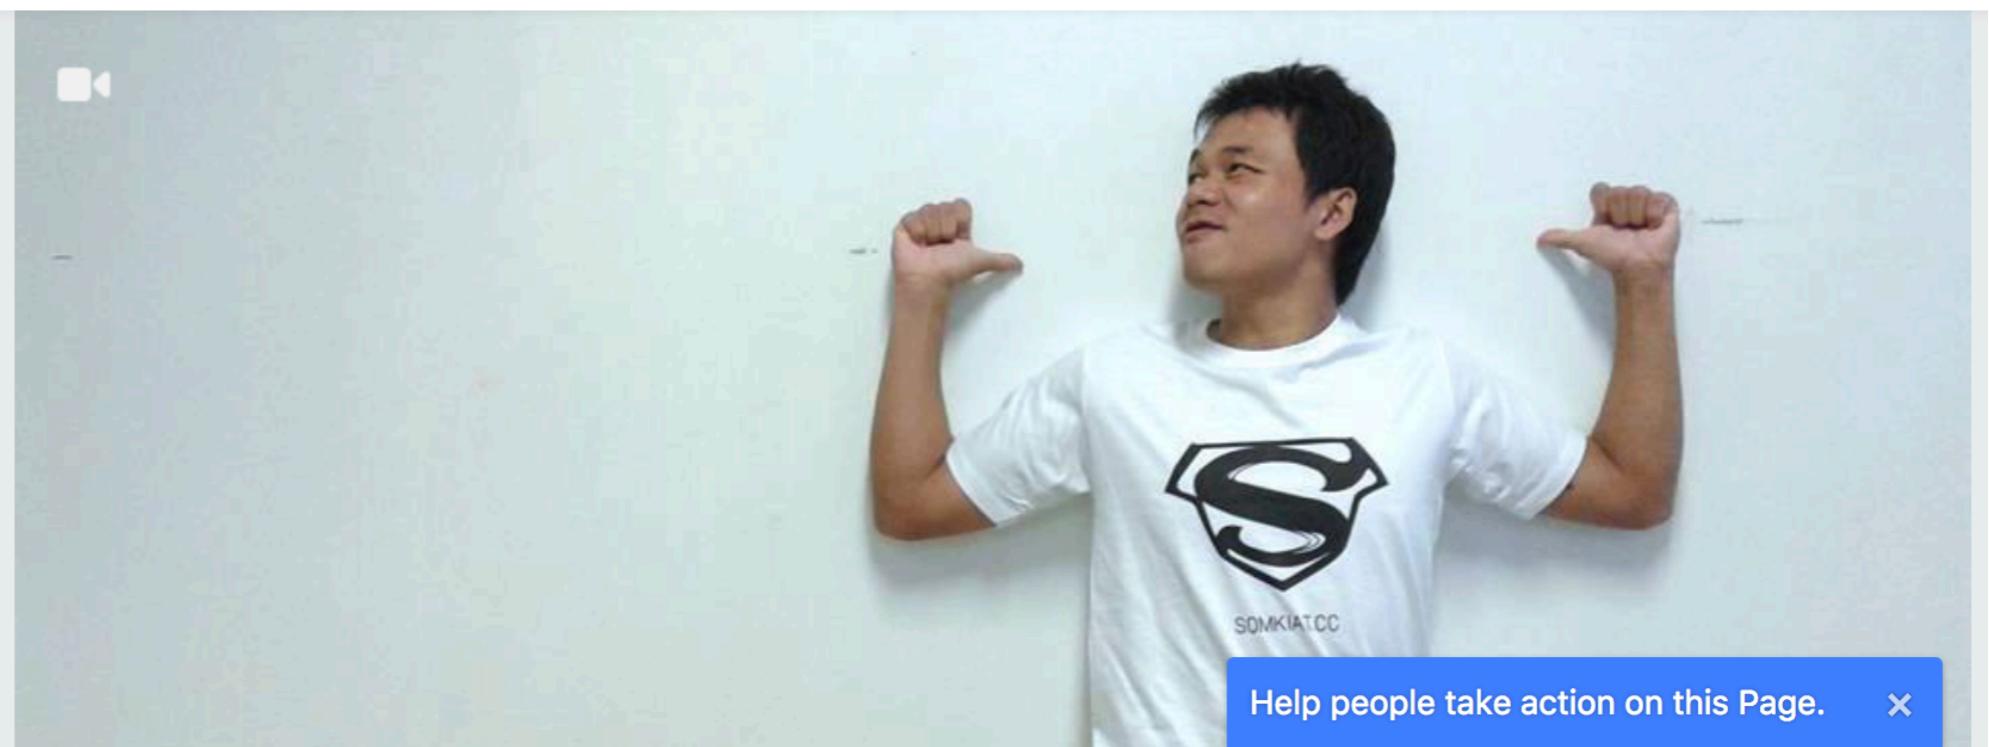
@somkiat.cc

Home

Posts

Videos

Photos



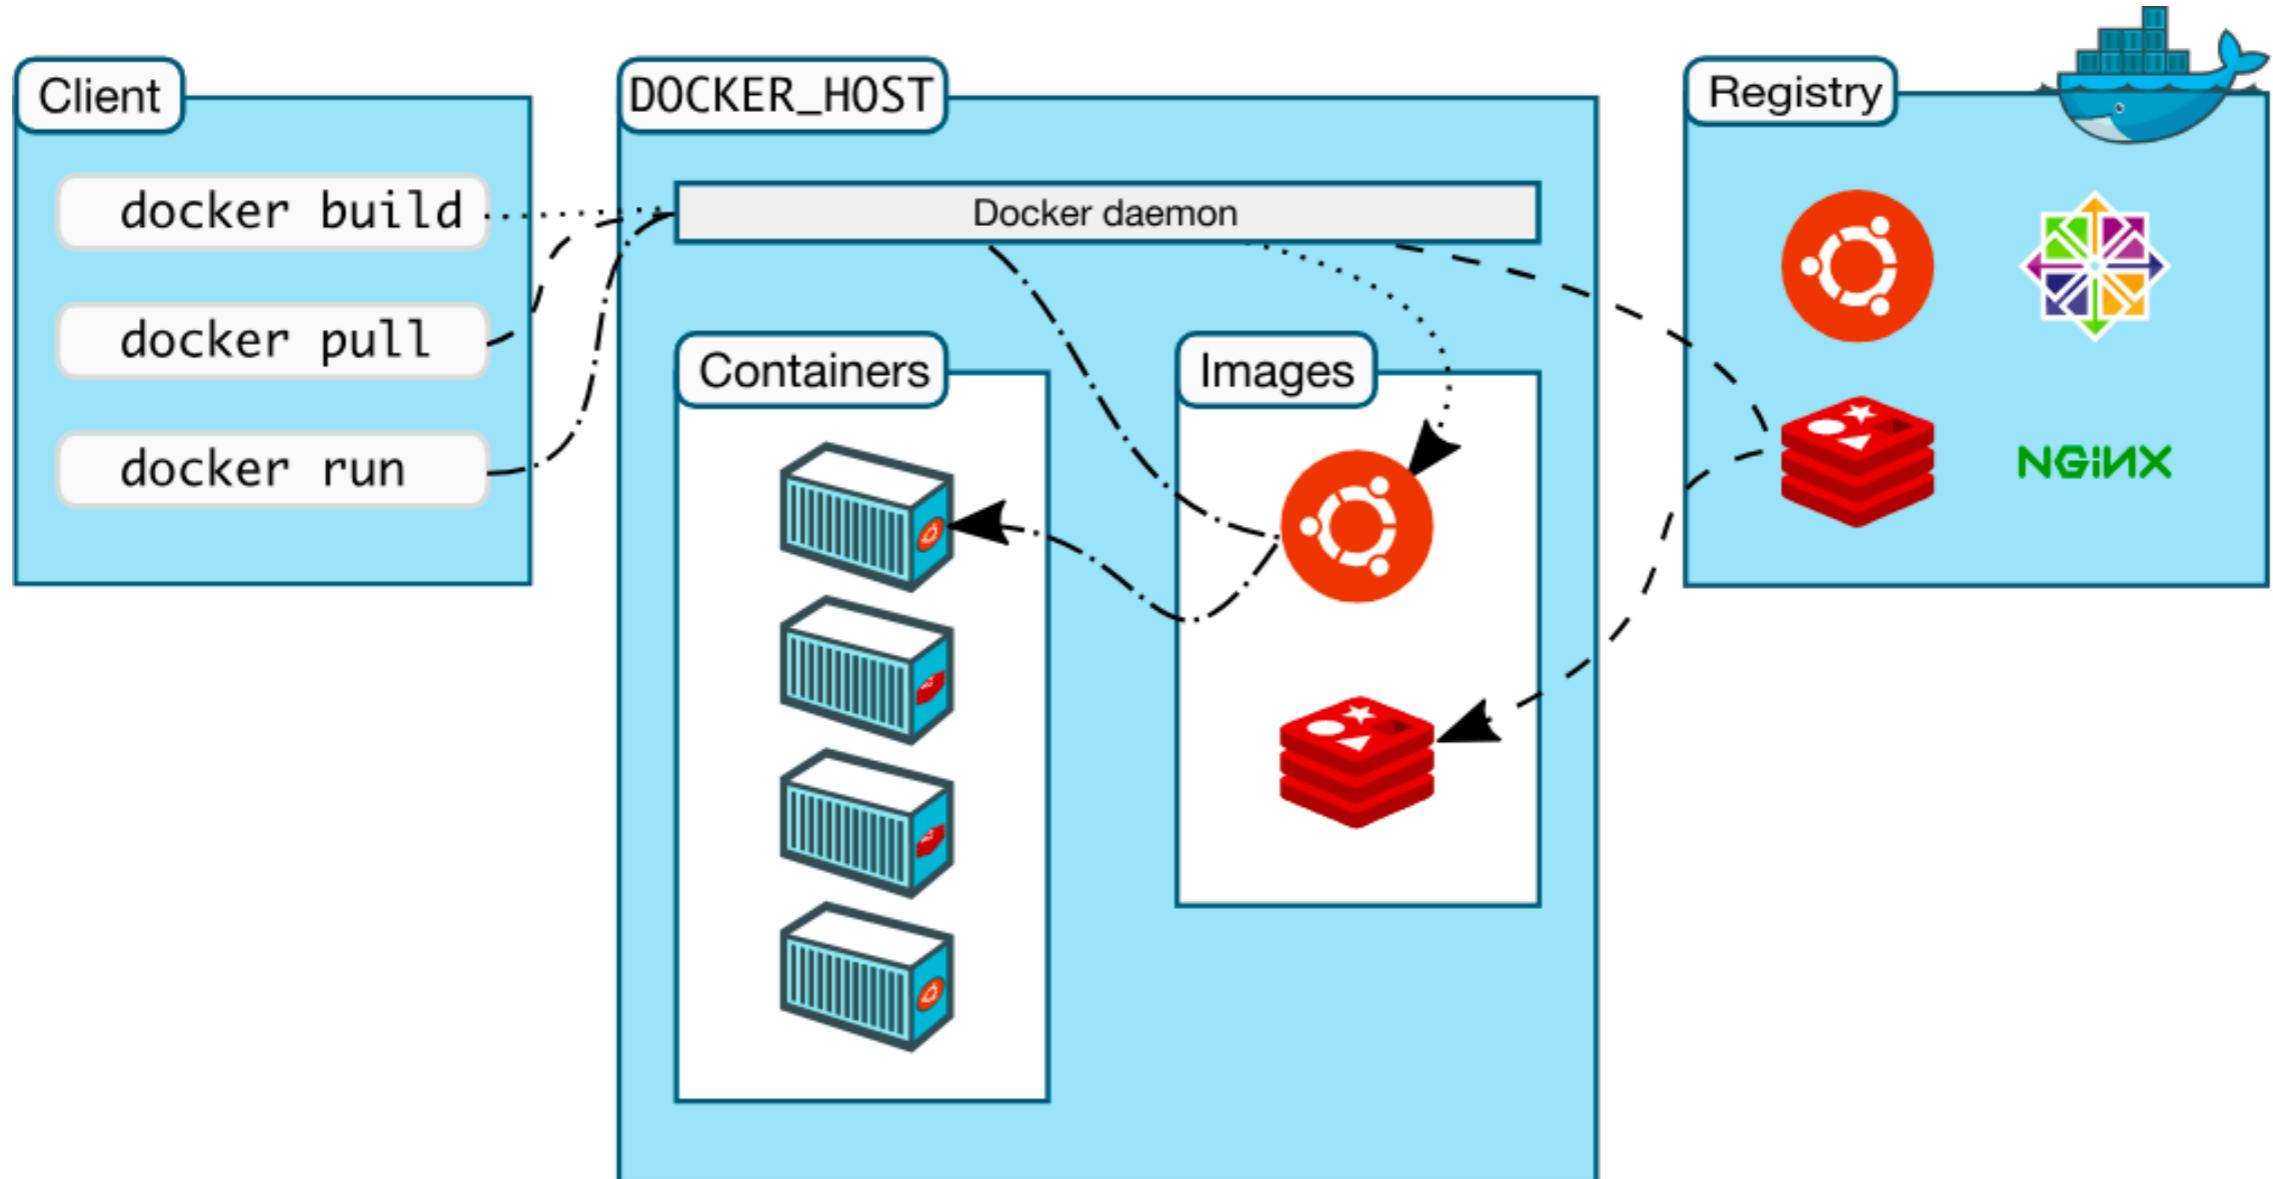
<https://github.com/up1/demo-docker-k8s>



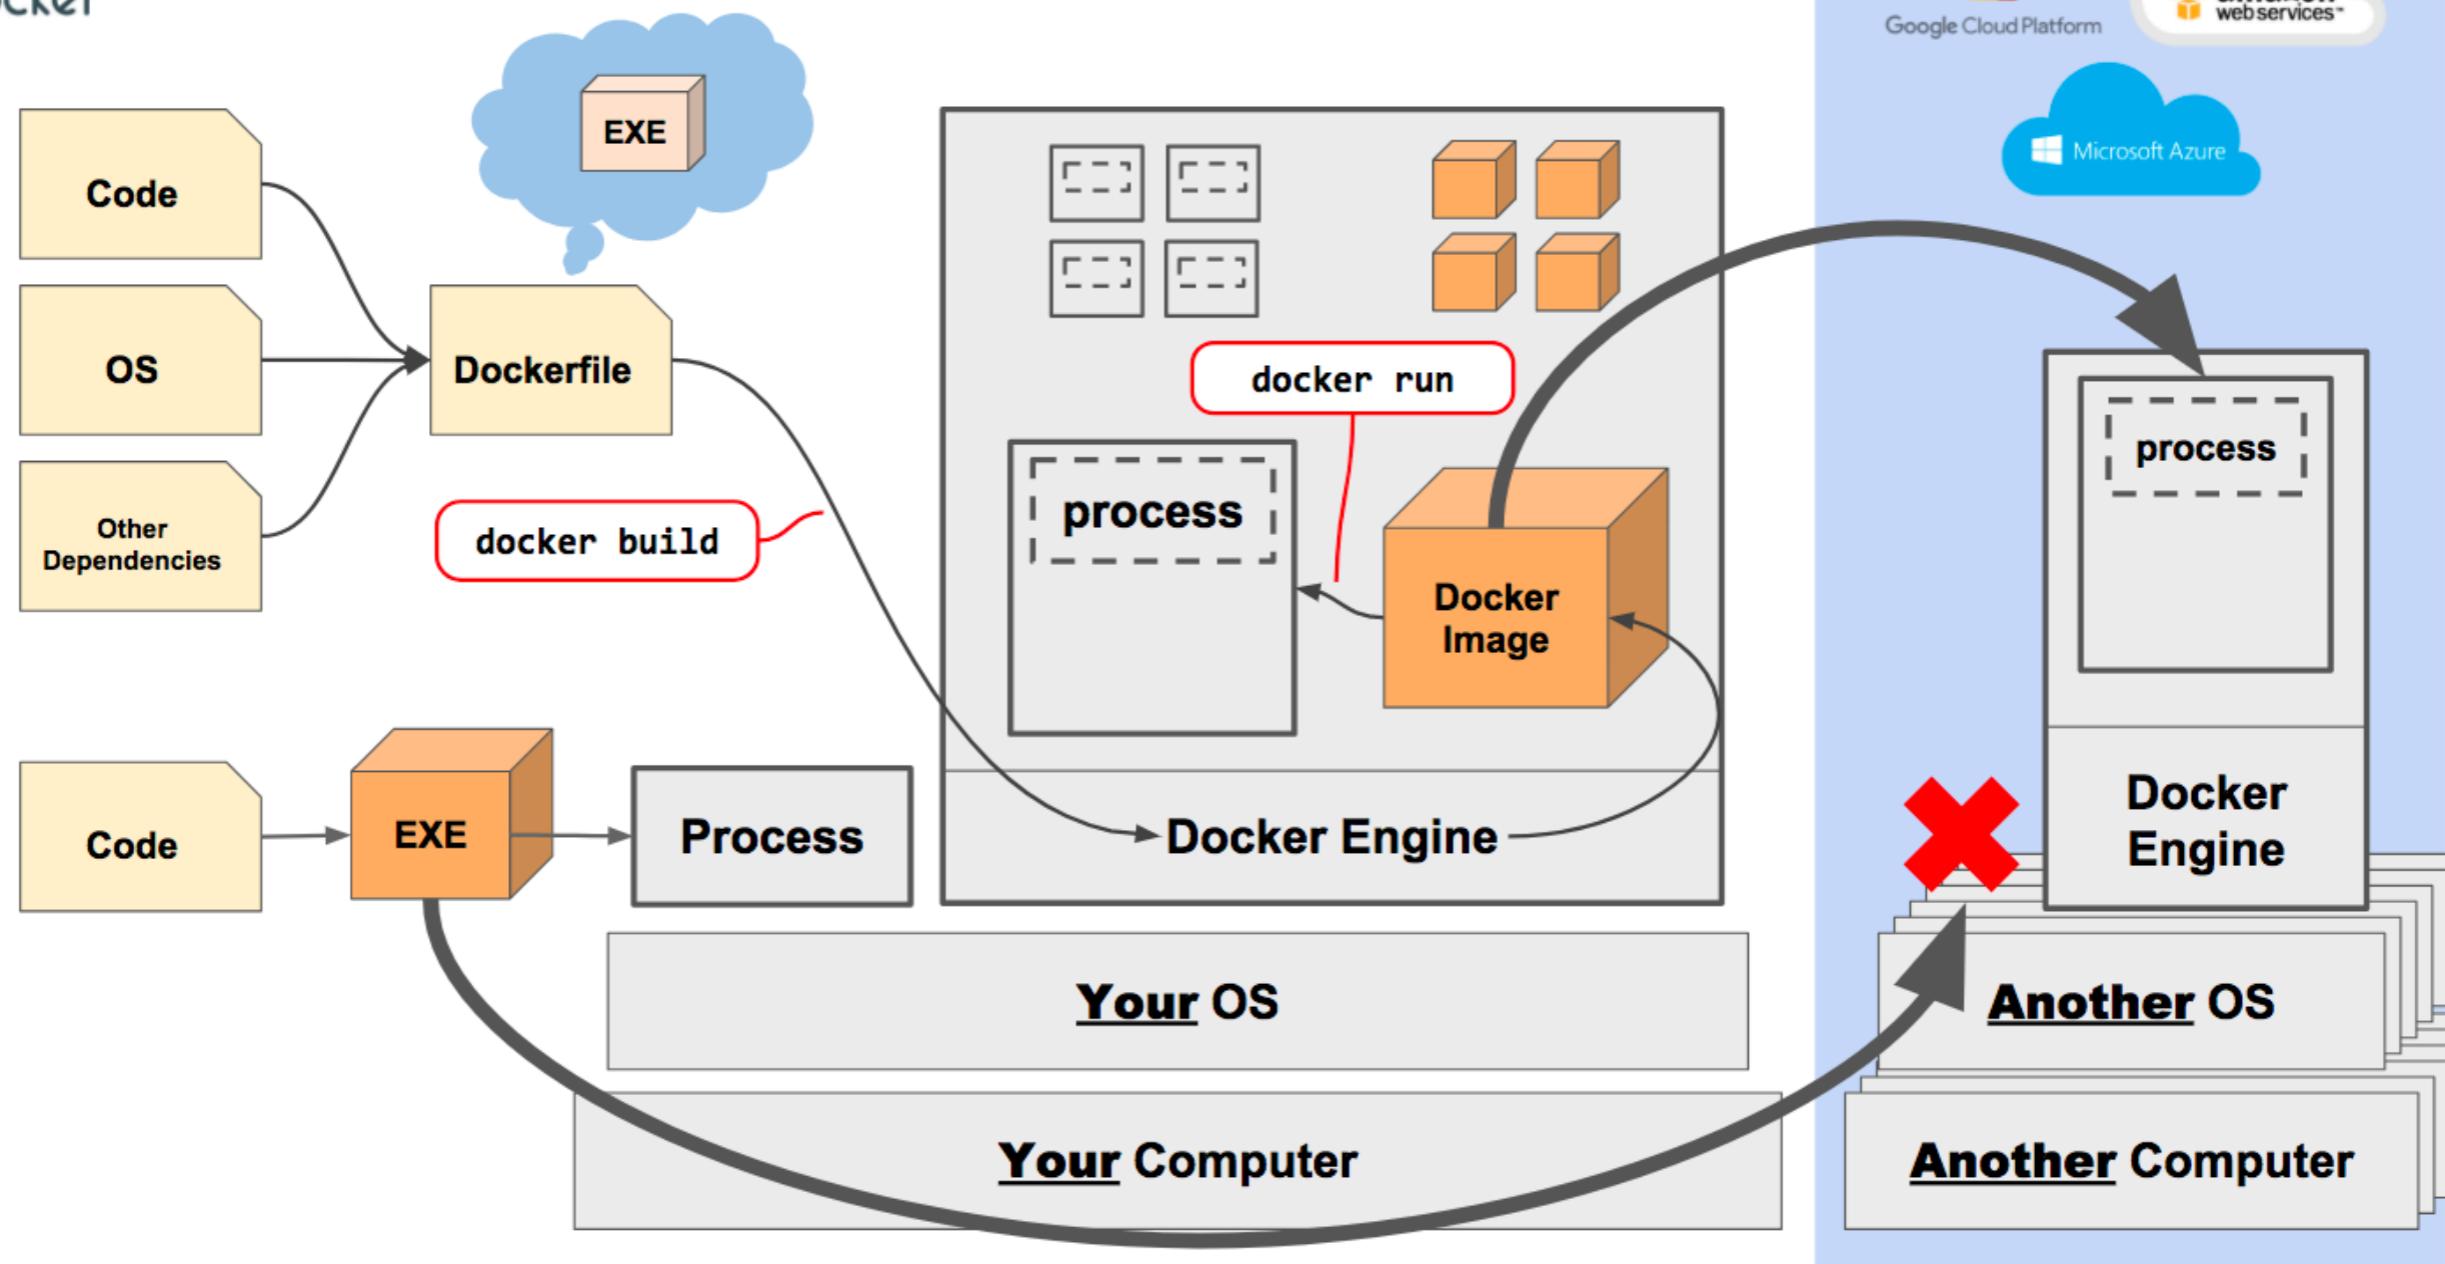
Goals



Docker overview



Docker overview



Docker command

\$docker-compose build

\$docker-compose up

\$docker-compose down



Building Docker image

From container
Dockerfile

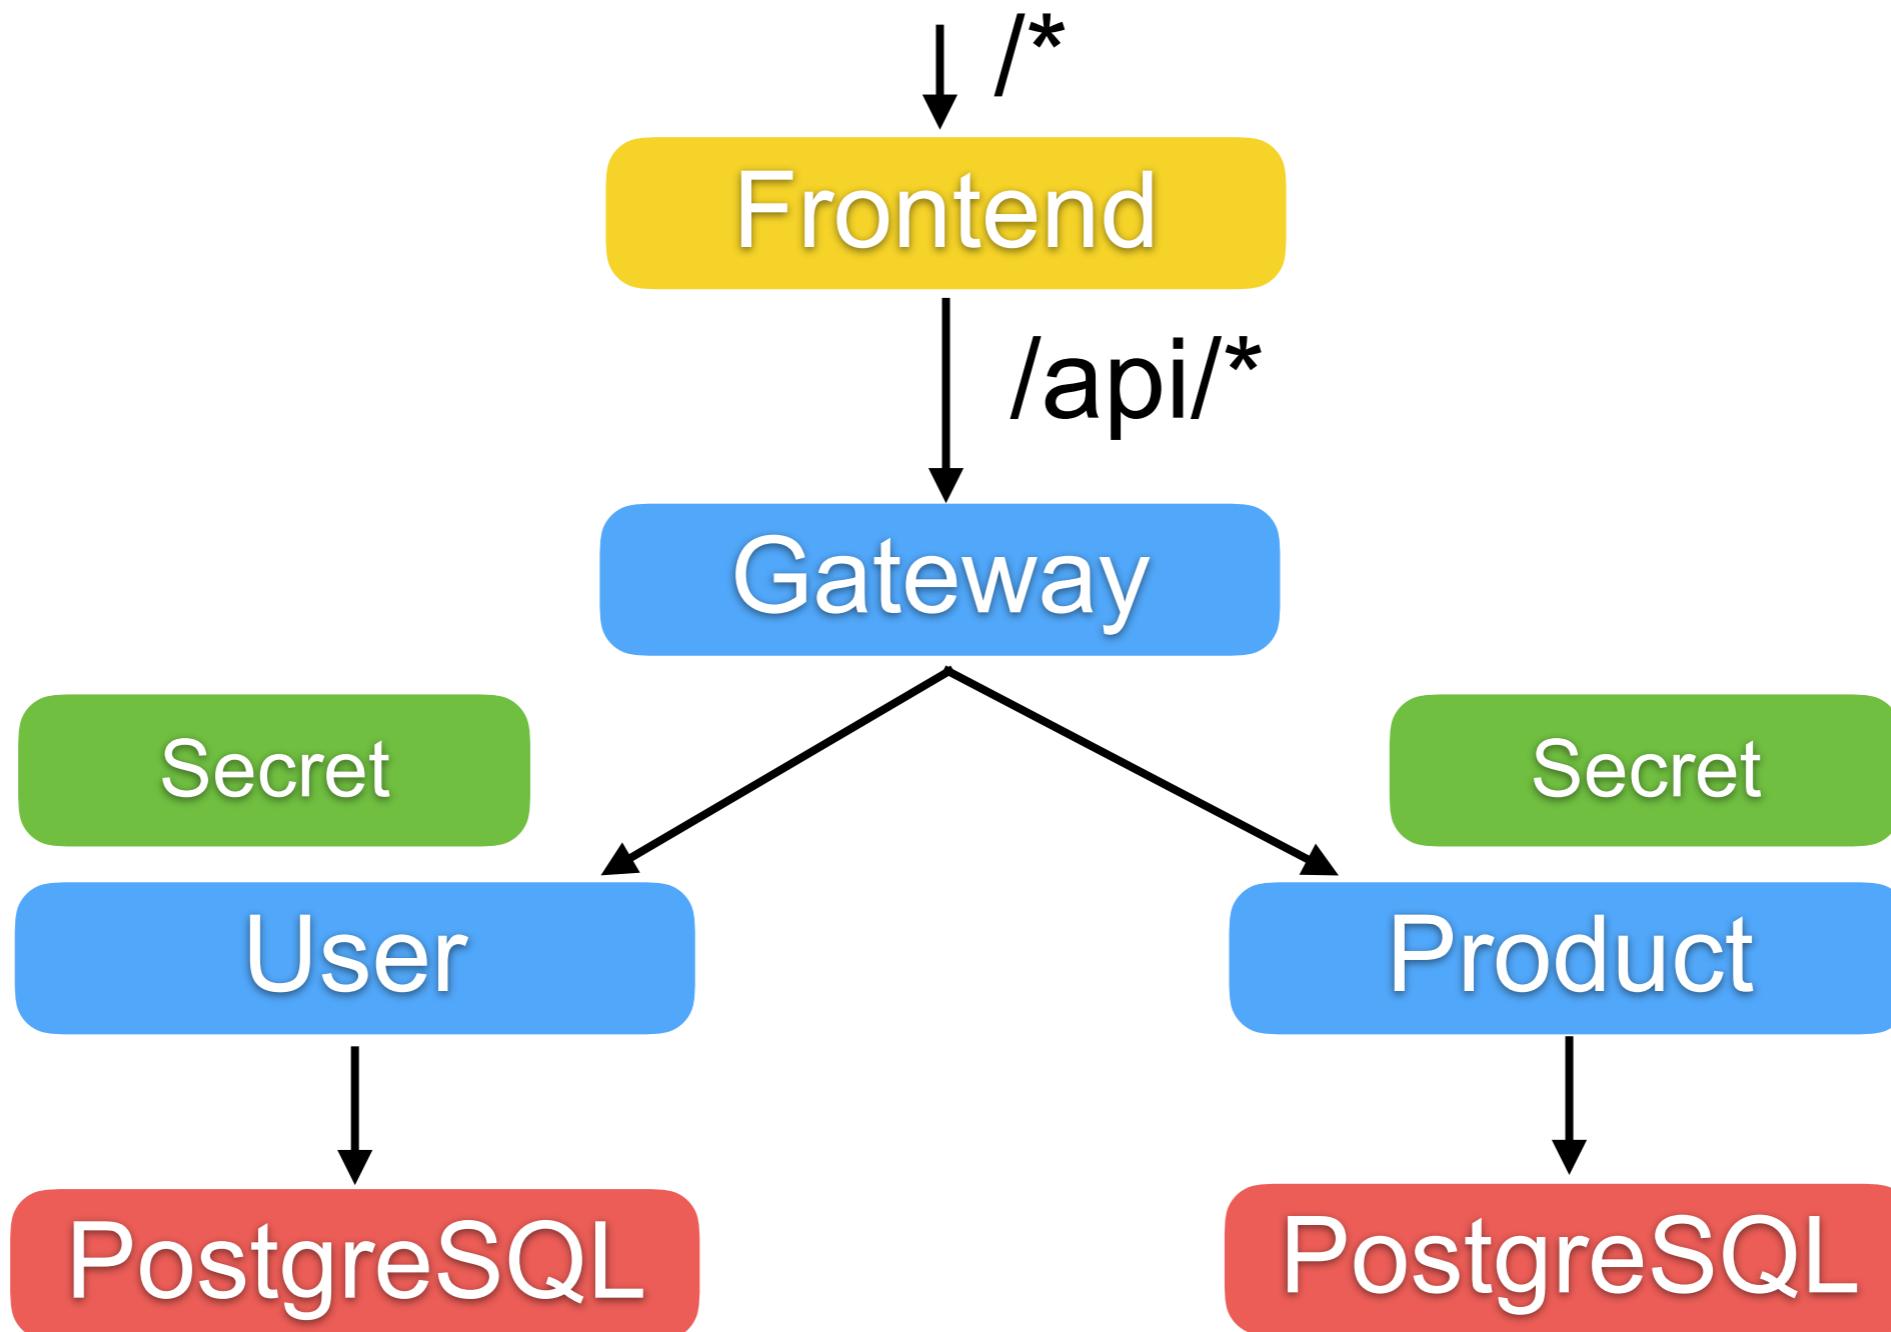


Docker compose

All-in-one with simple command
\$docker-compose build/up/down



Structure of service



Workshop

<https://github.com/up1/demo-docker-k8s>

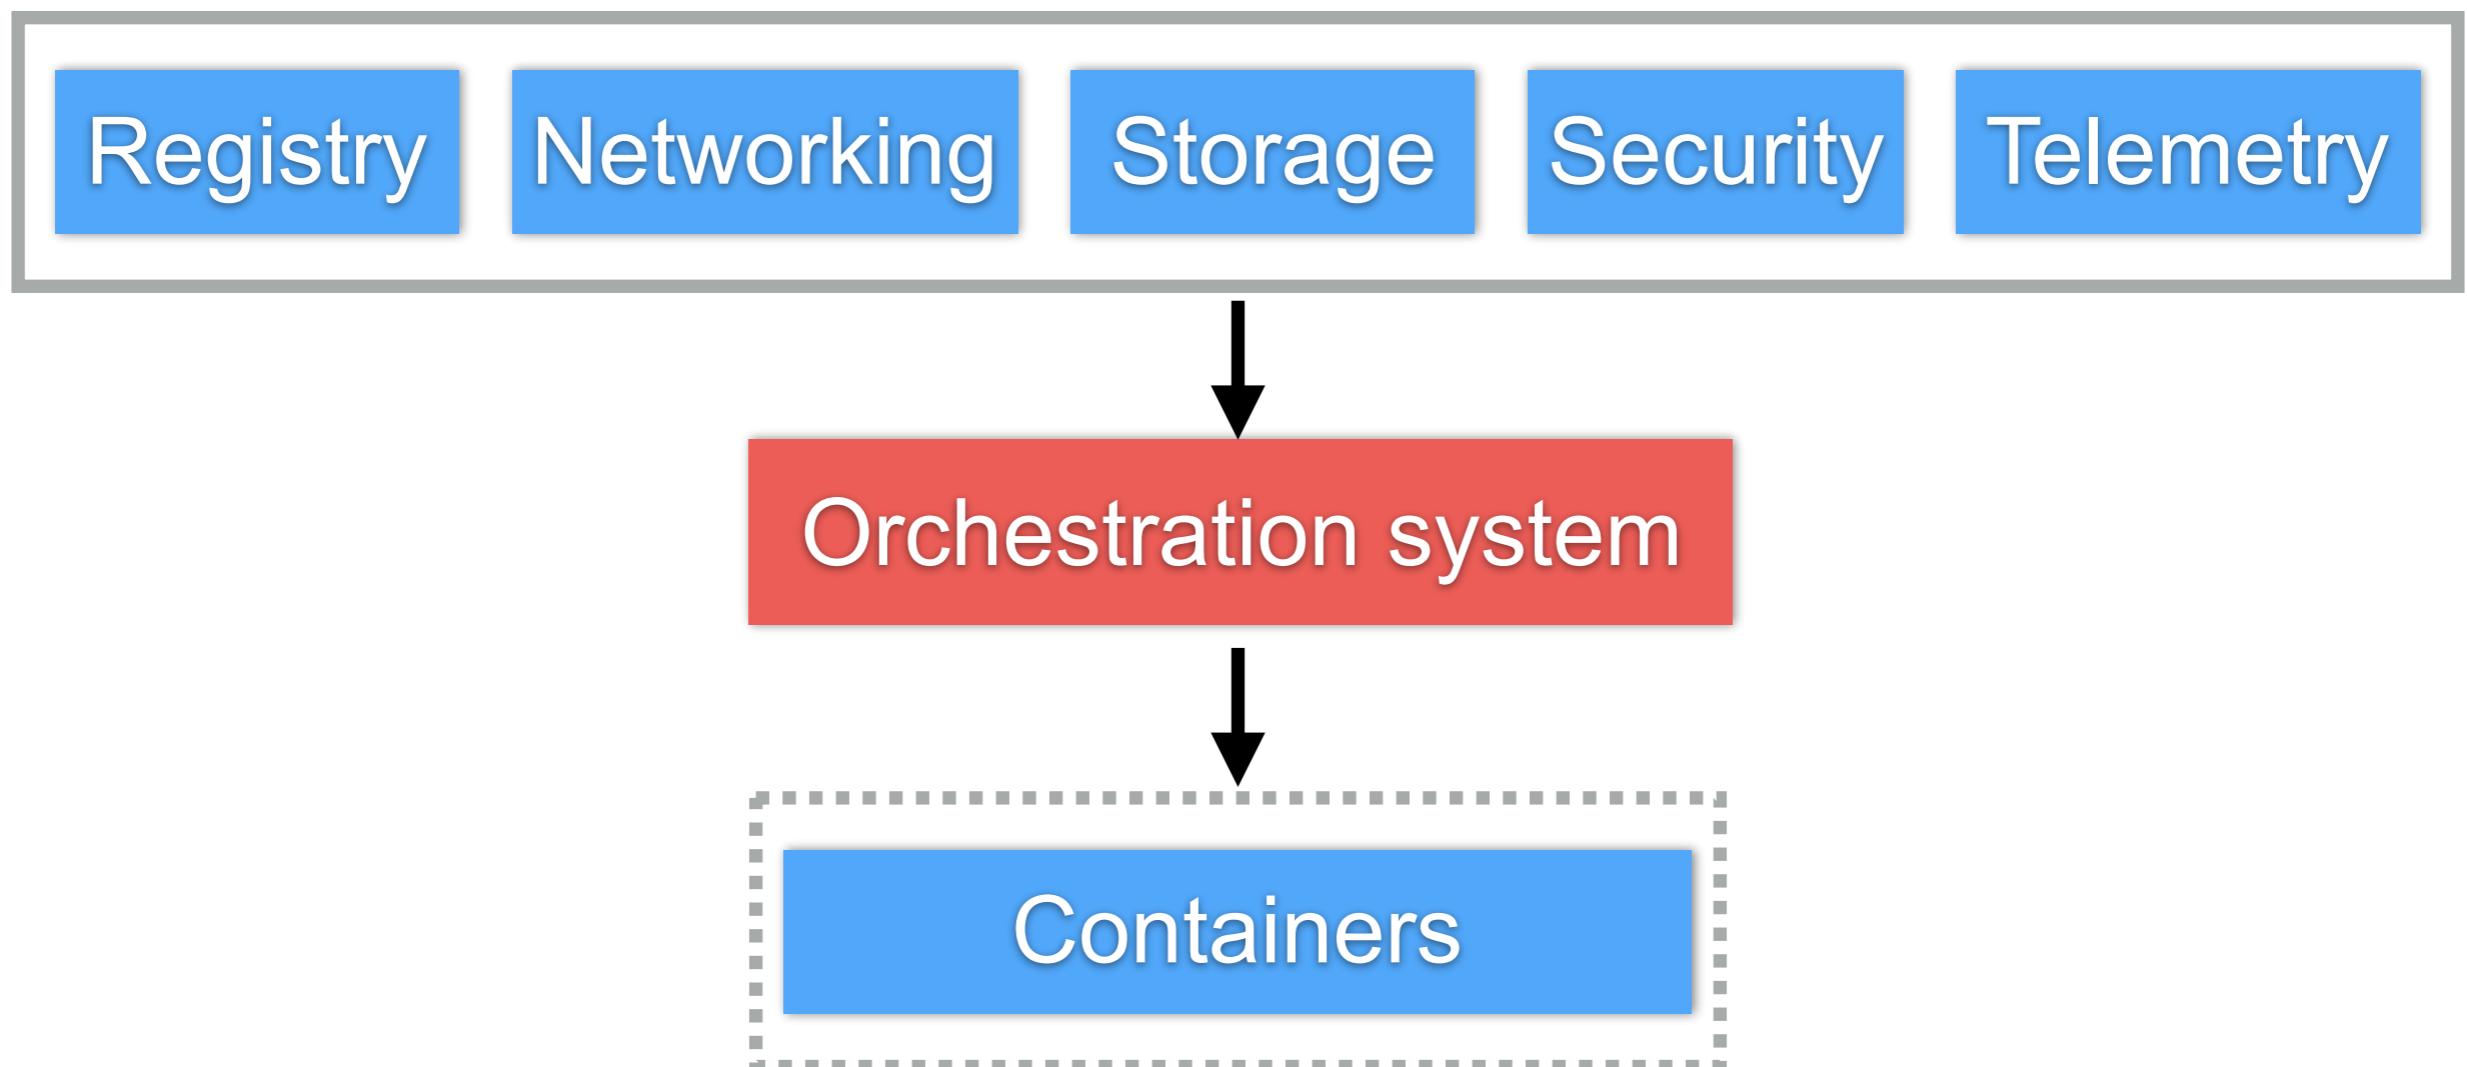


Introduction Kubernetes



Kubernetes

Open source **orchestration** for containers
Docker and rkt



Kubernetes

Provide declarative primitives for **desired state**

- Self-healing

- Horizontal scaling

- Automatic binpacking

- Service discovery

- Load balancing



Benefits

Deploy application quickly and predictable

Scaling in the fly

Release new feature seamless

Fail-proofing

Limit hardware usage

Agile software development

Portability between OS, host and cloud provider



Kubernetes concepts

Pods

Replication controller

Deployment/ReplicaSet

Service

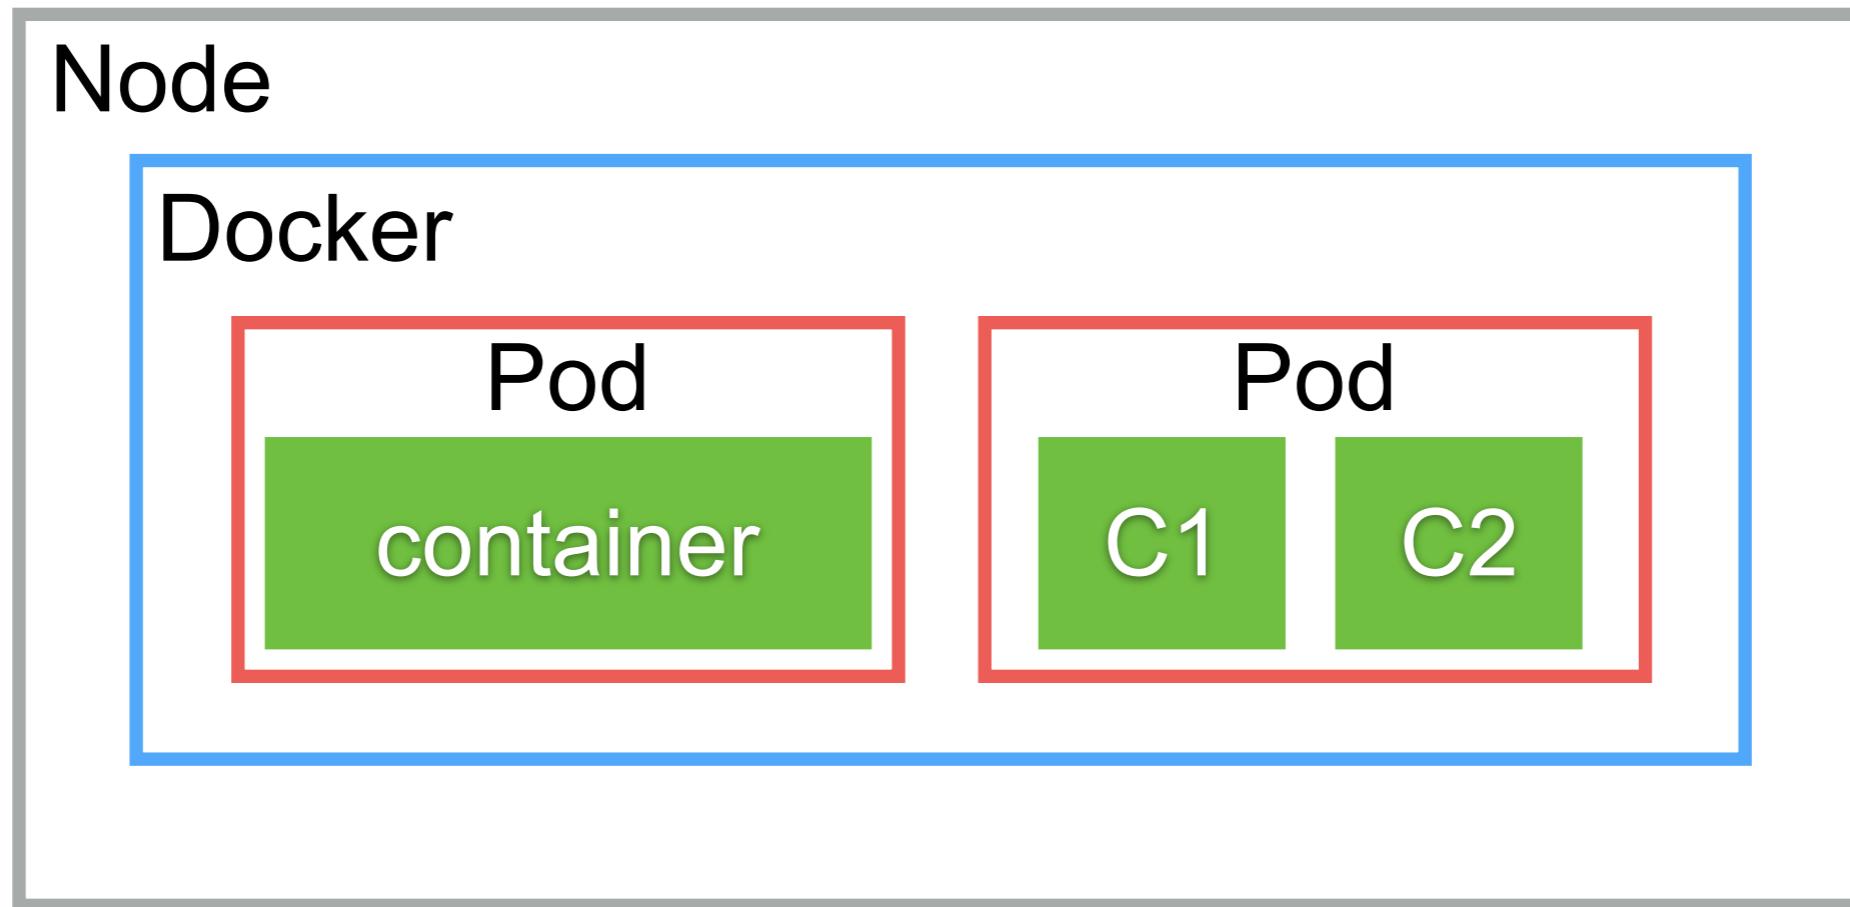
Label

Namespace



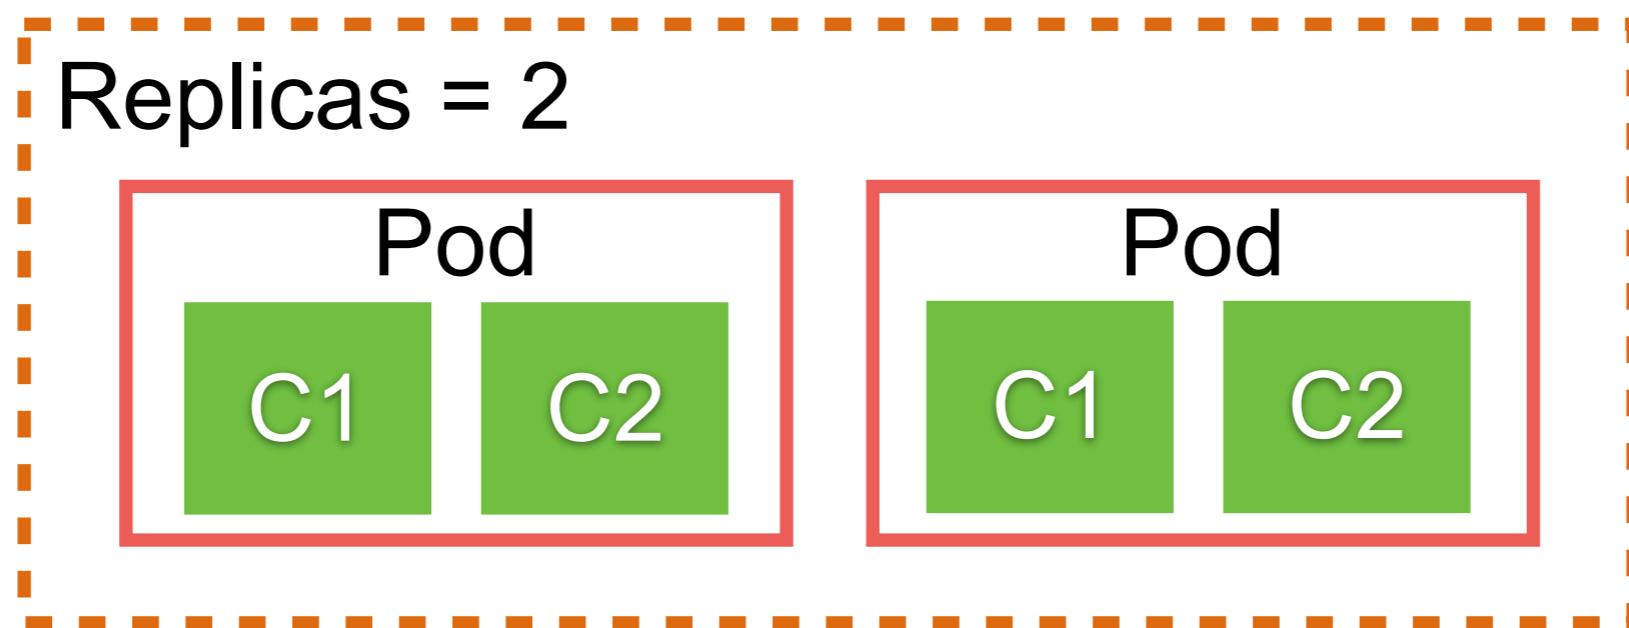
Pods

Colocated group of containers
Share IP, storage and namespace



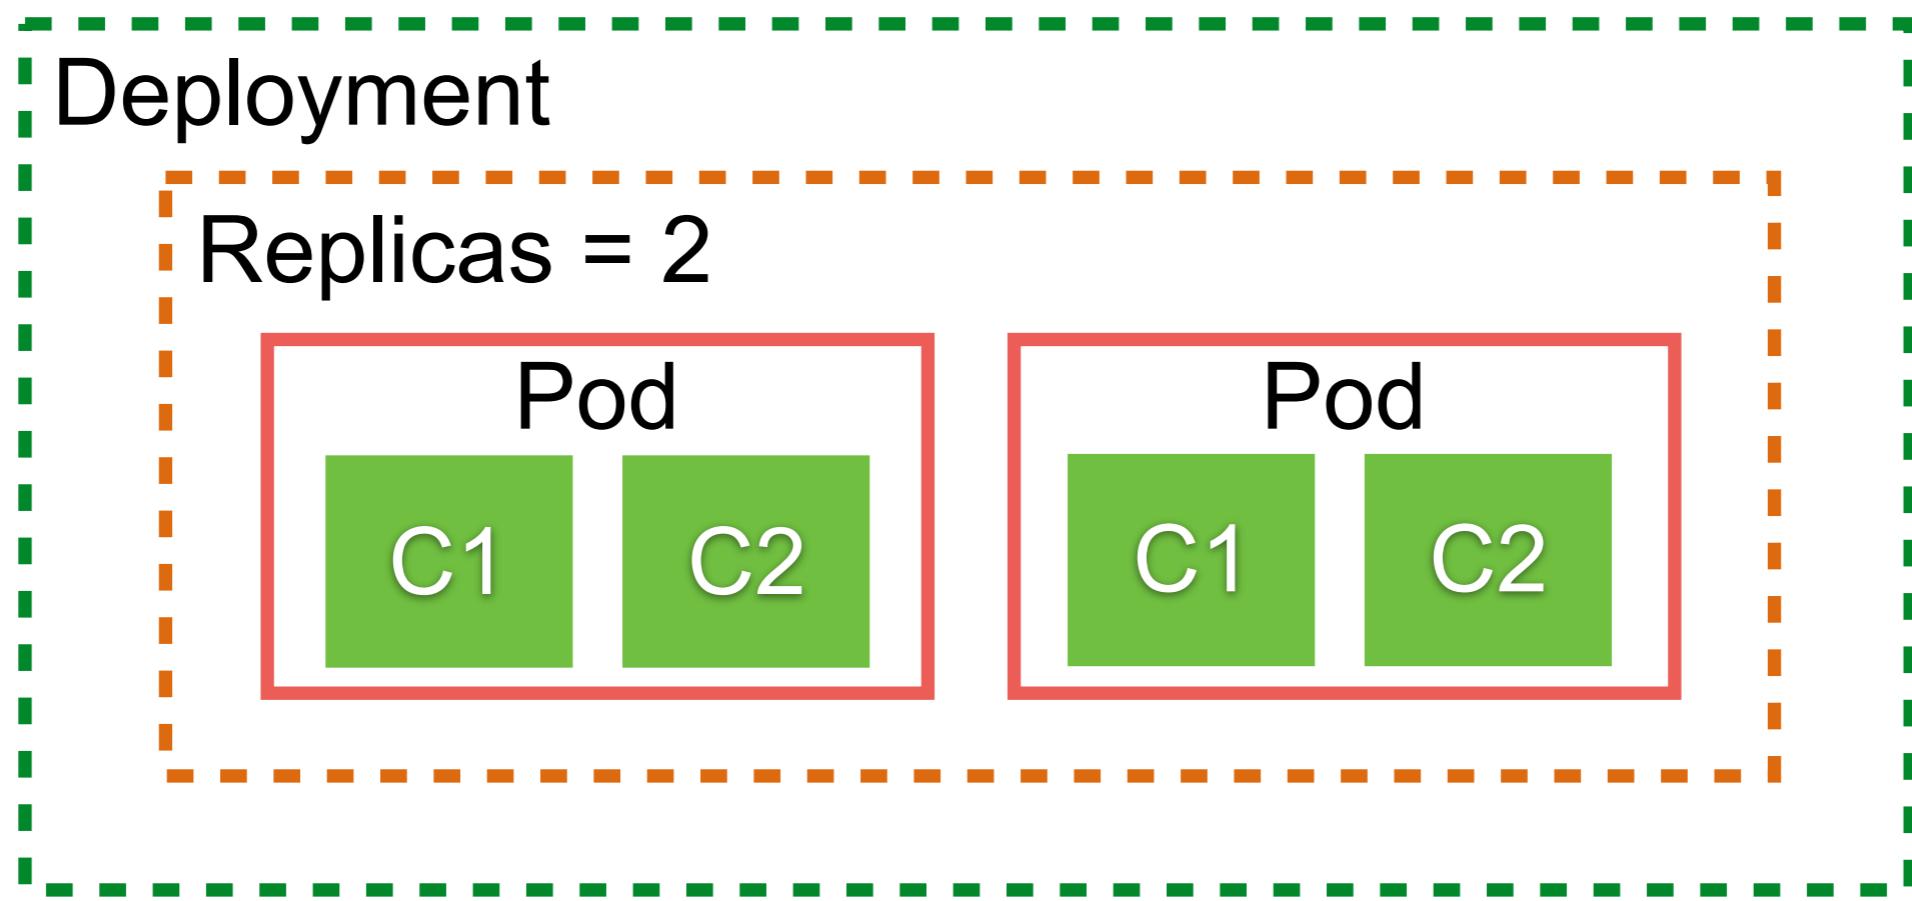
Replication Controller (RC)

Manage the lifecycle of Pod
Ensure specified number are running



ReplicaSet (RS)

Next generation of Replication Controller
Generally create with Deployment
Enable Horizontal Pod Autoscaling (HPA)



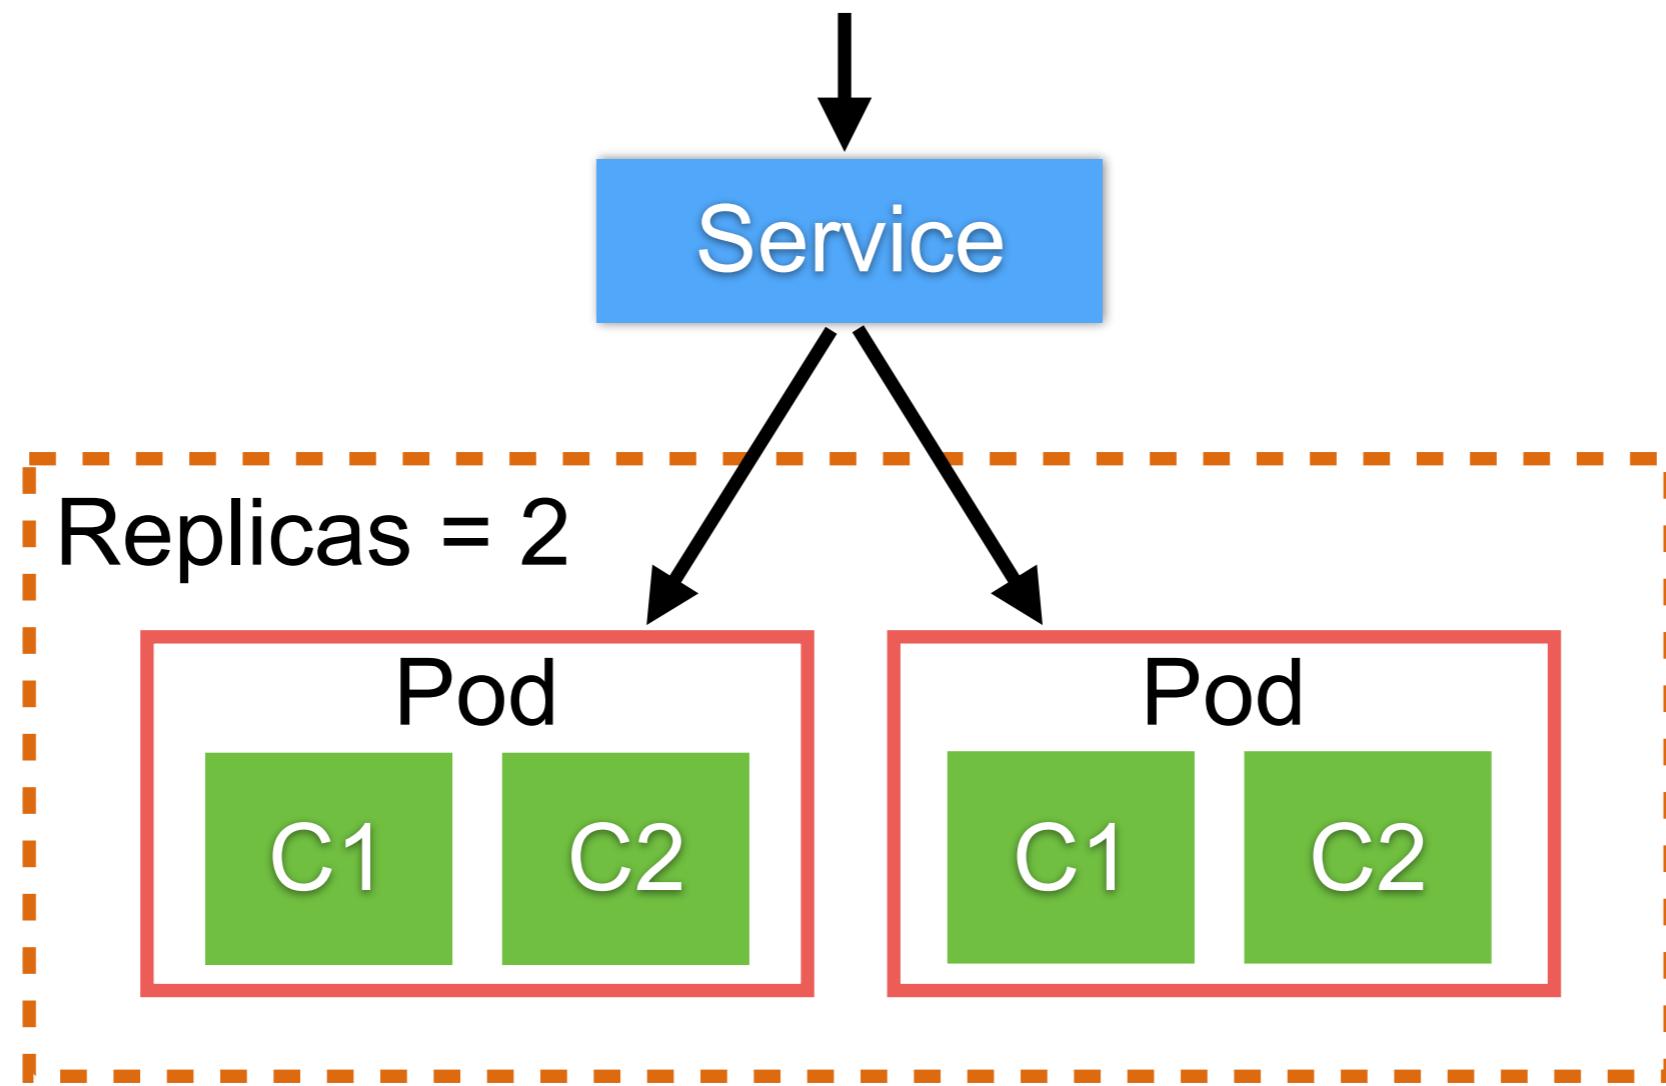
Deployment

Declarative update for Pod and RS
Rolling update
Rollback to previous version



Service

Single and stable name of the set of Pods
Act as a Load Balancer



Publish services

ClusterIP

NodePort

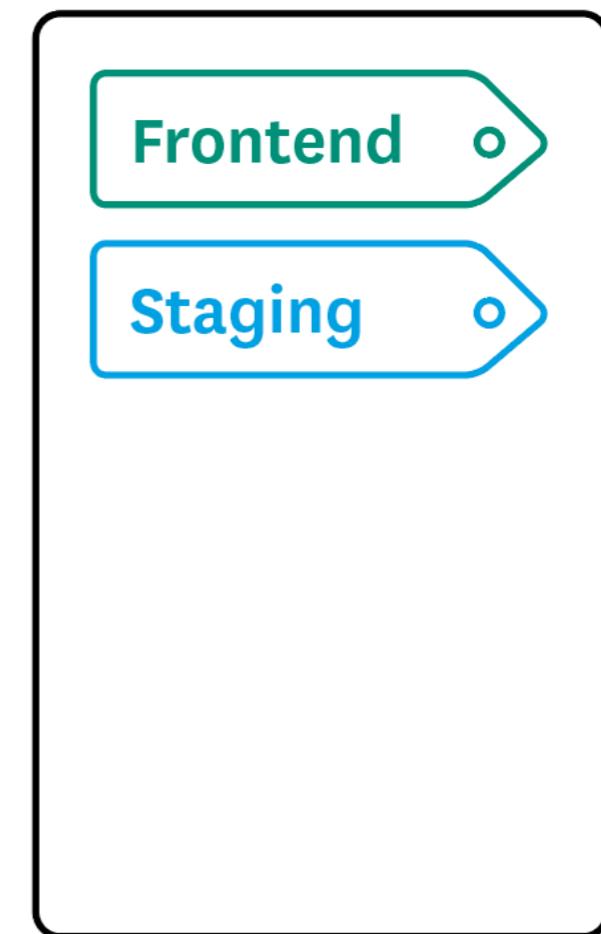
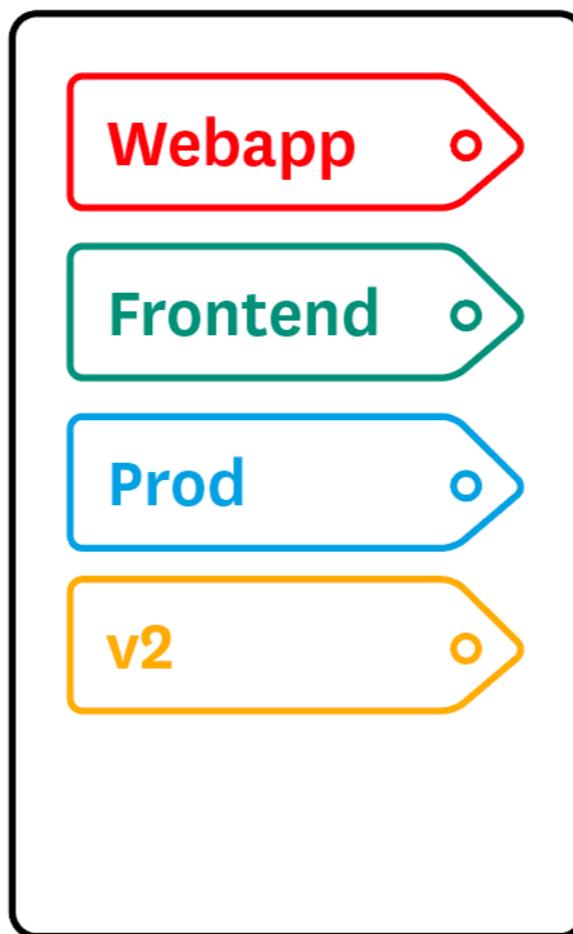
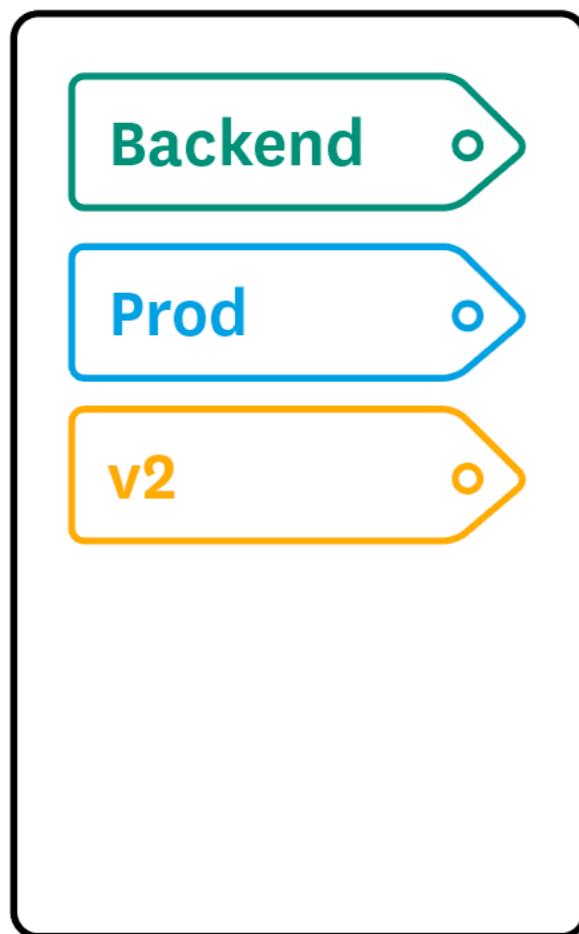
LoadBalancer

ExternalName



Label

Use to organize and select group of objects



Pod 1

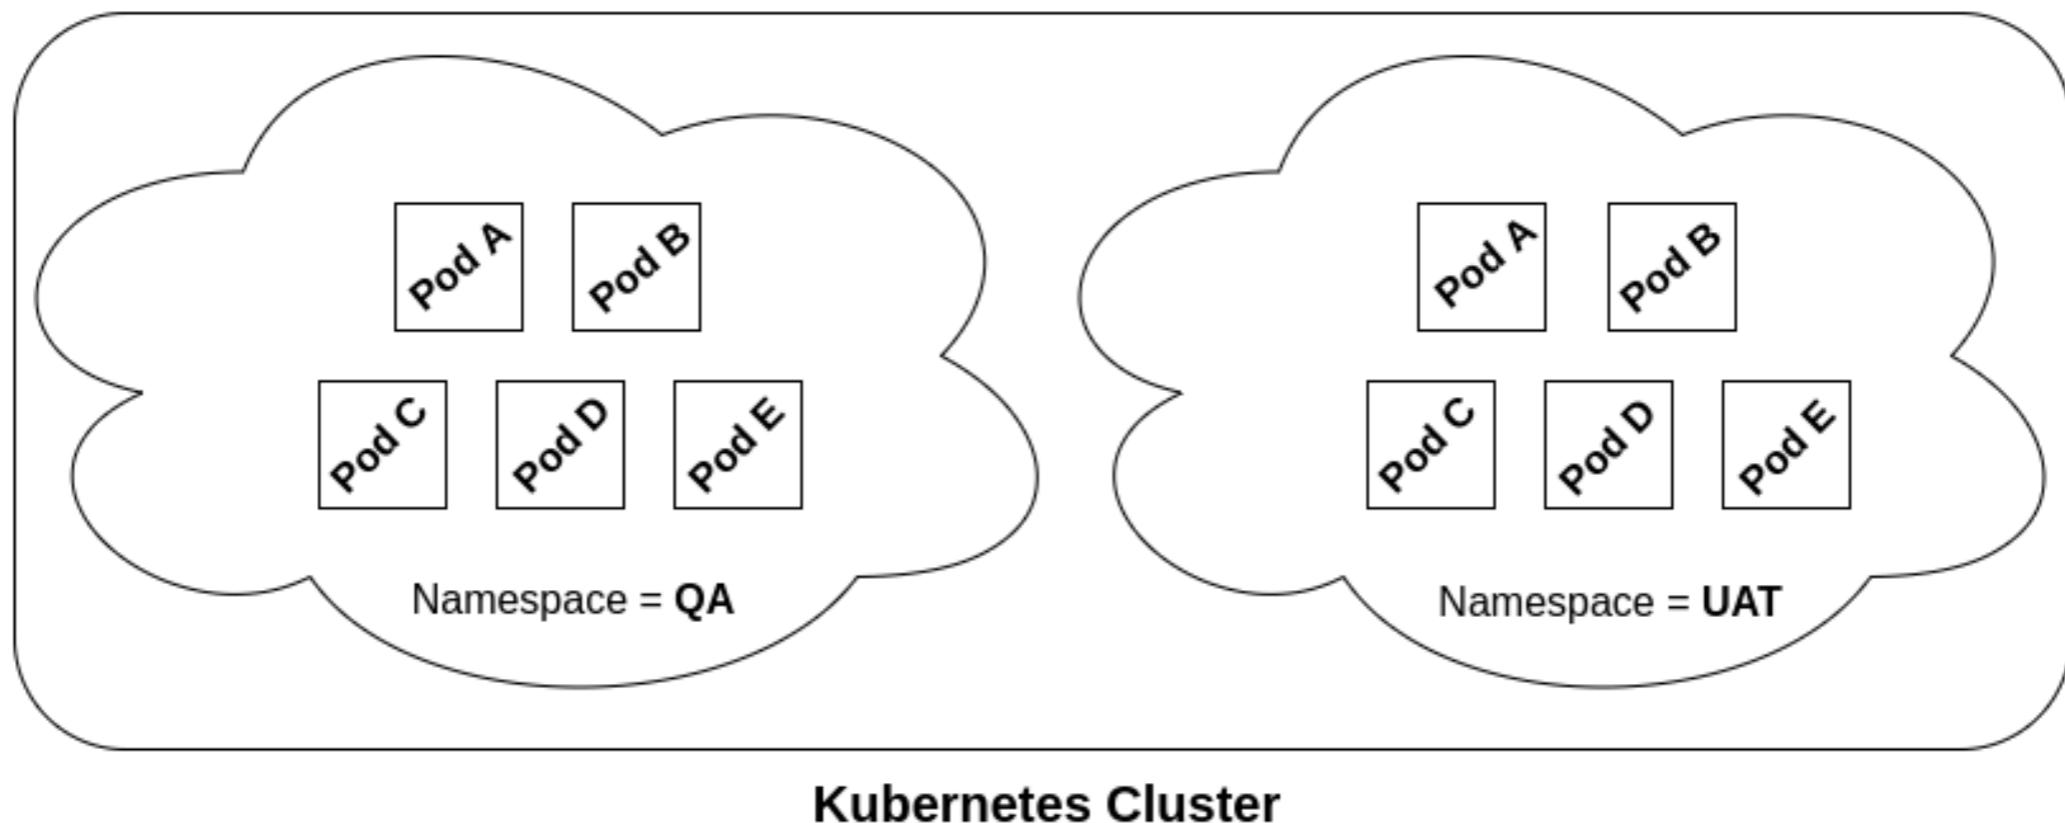
Pod 2

Pod 3



Namespace

Namespace sets up virtual clusters over the same physical cluster

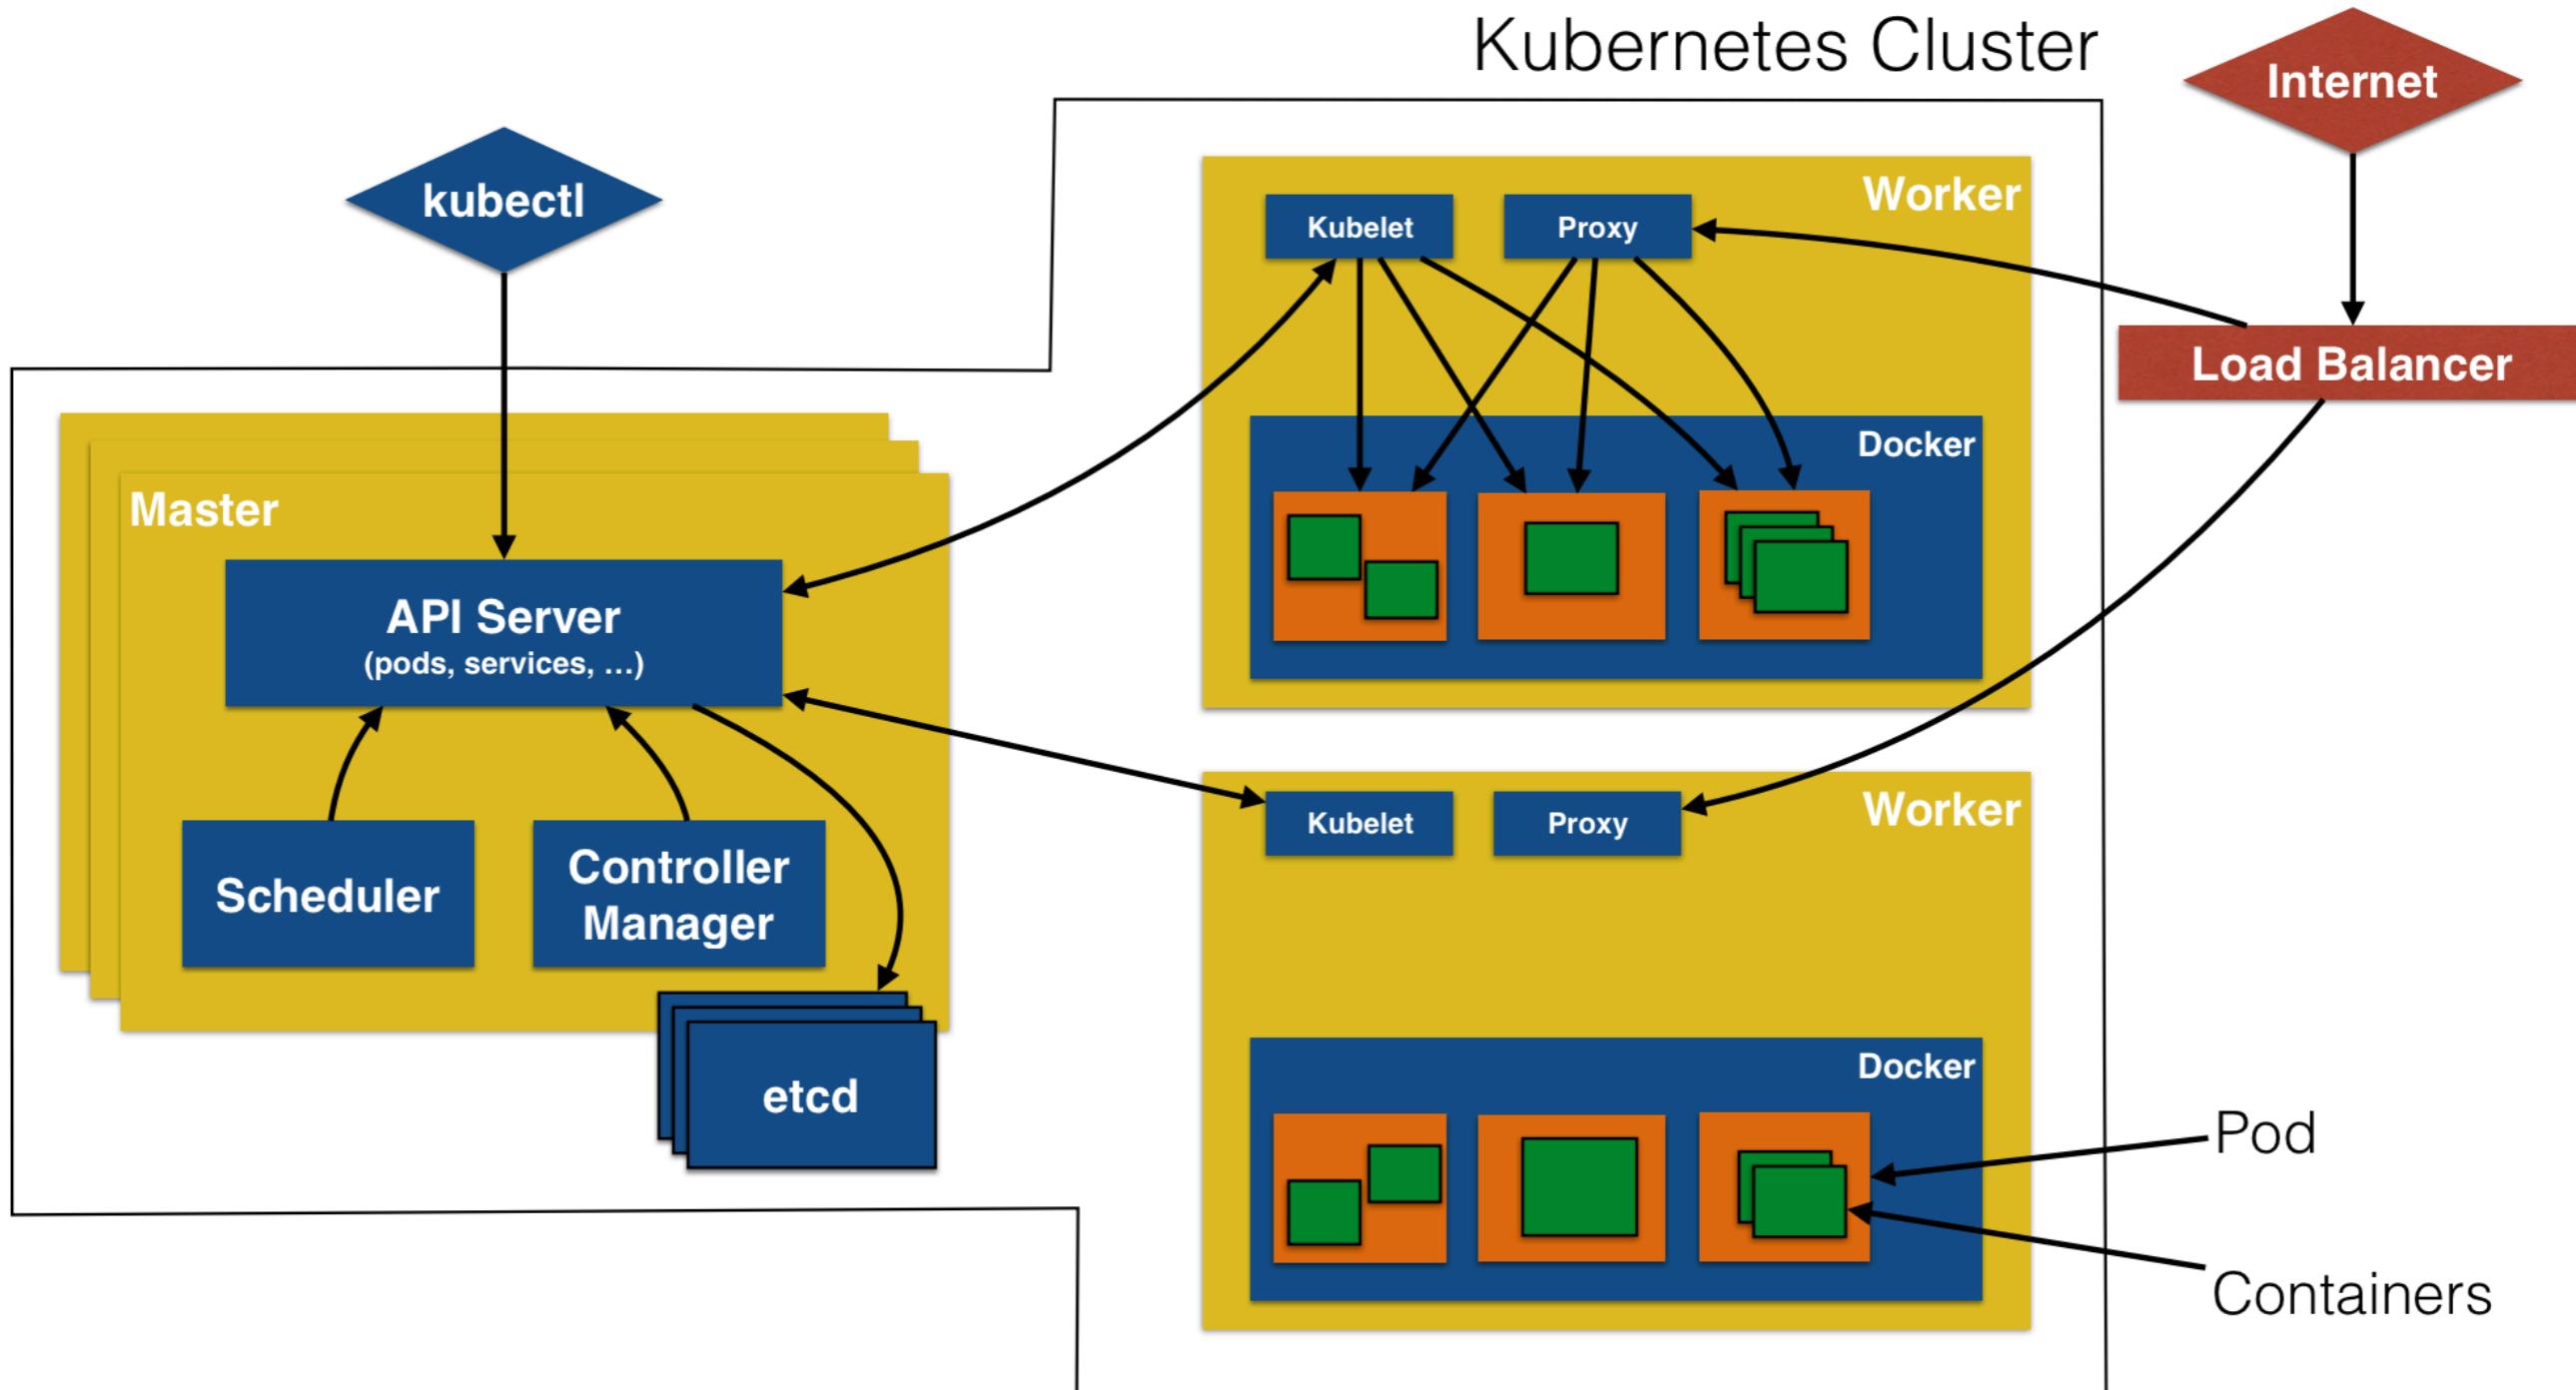


Kubernetes components

Node
Master node
Worker node



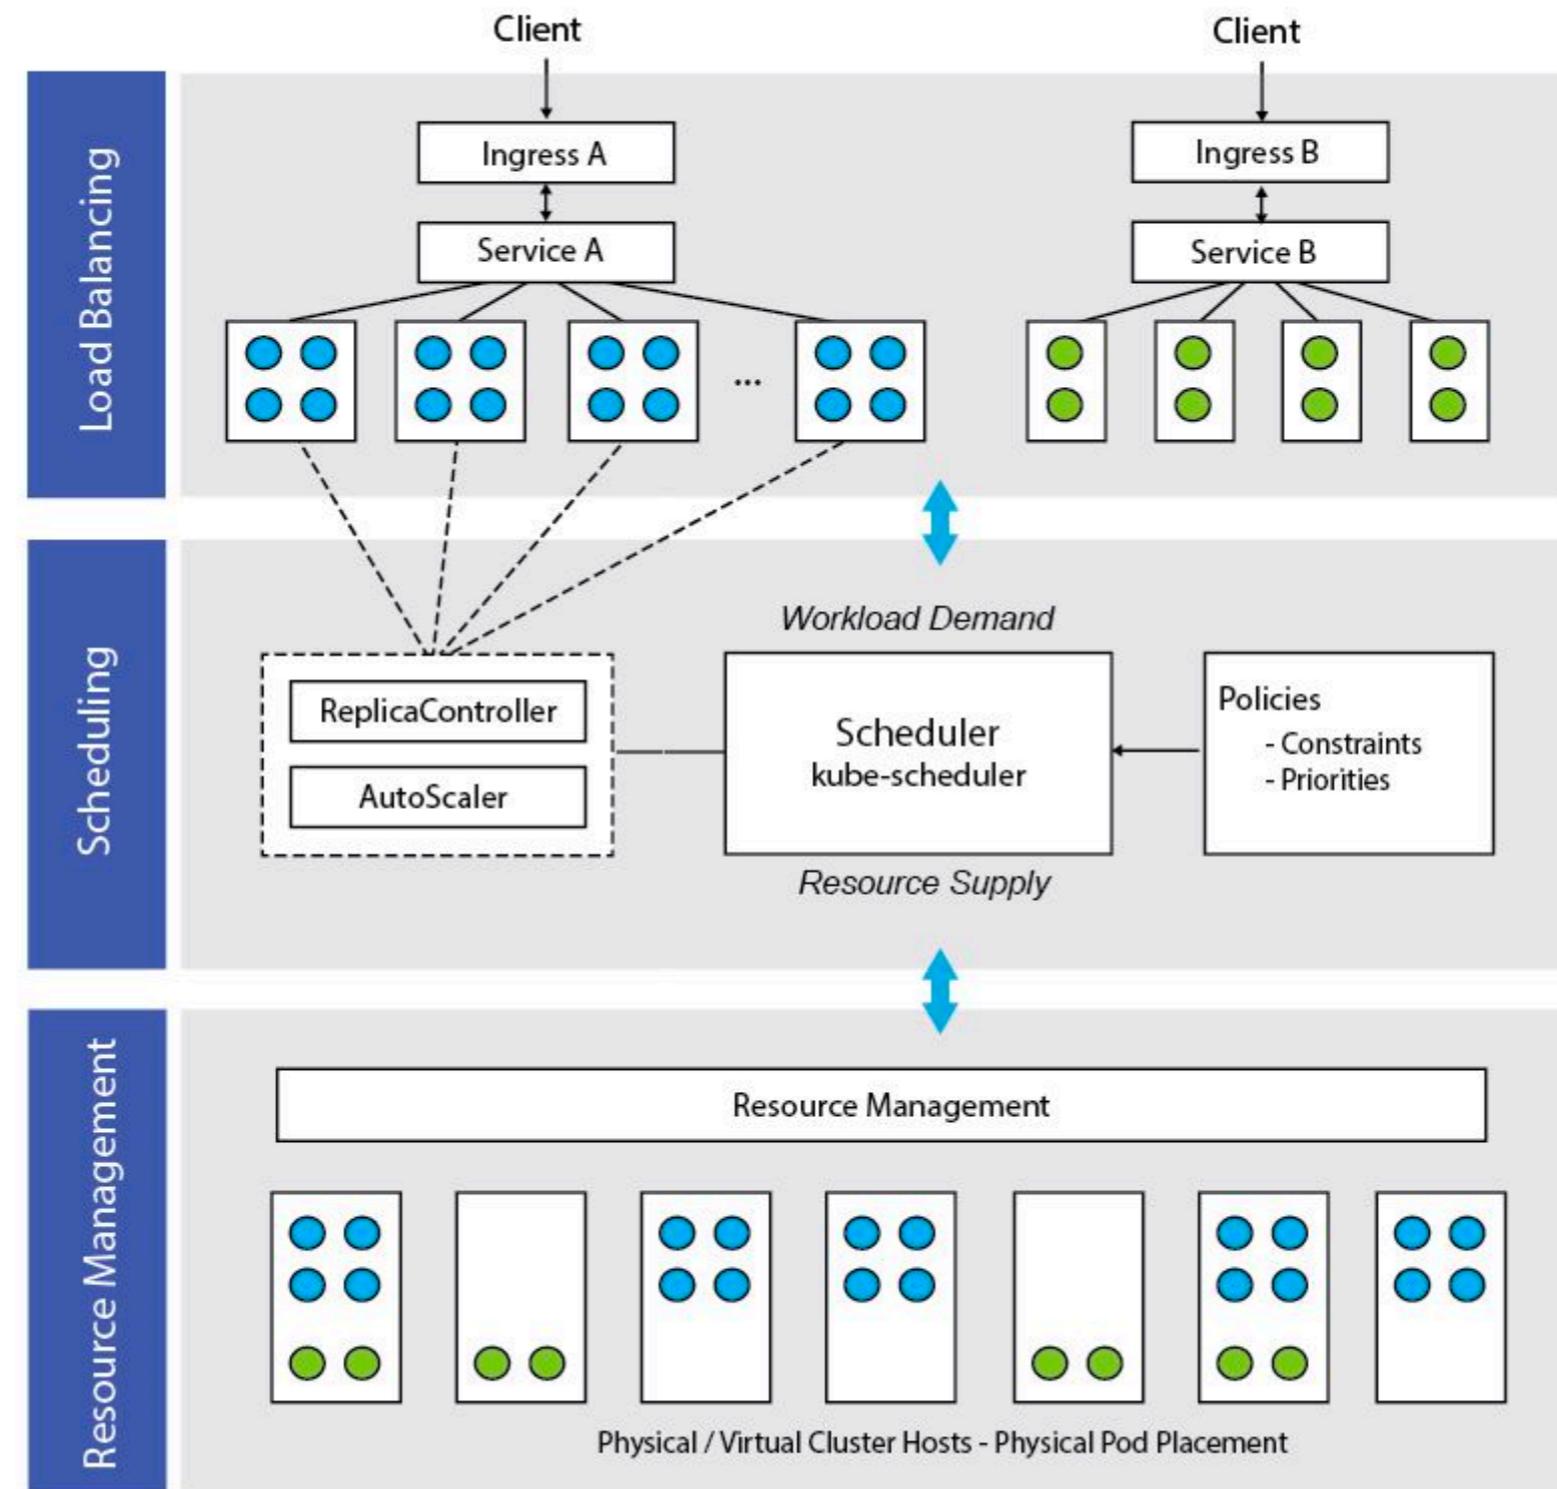
Kubernetes components



6



Kubernetes Container Orchestration



Kubernetes Scheduling Algorithm



kubectl

Command-line

Control the Kubernetes cluster manager

CRUD Kubernetes resources



kubectl

\$kubectl create

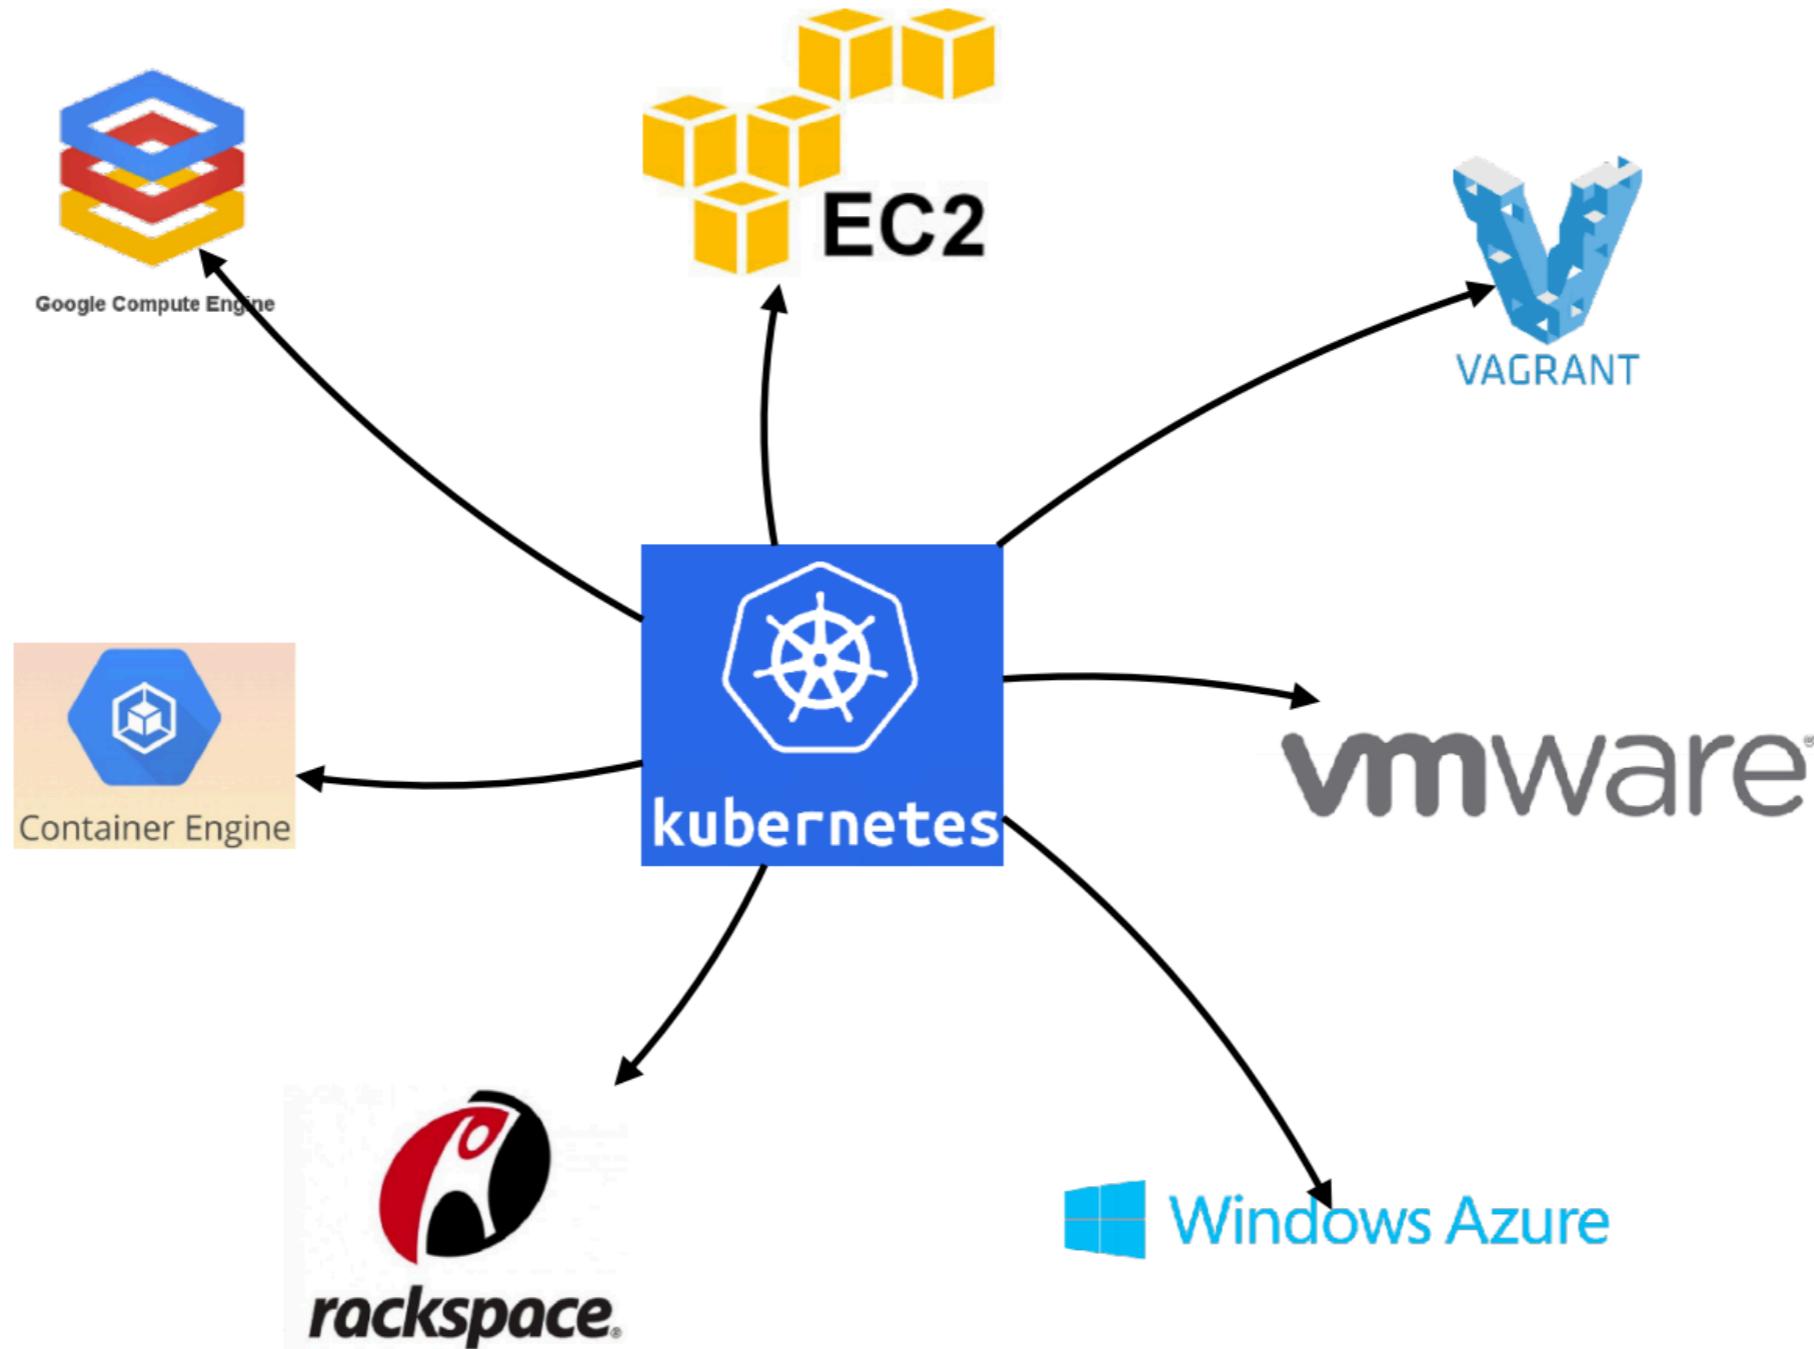
\$kubectl get

\$kubectl describe

\$kubectl delete



Kubernetes on Cloud



Let's start on Google Cloud

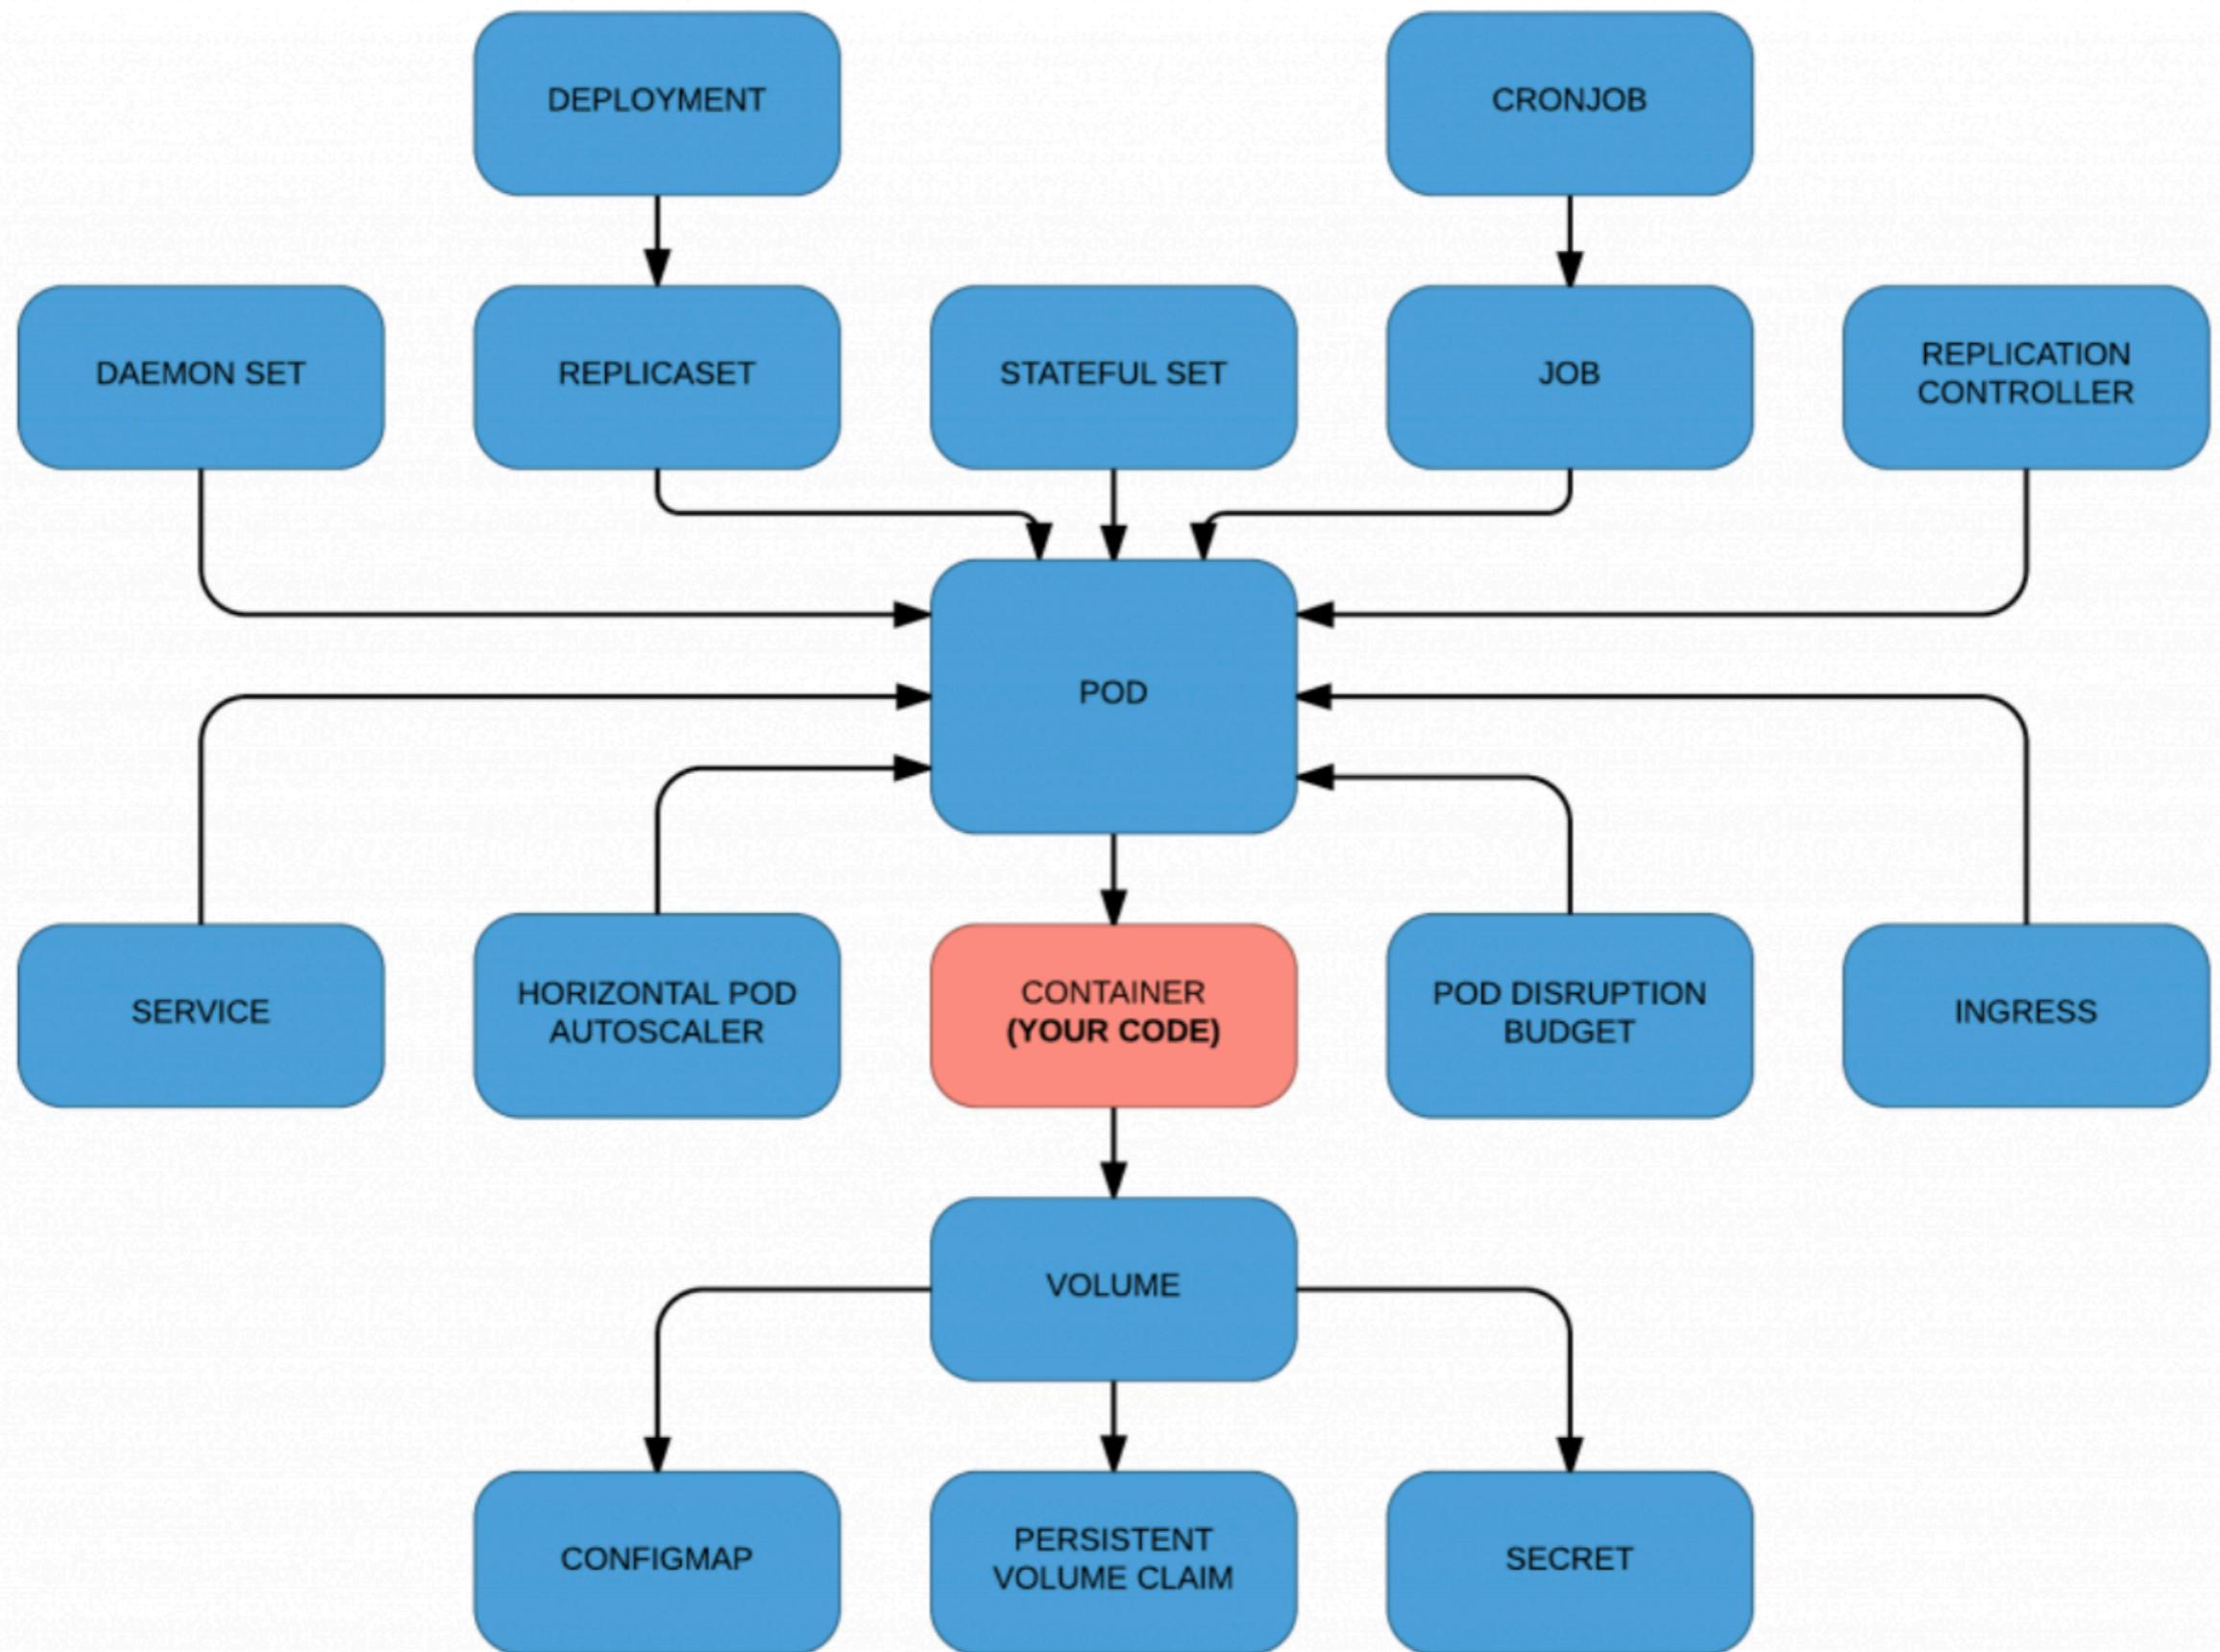


Workshop



Components



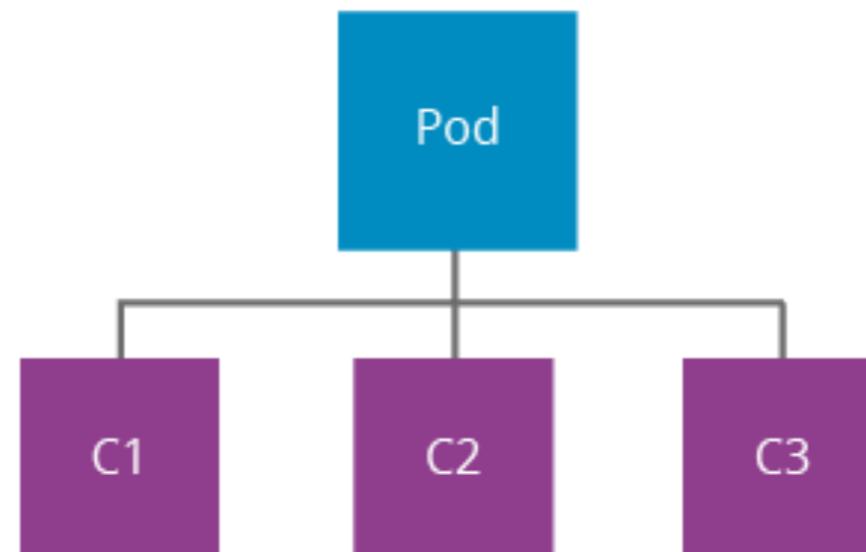


Pods



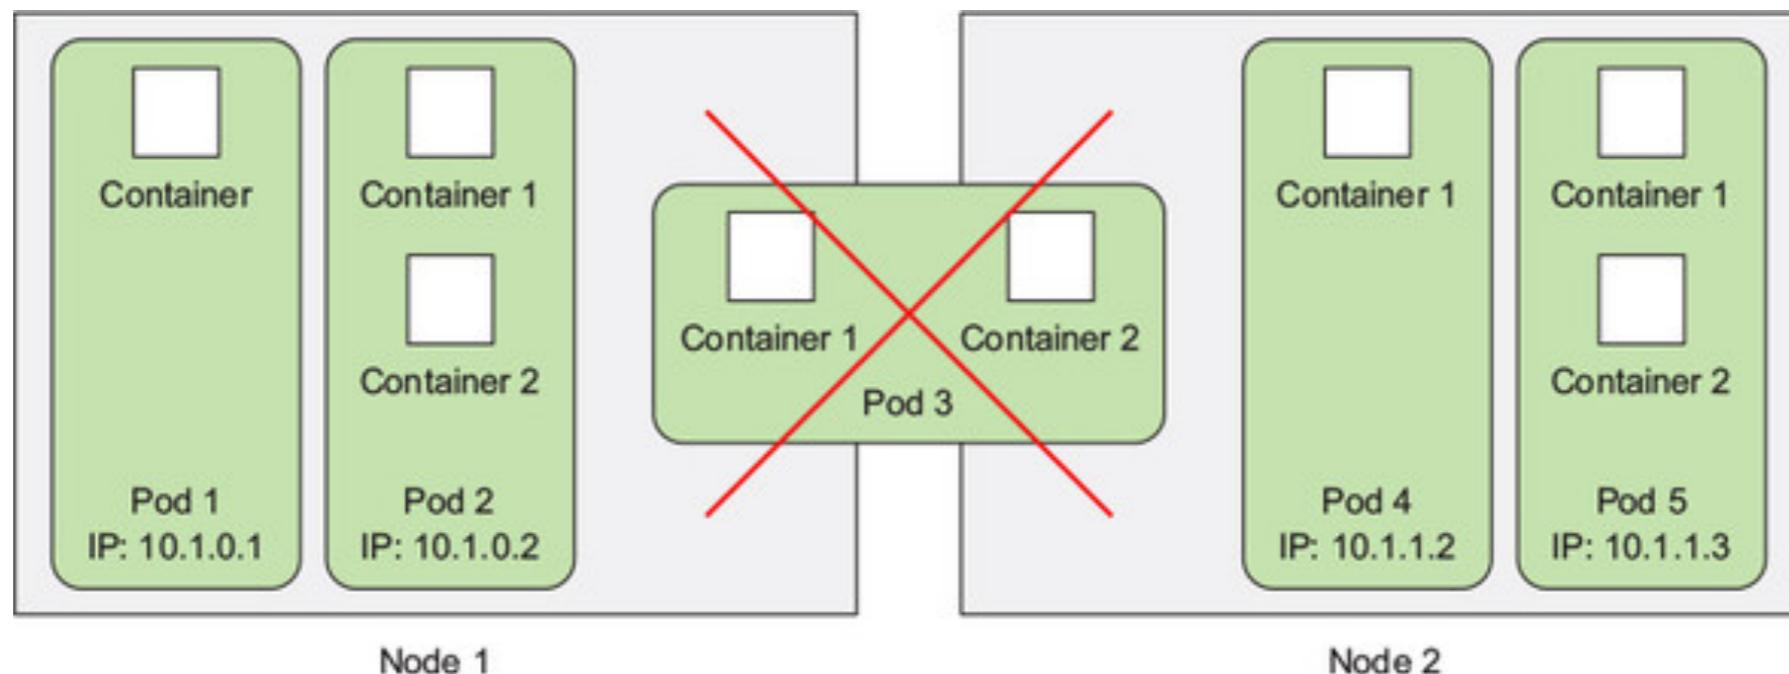
Pods

Small group of co-located containers
Optionally shared volume between containers
Basic deployment unit in Kubernetes



Pods

1 pods = 1 container
1 pods = N containers



All containers in same Pods

Share process ID

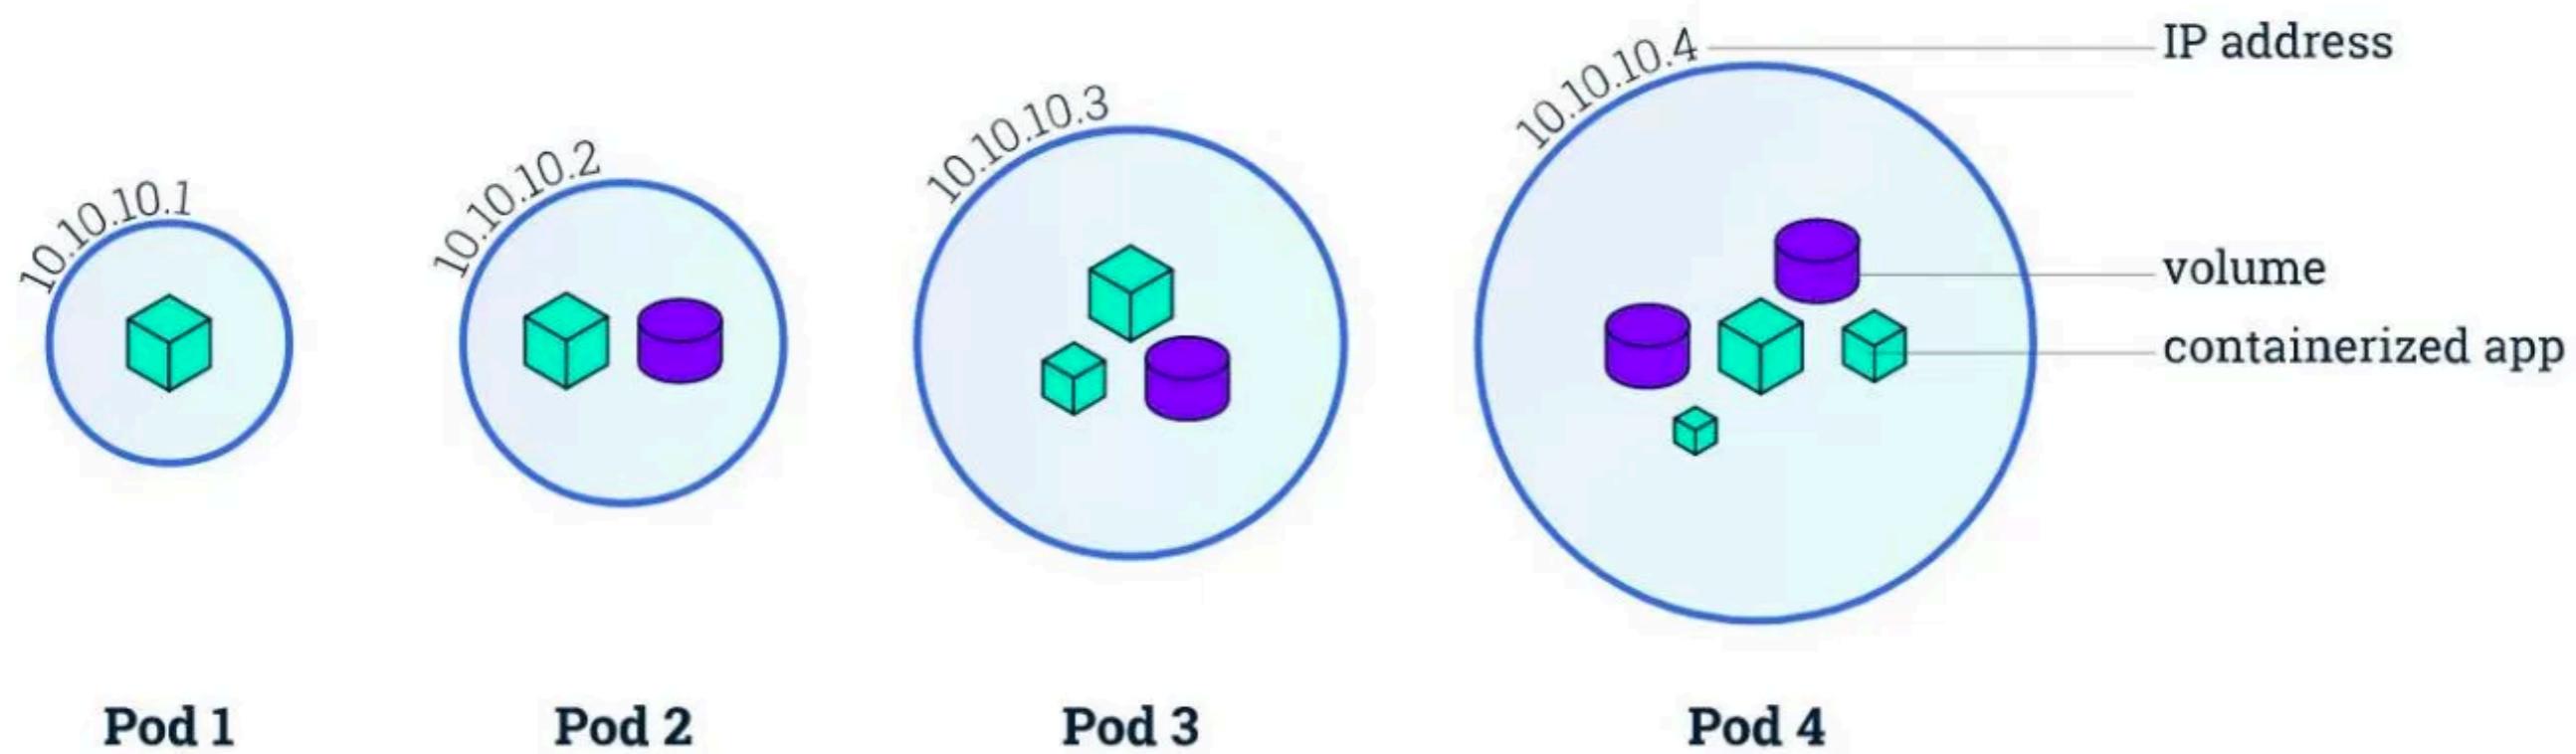
Share network interface

Share Hostname/IP/Port

Share Unix Time Sharing (UTS)

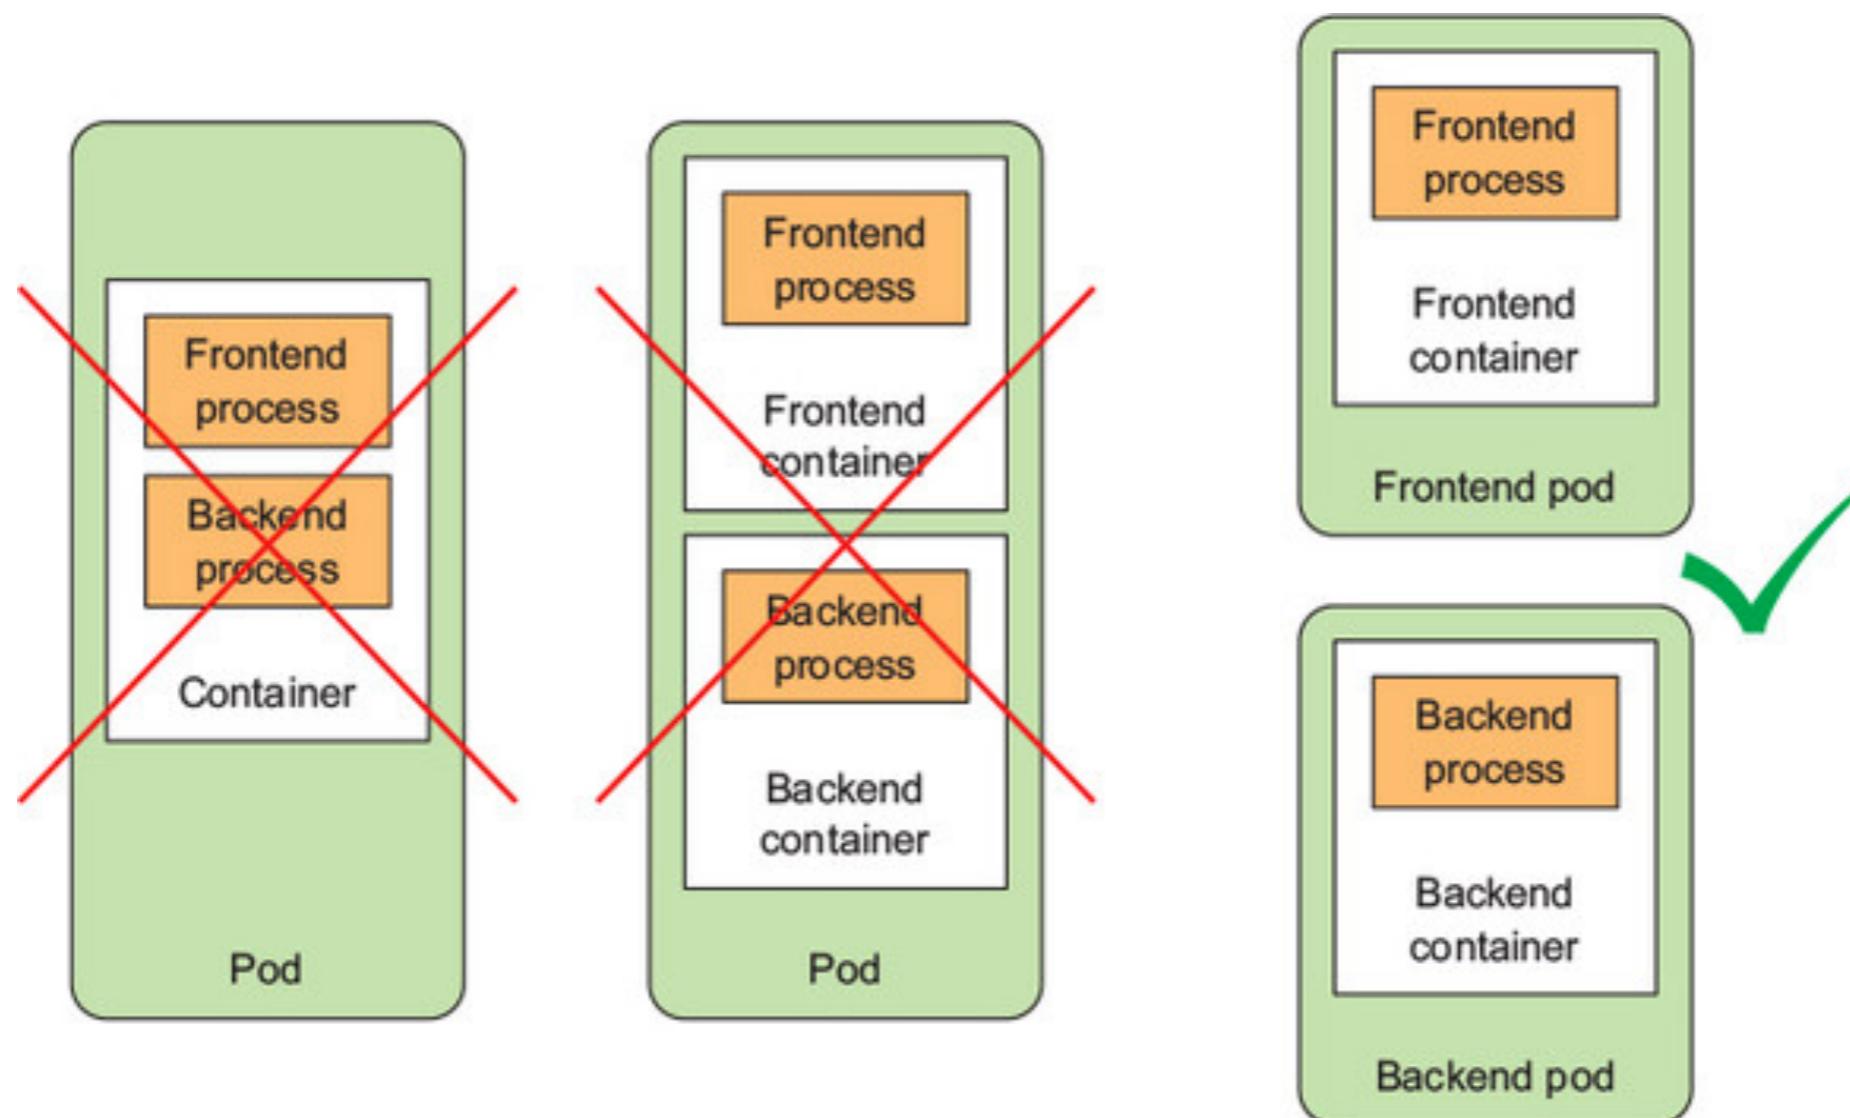


Pods



Organize container across Pods

Split multi-tiers app into multiple pods

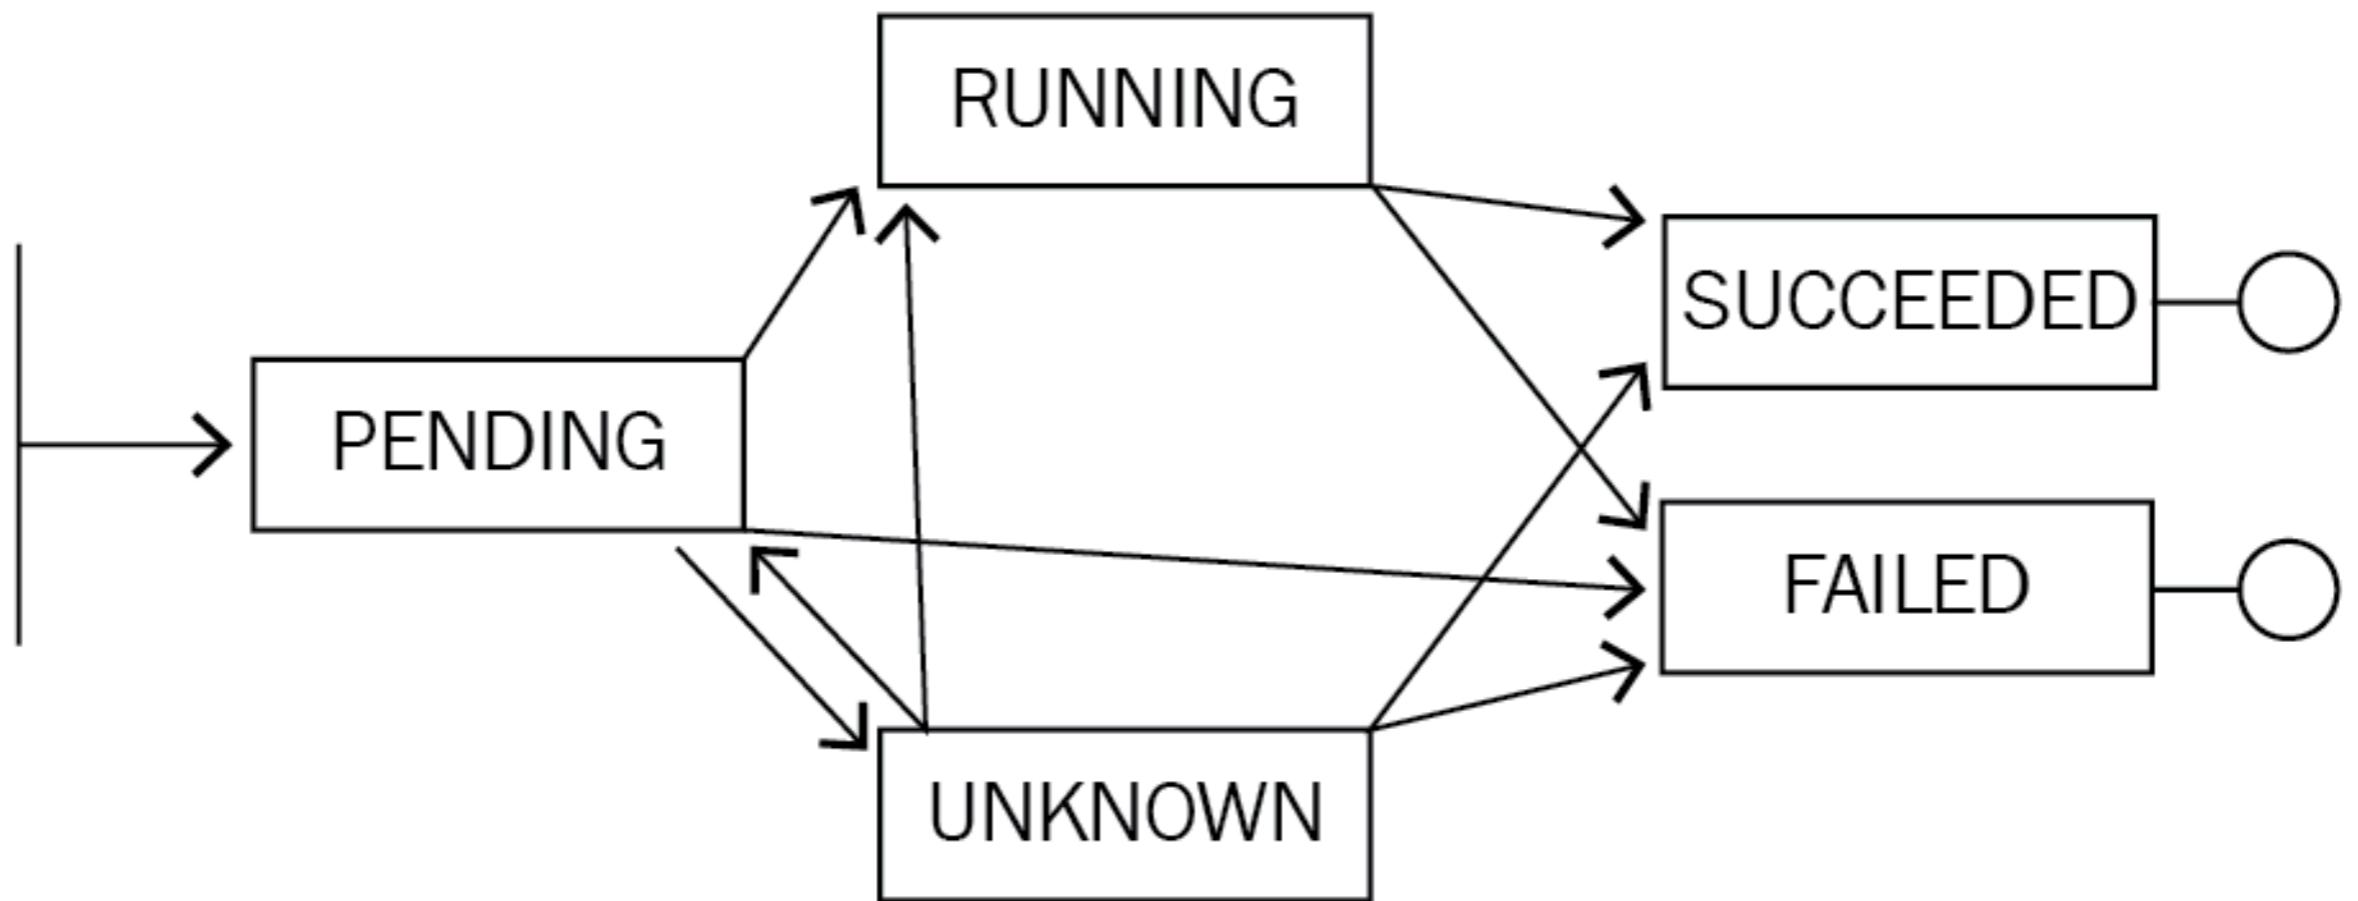


Pods lifecycle

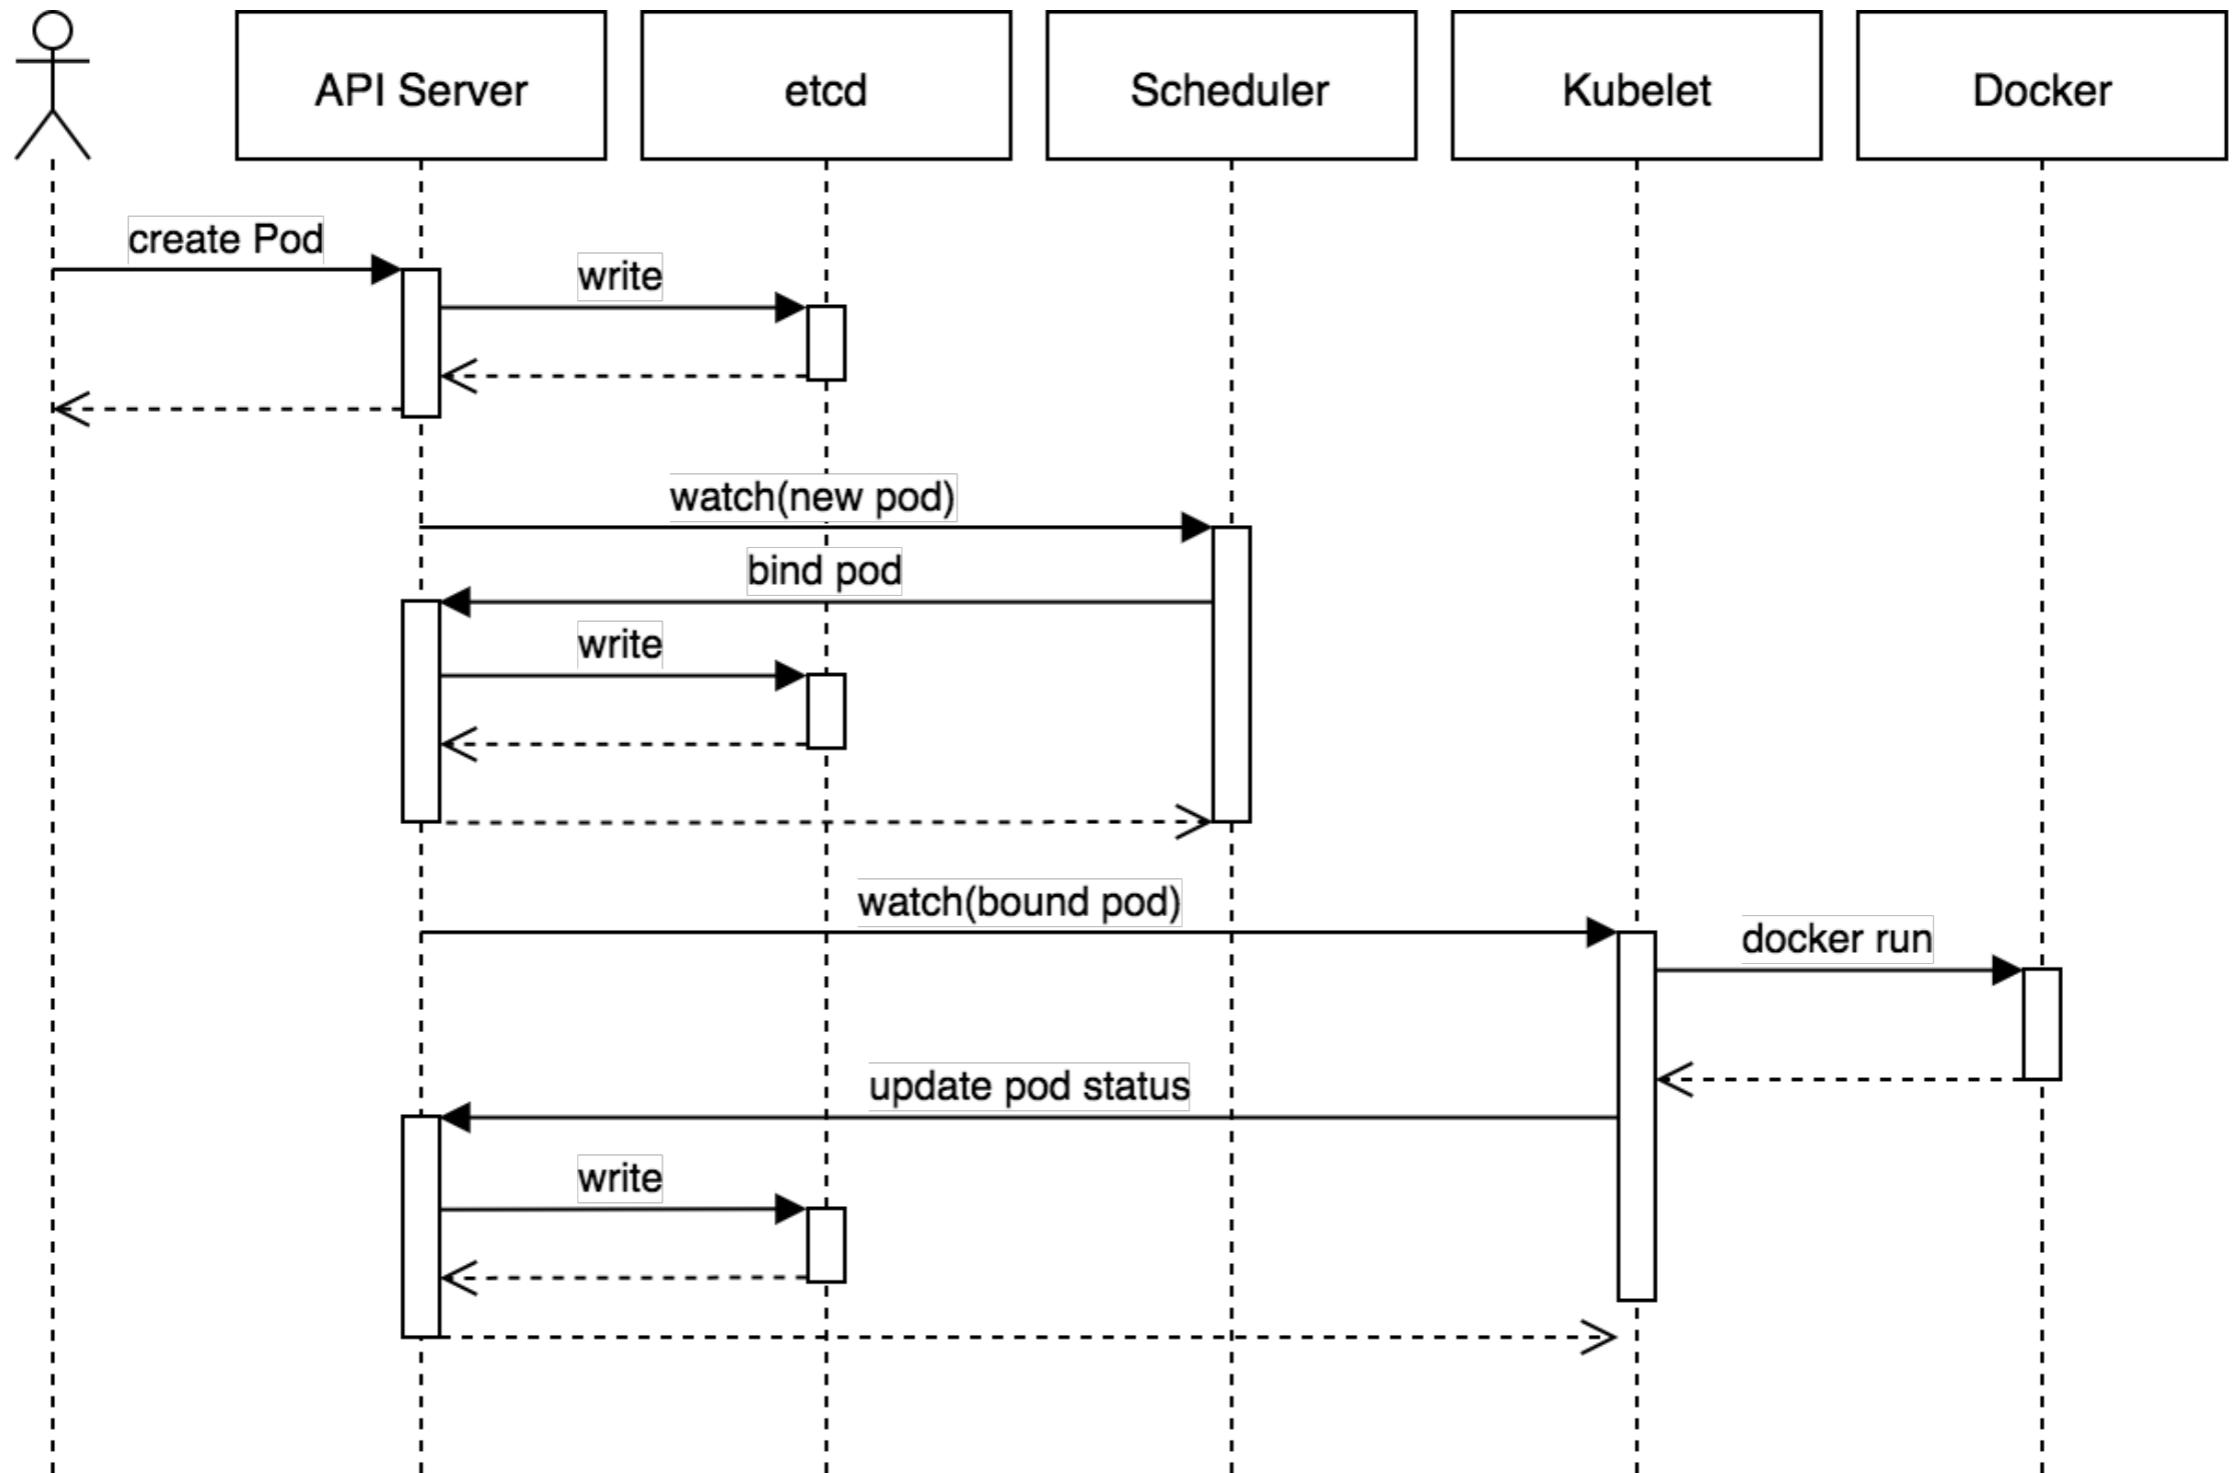
Phase	Description
Pending	Accepted by kubernetes but container not created yet
Running	Pods bound to the node, all containers created and at least one container is running/start/restarting
Successed	Containers exited with status 0
Failed	All containers exit and at least one exited with non-zero status
Unknown	State of Pods can't be determined due to communication issues with its node



Pods lifecycle



How to create Pods ?



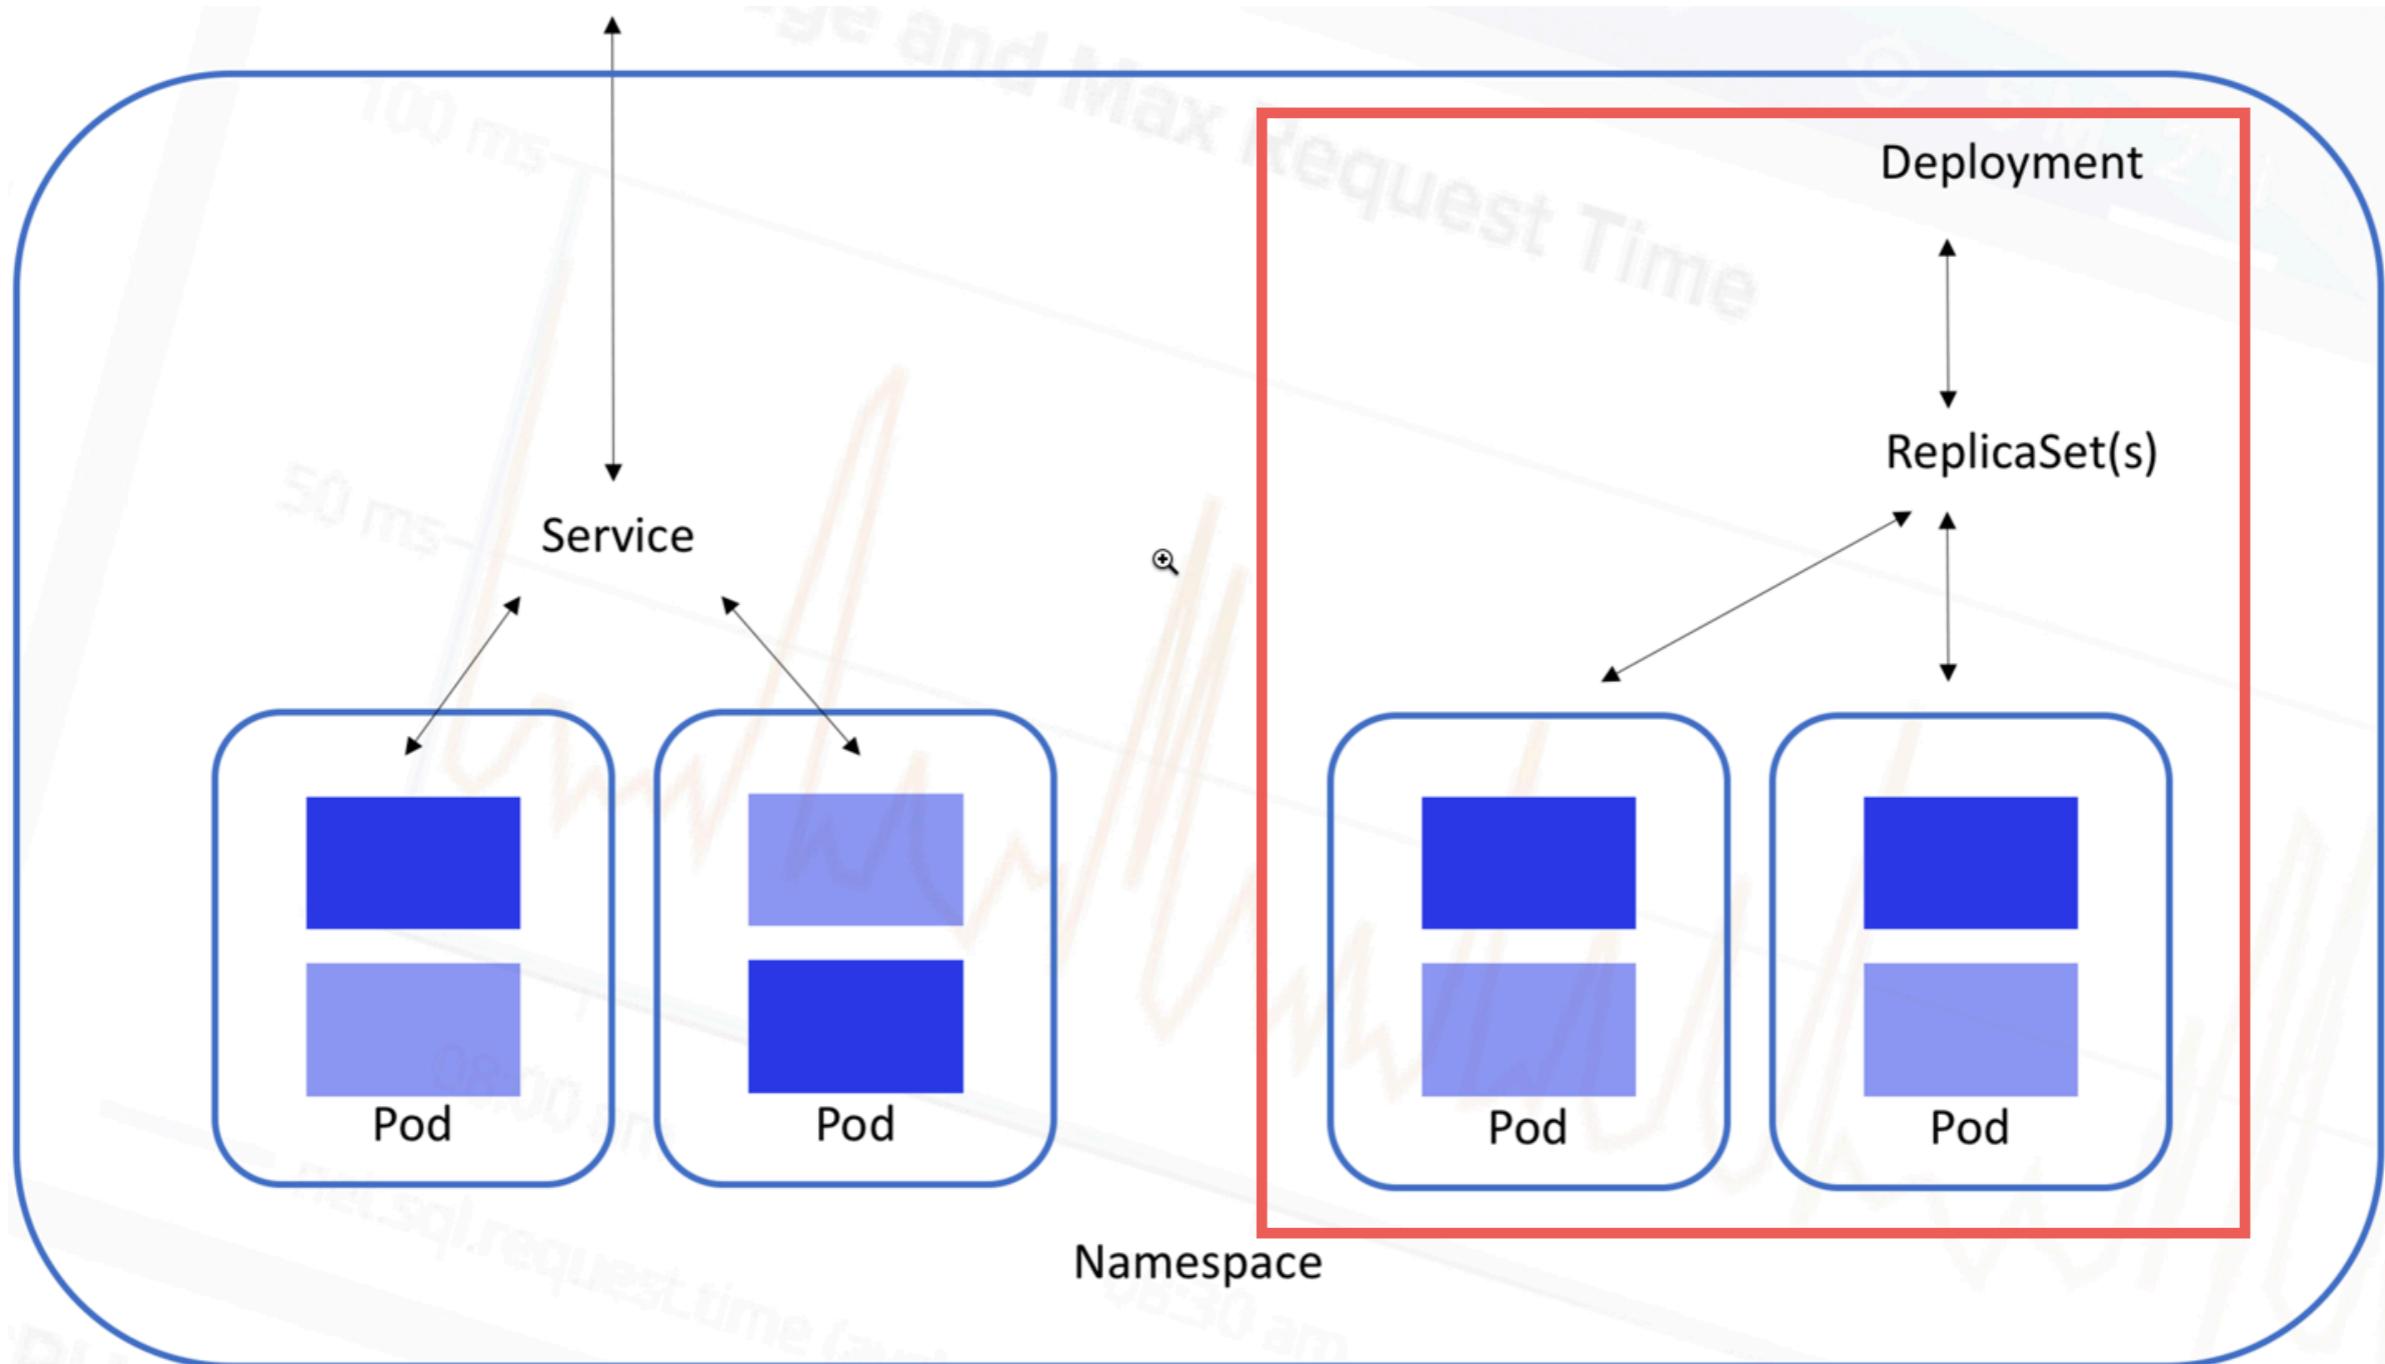
<https://blog.heptio.com/core-kubernetes-jazz-improv-over-orchestration-a7903ea92ca>



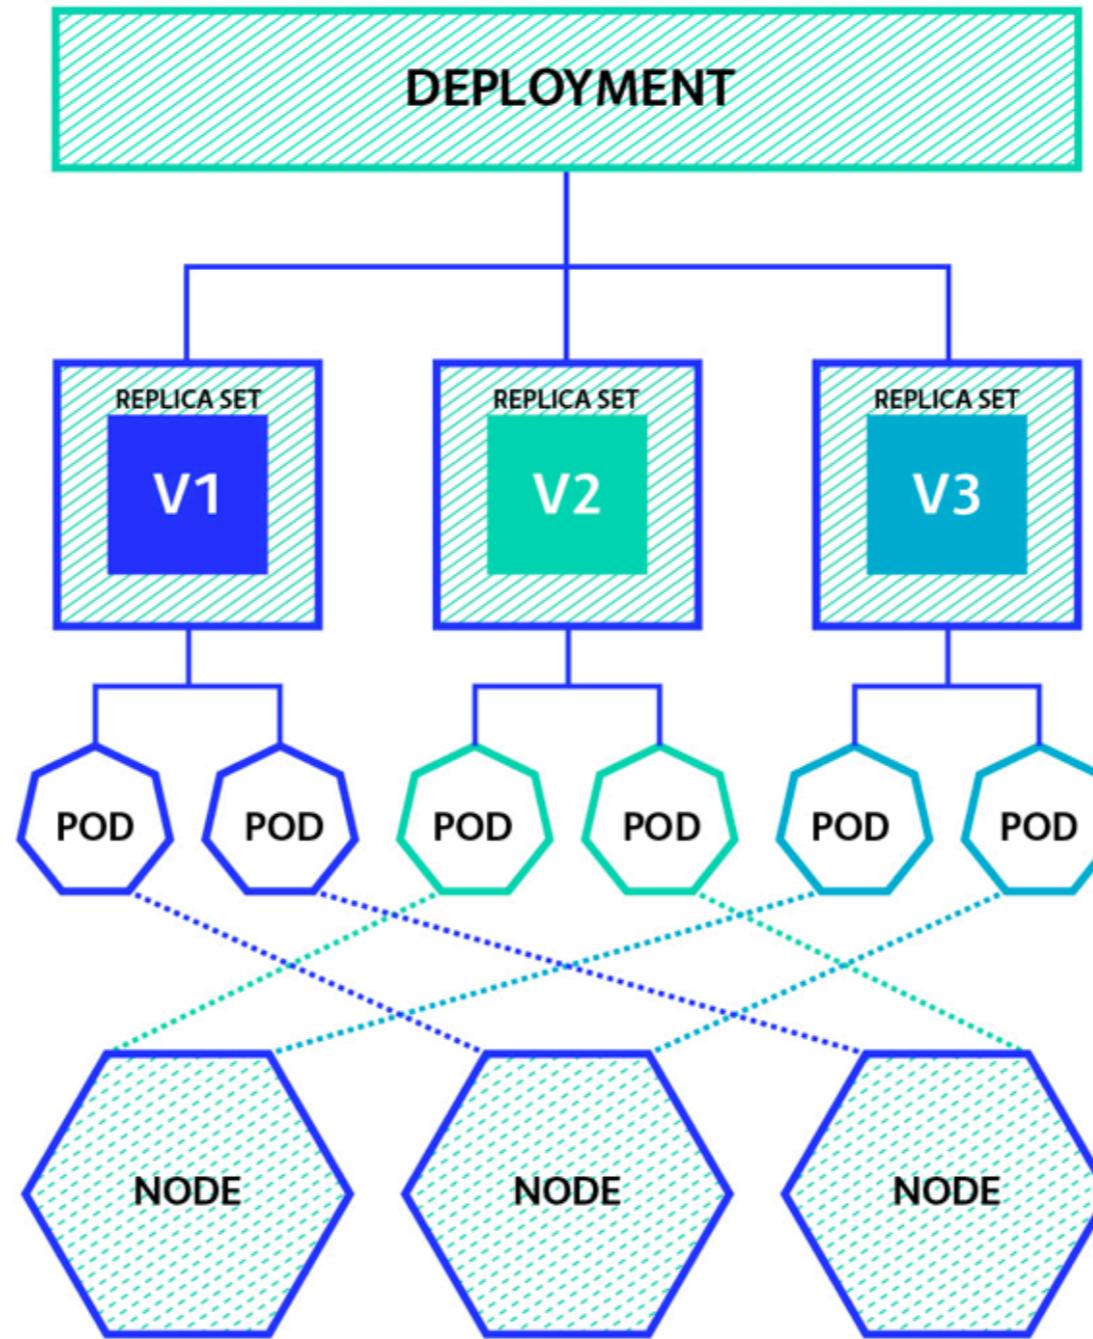
Deployment



Deployment and ReplicaSet



Deployment and ReplicaSet



<https://thenewstack.io/kubernetes-deployments-work/>



Deployment and ReplicaSet

Next-generation of RC

Provide function to maintain versioning of Pods

Update new version (Rollout)

Revert to old version (Rollback)

Scale a deployment

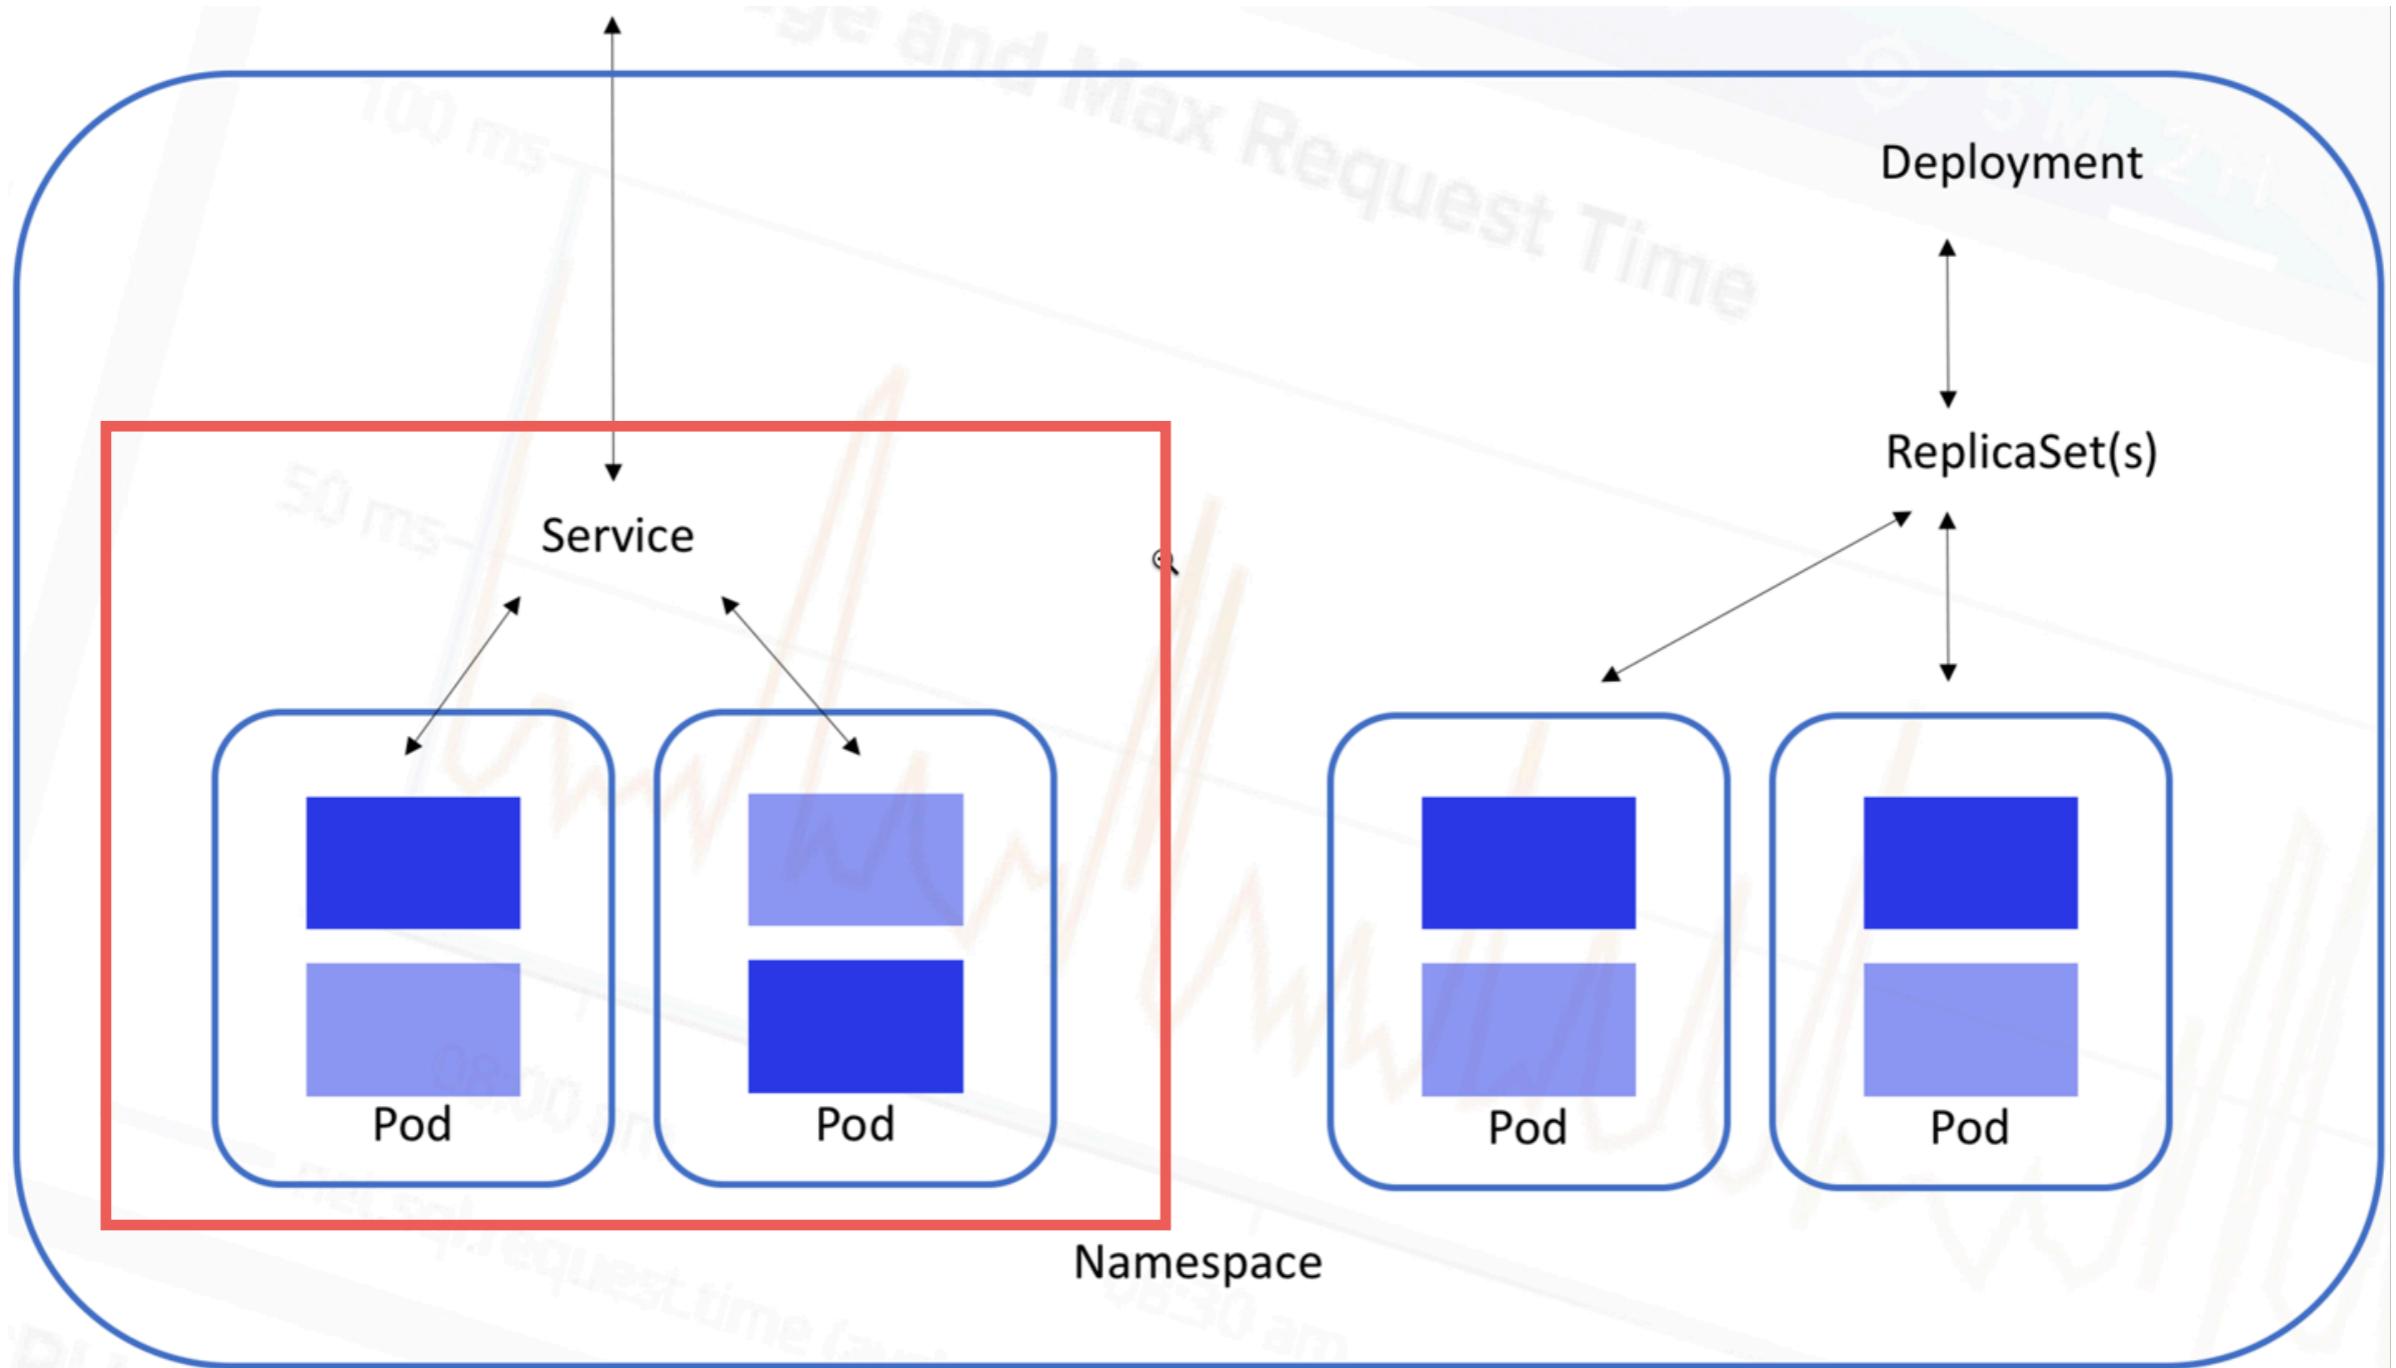
Pause/Resume process



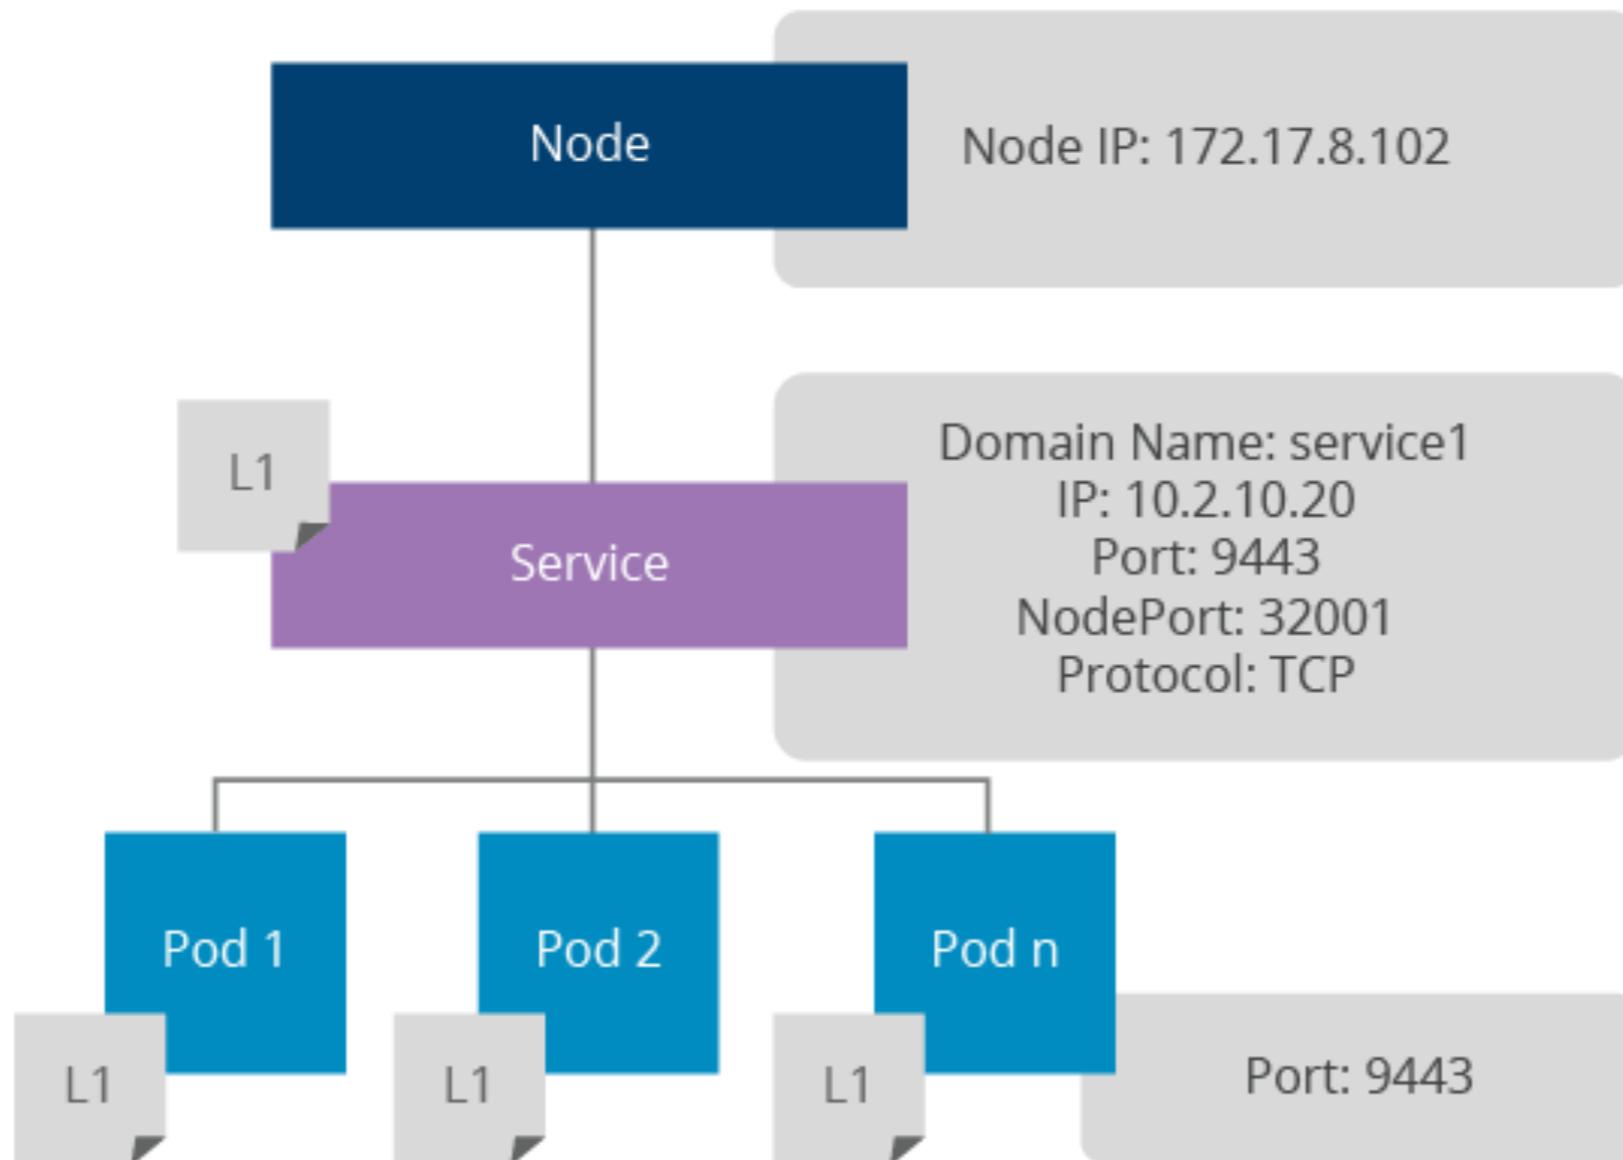
Services



Services



Services

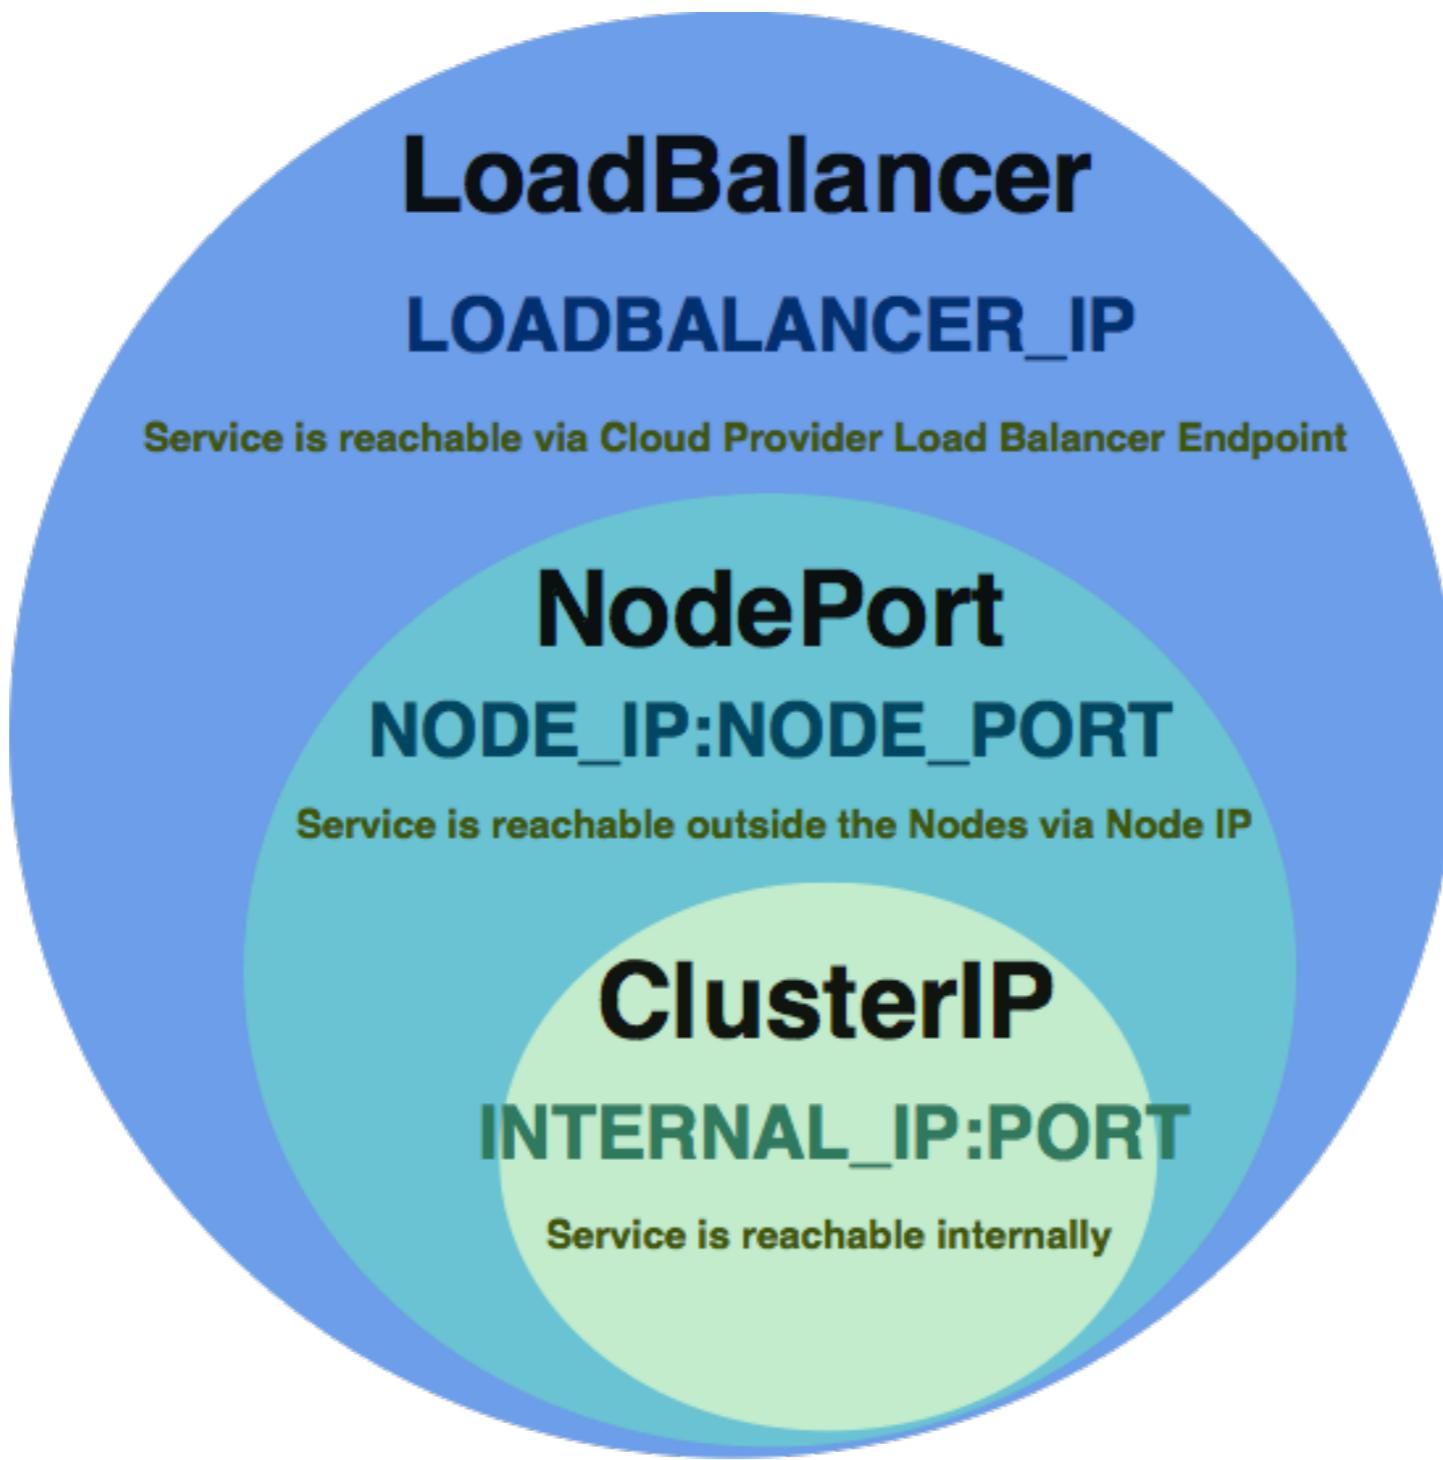


Services

- Independent from Pods
- Abstraction layer of Pods
- Provide load balance
- Expose access Pods/Load balance
- Find Pods by label selector



3 types of services

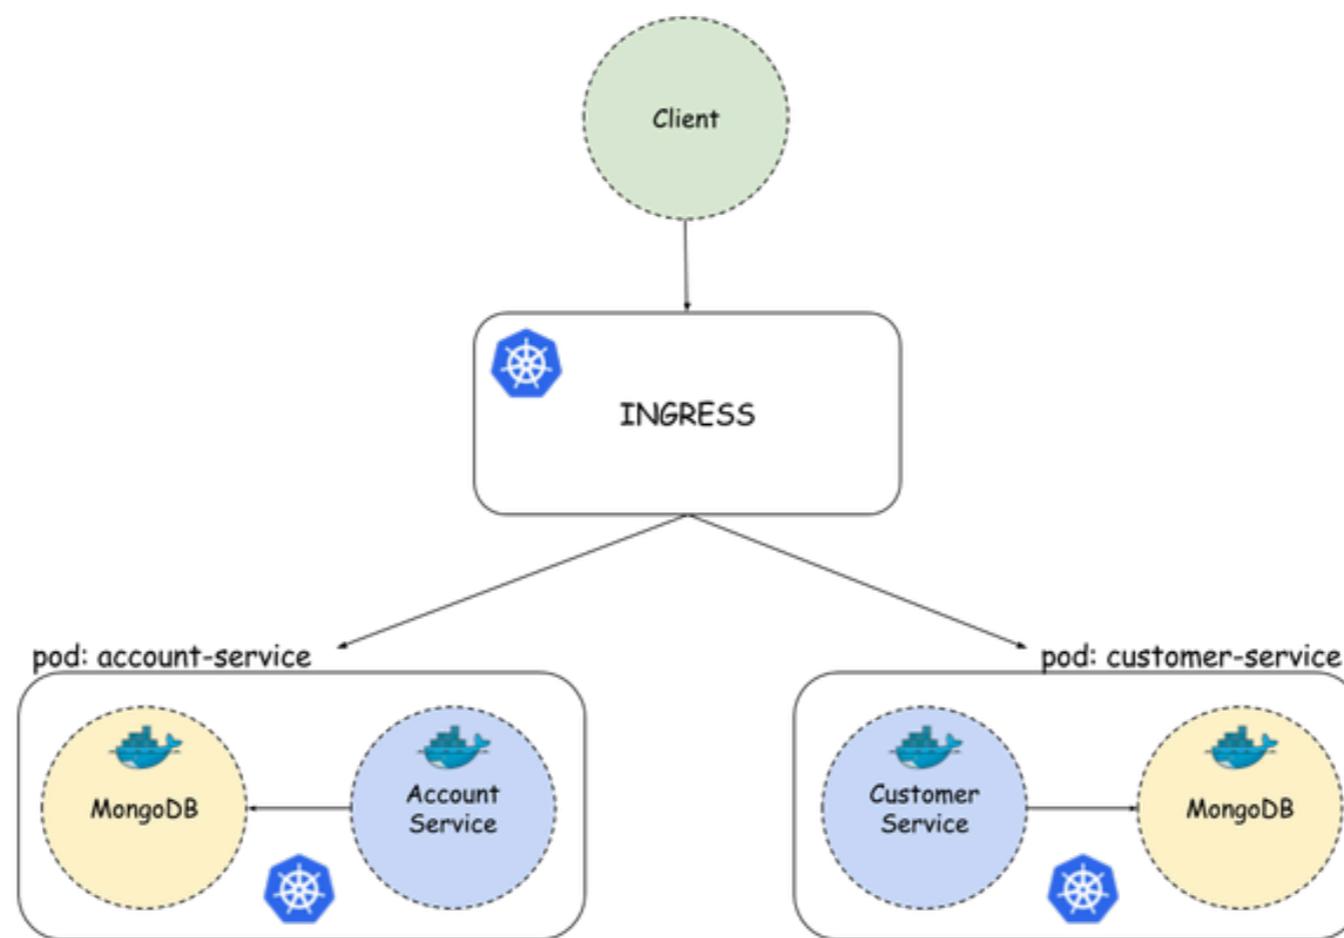


Ingress controller



Ingress controller

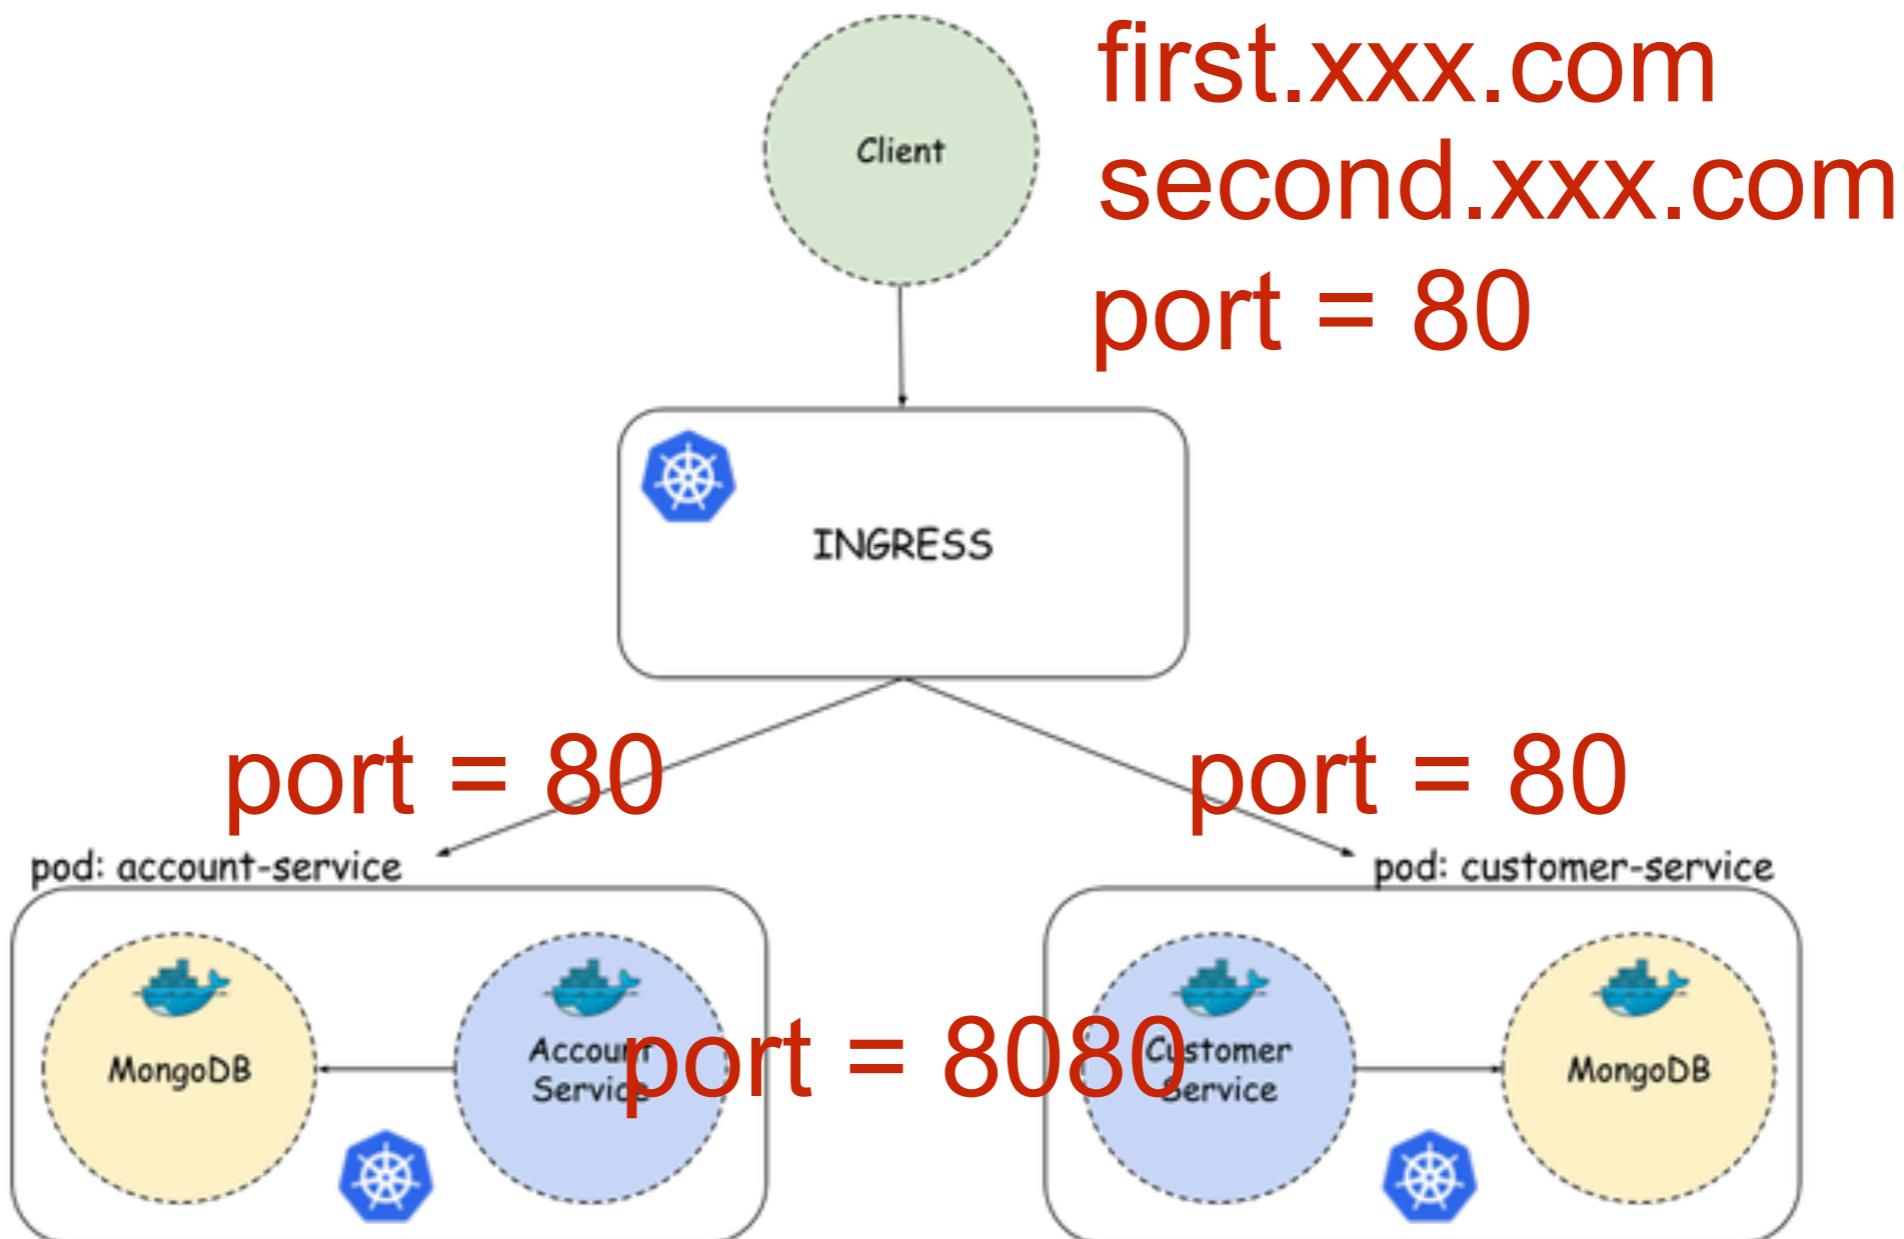
How to handle multiple services in same port ?
How to limit protocol to access ?



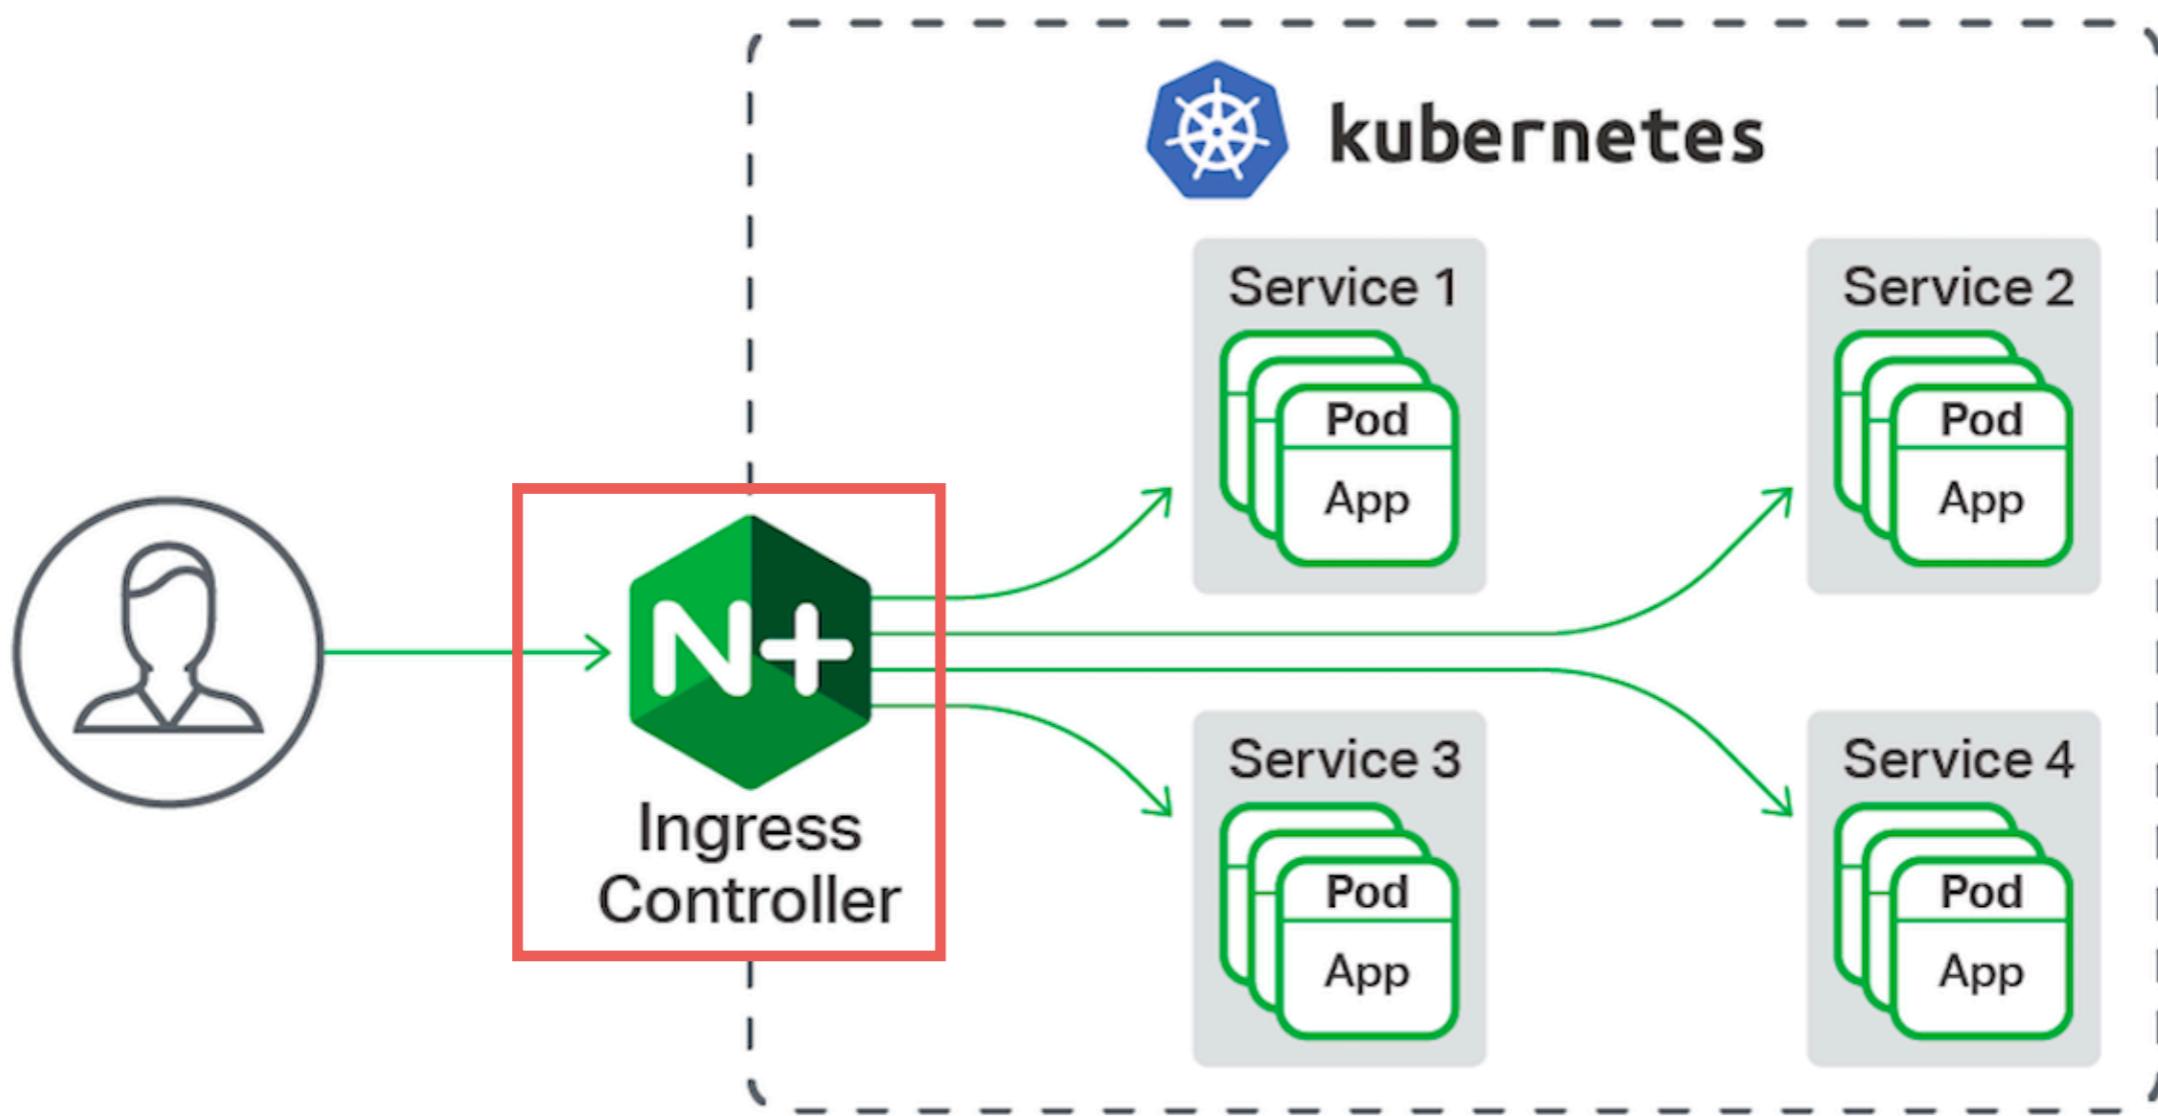
<https://kubernetes.io/docs/concepts/services-networking/ingress/>



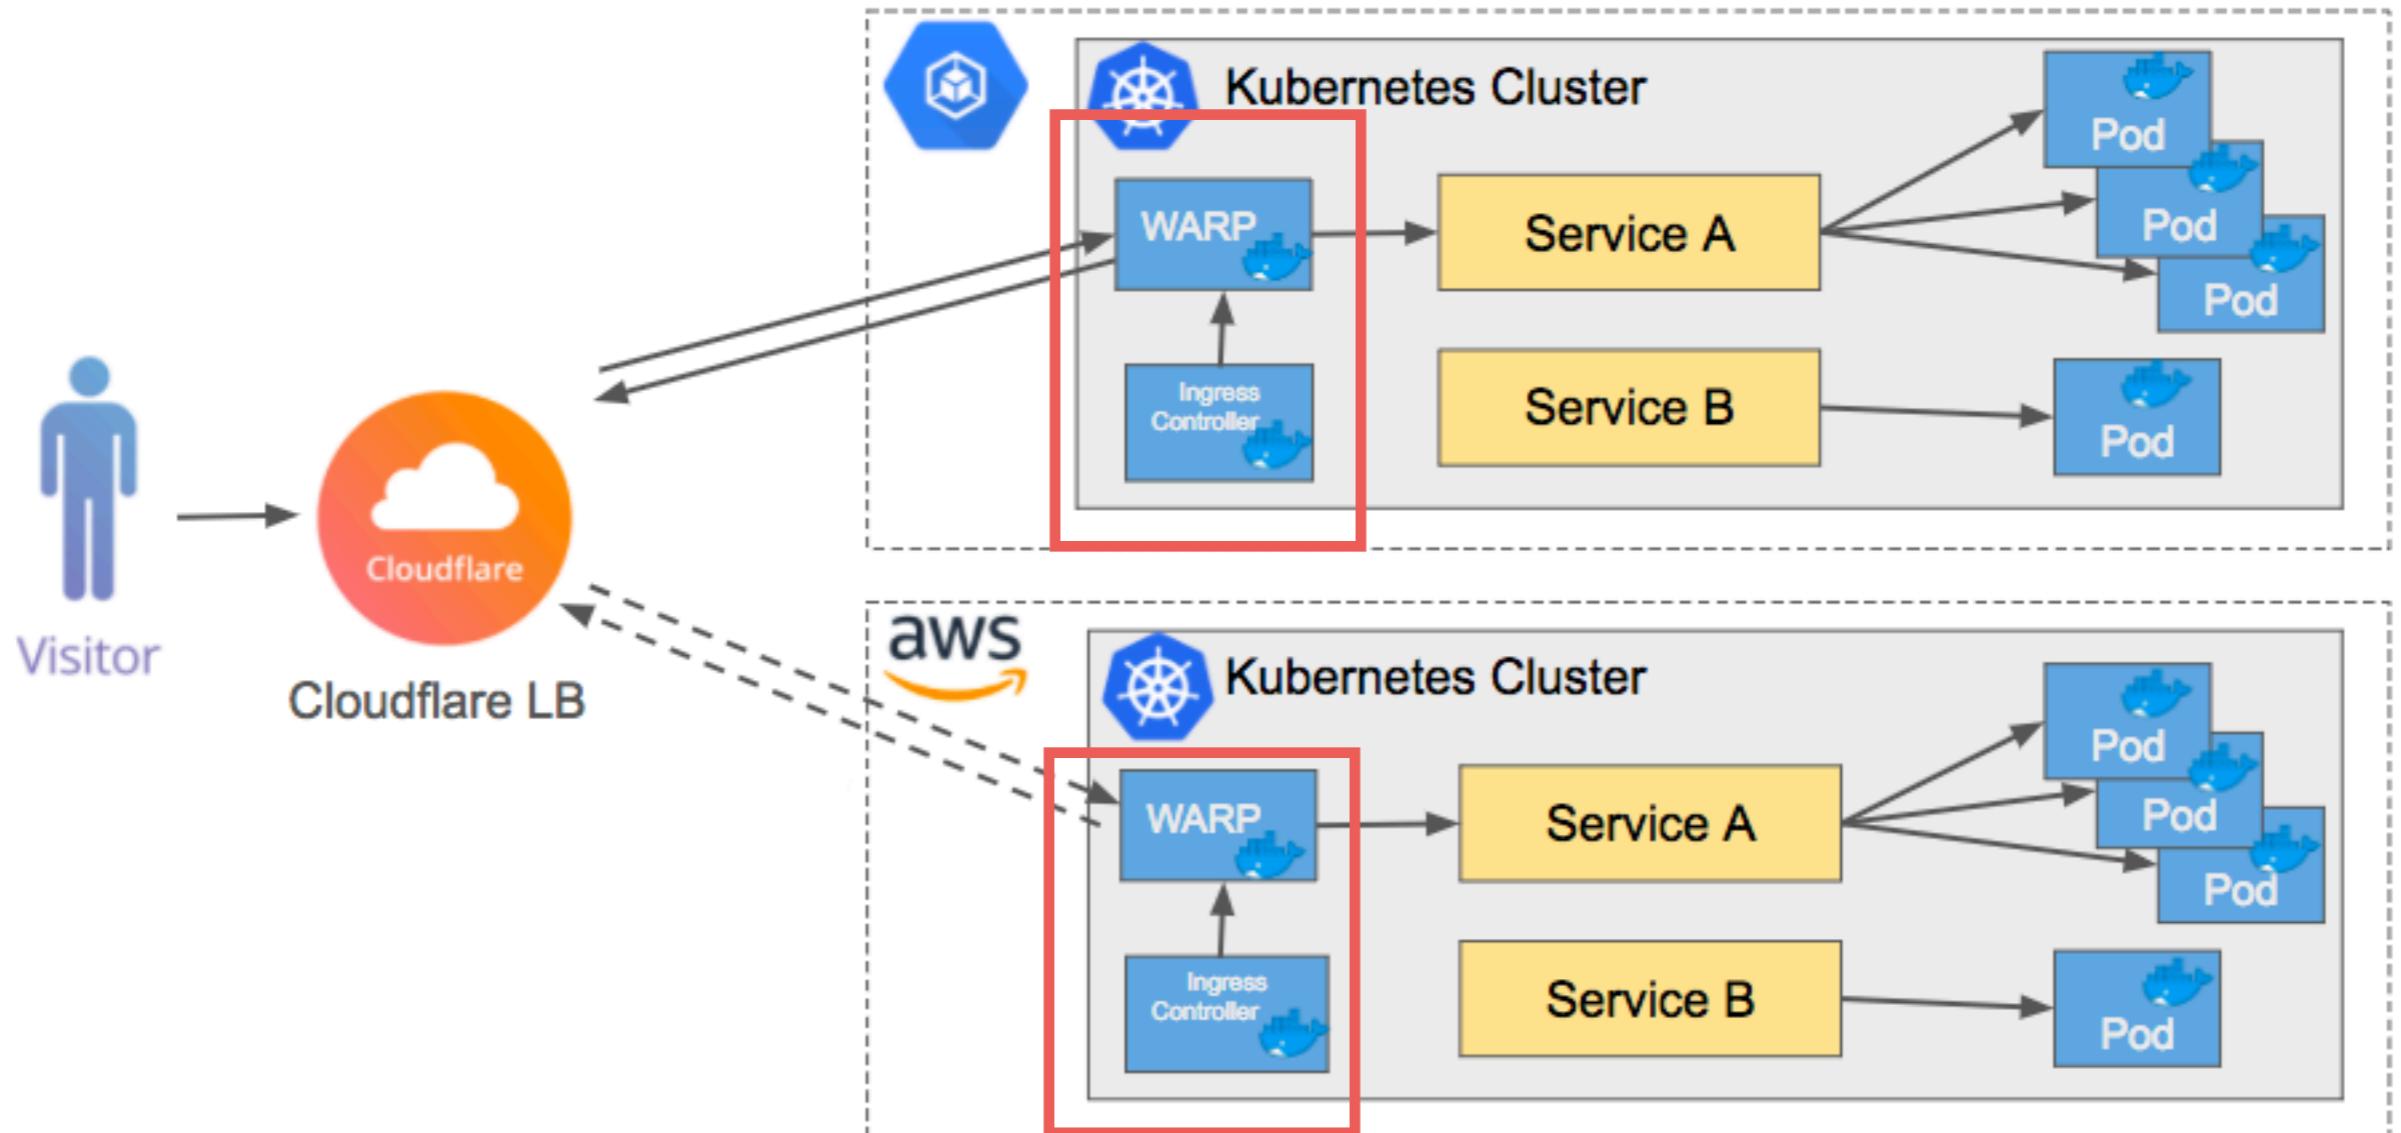
Ingress controller



Ingress Network



Ingress Network



ConfigMap and Secret



ConfigMap

Configuration parameters
Key-value pairs

Environment
variables

Container
command line
arguments

Volume



Secret

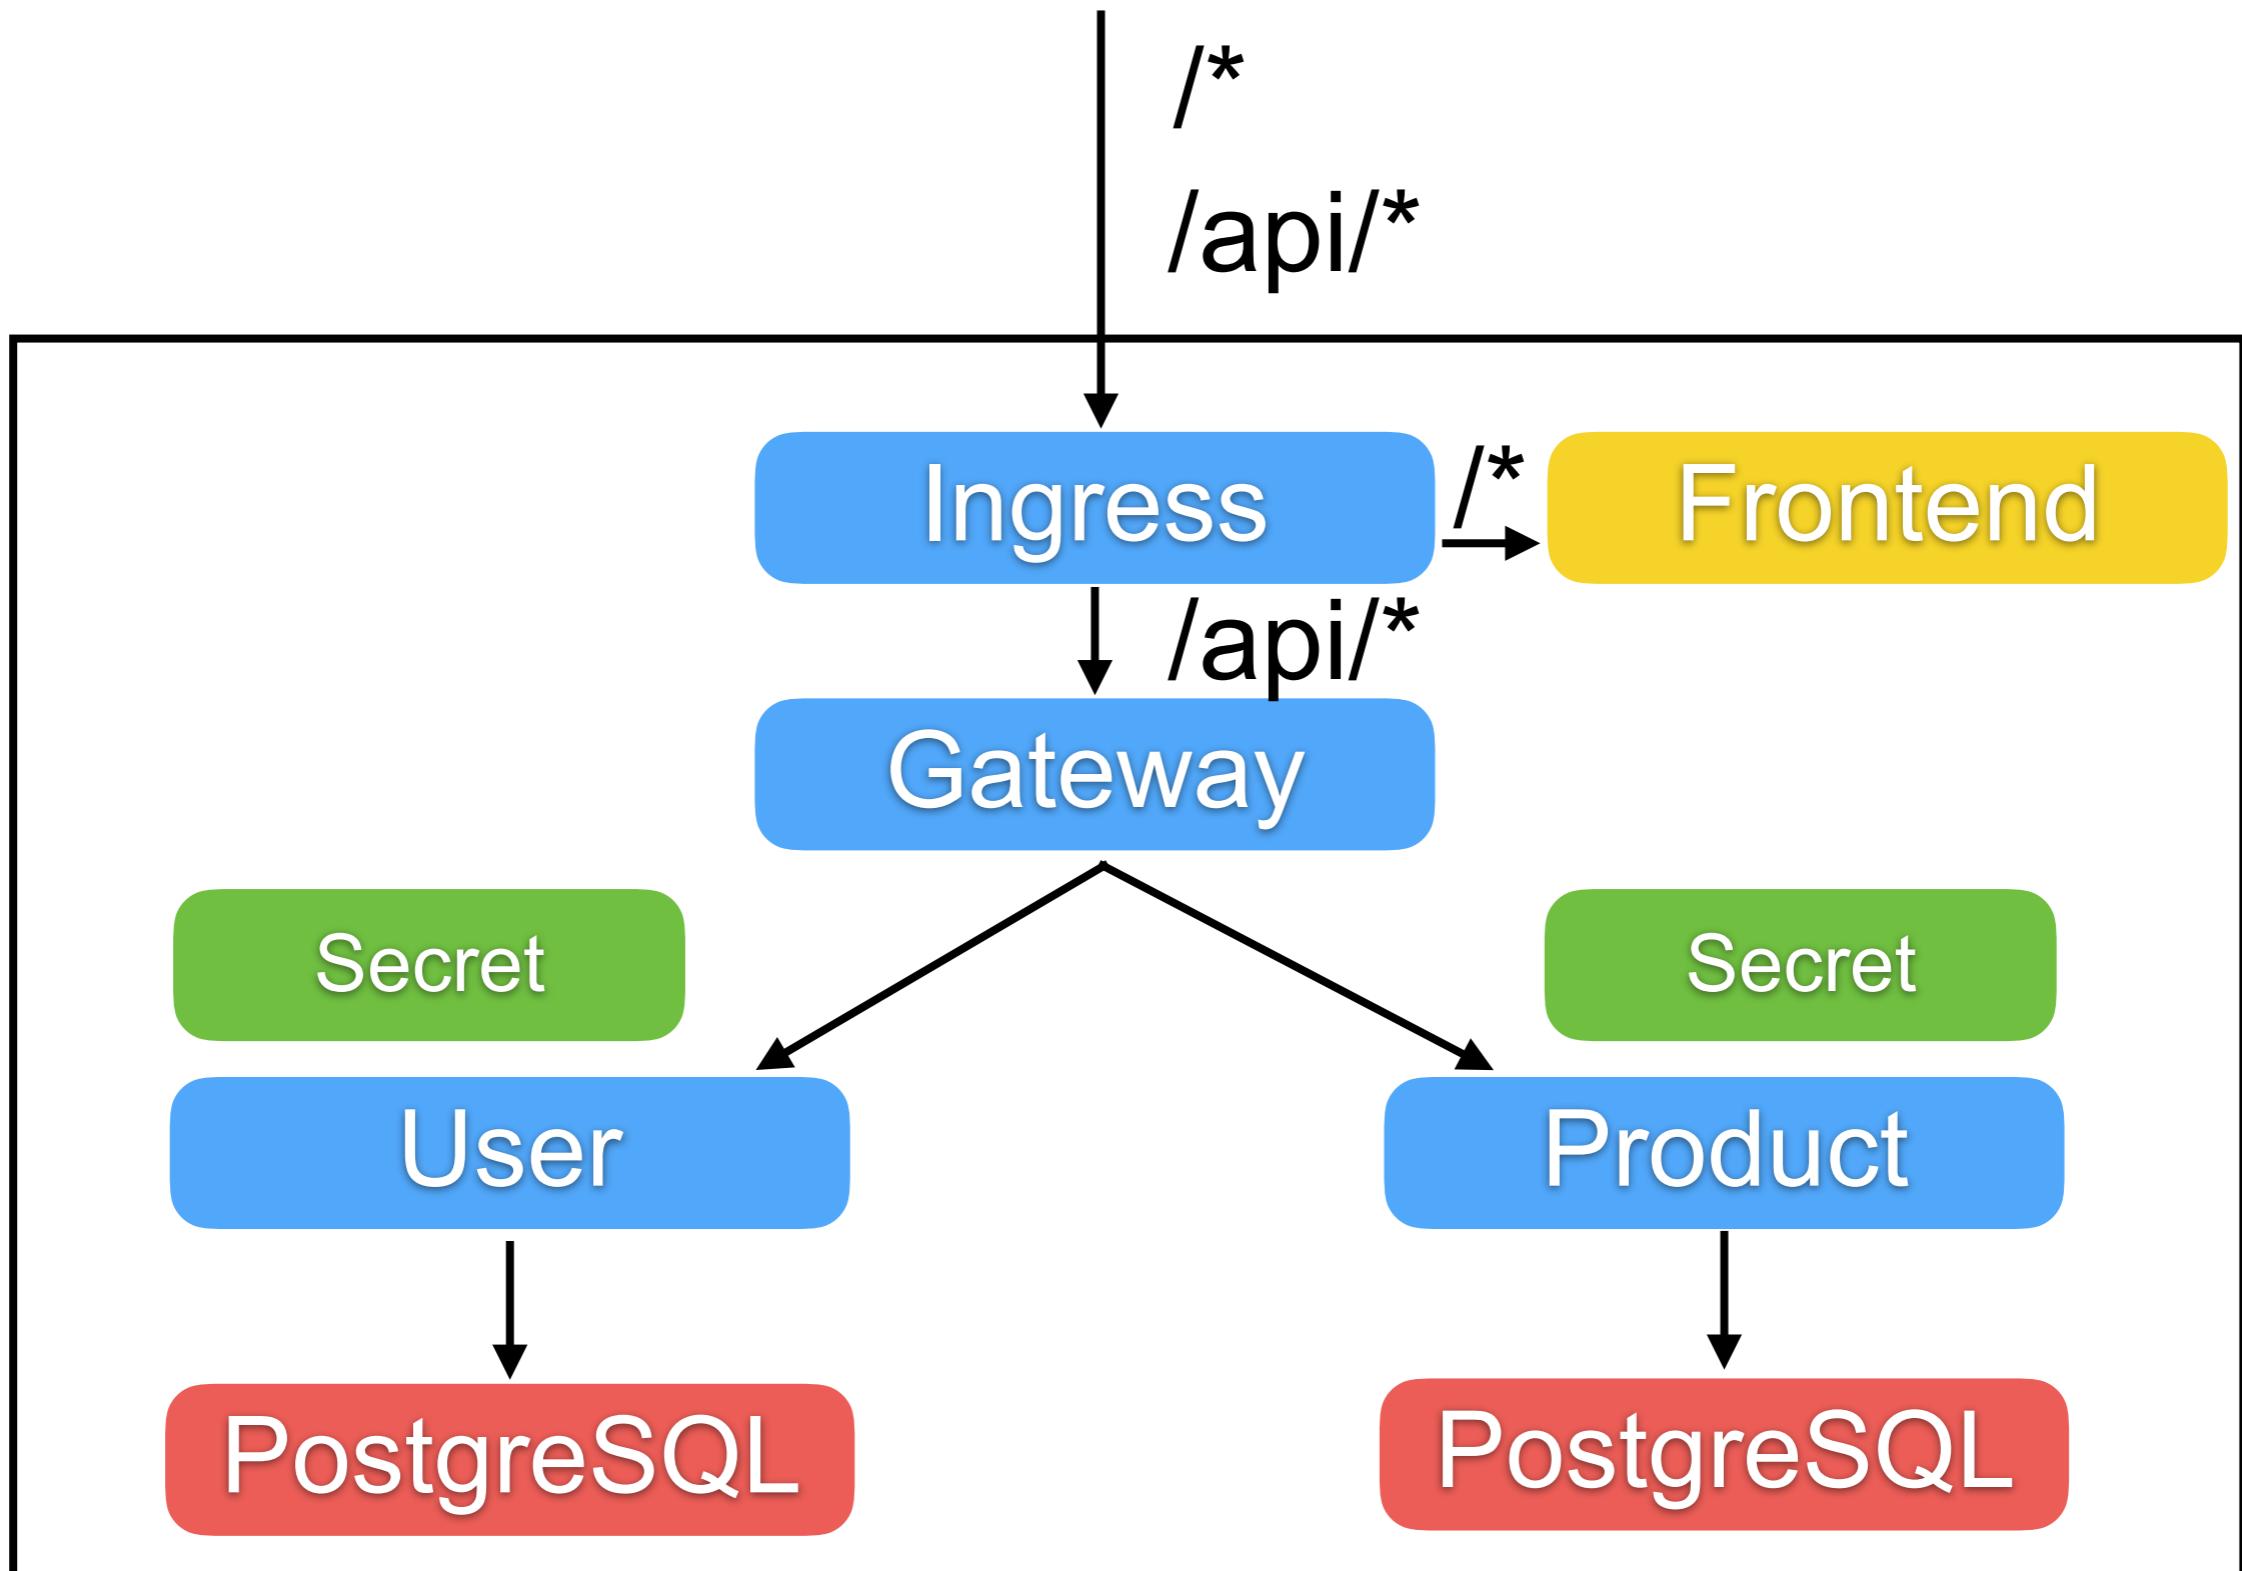
**Way to Kubernetes to distributed credentials
keys, passwords or secret data to Pod**

Native way of Kubernetes

Other way use external vault services



Structure of service



Workshop

<https://github.com/up1/demo-docker-k8s>



Application Deployment Strategies



Strategies to deploy

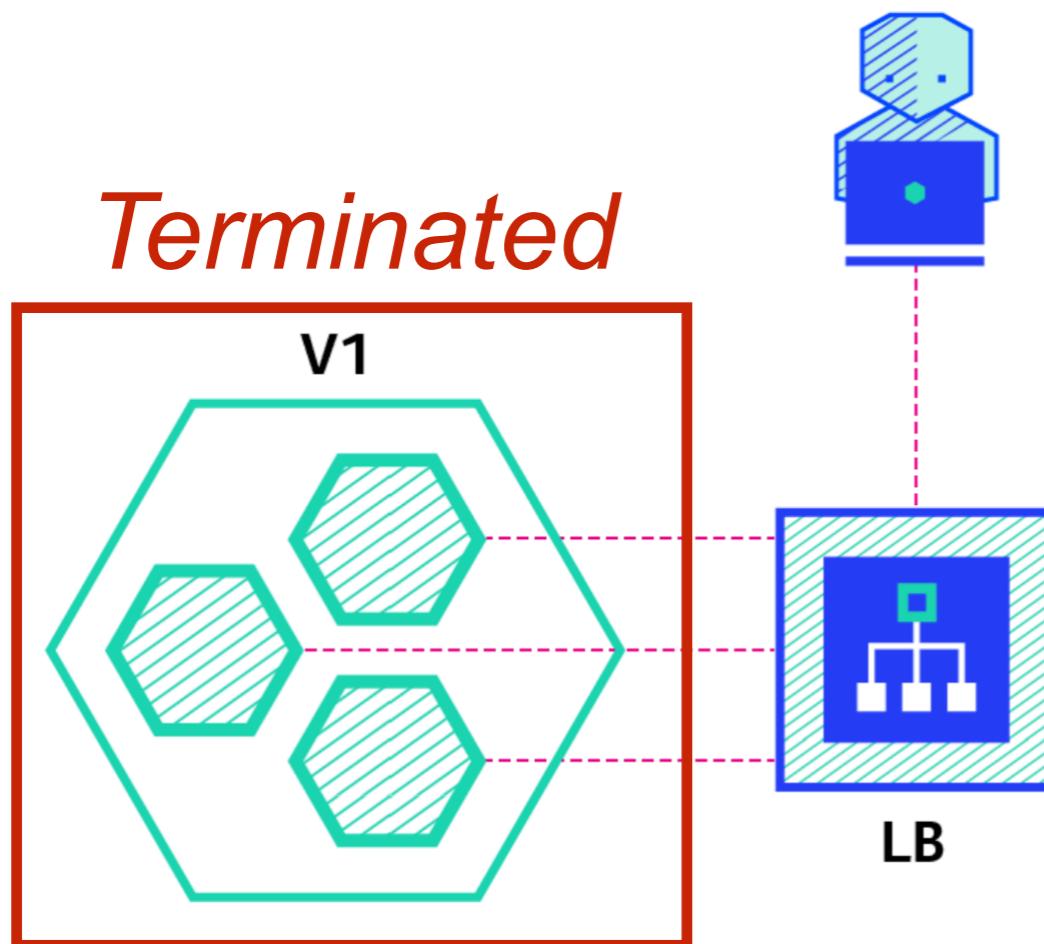
Recreate
Ramped
Blue/Green
Canary
A/B testing
Shadow

<https://thenewstack.io/deployment-strategies/>



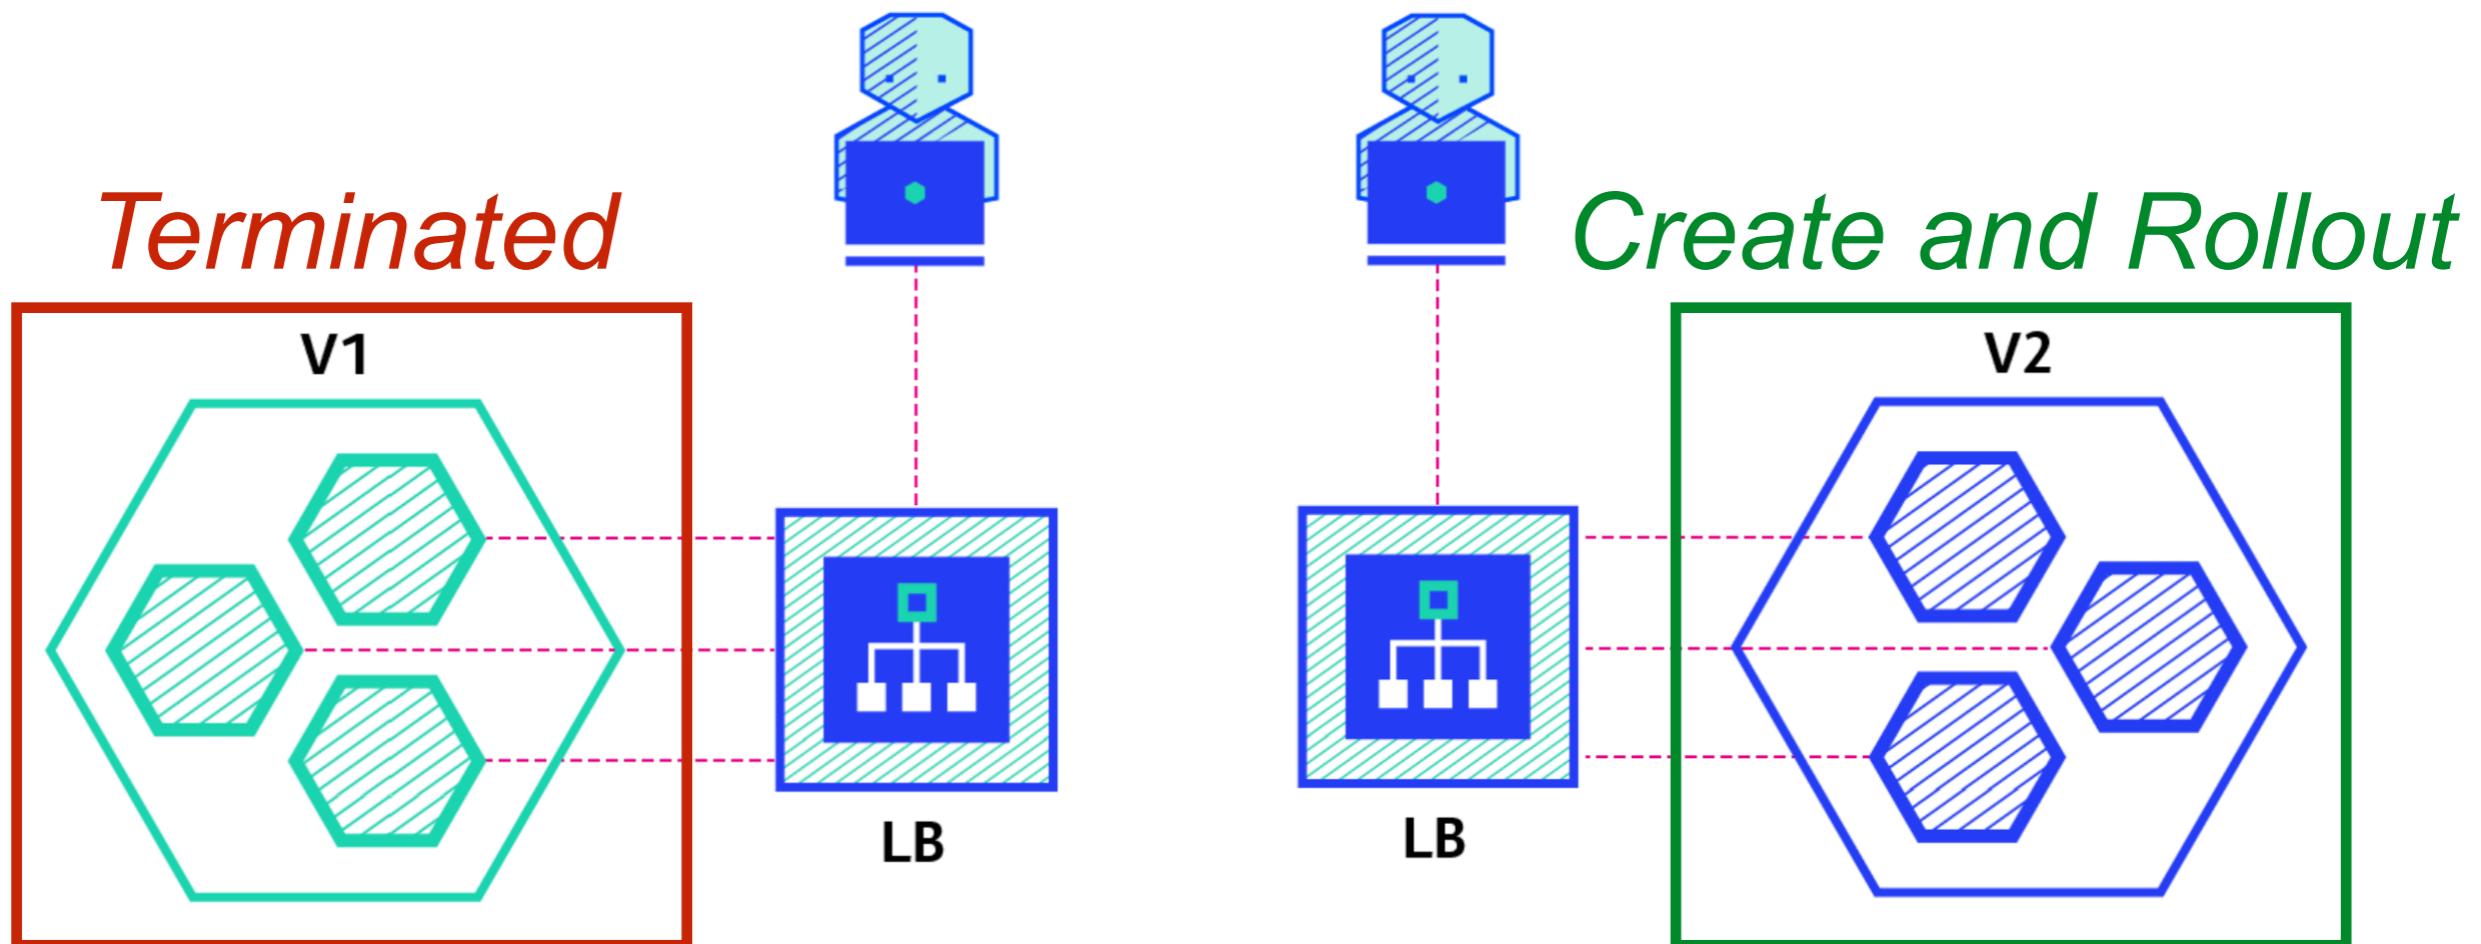
1. Recreate

Version A is terminated then version B is rollout



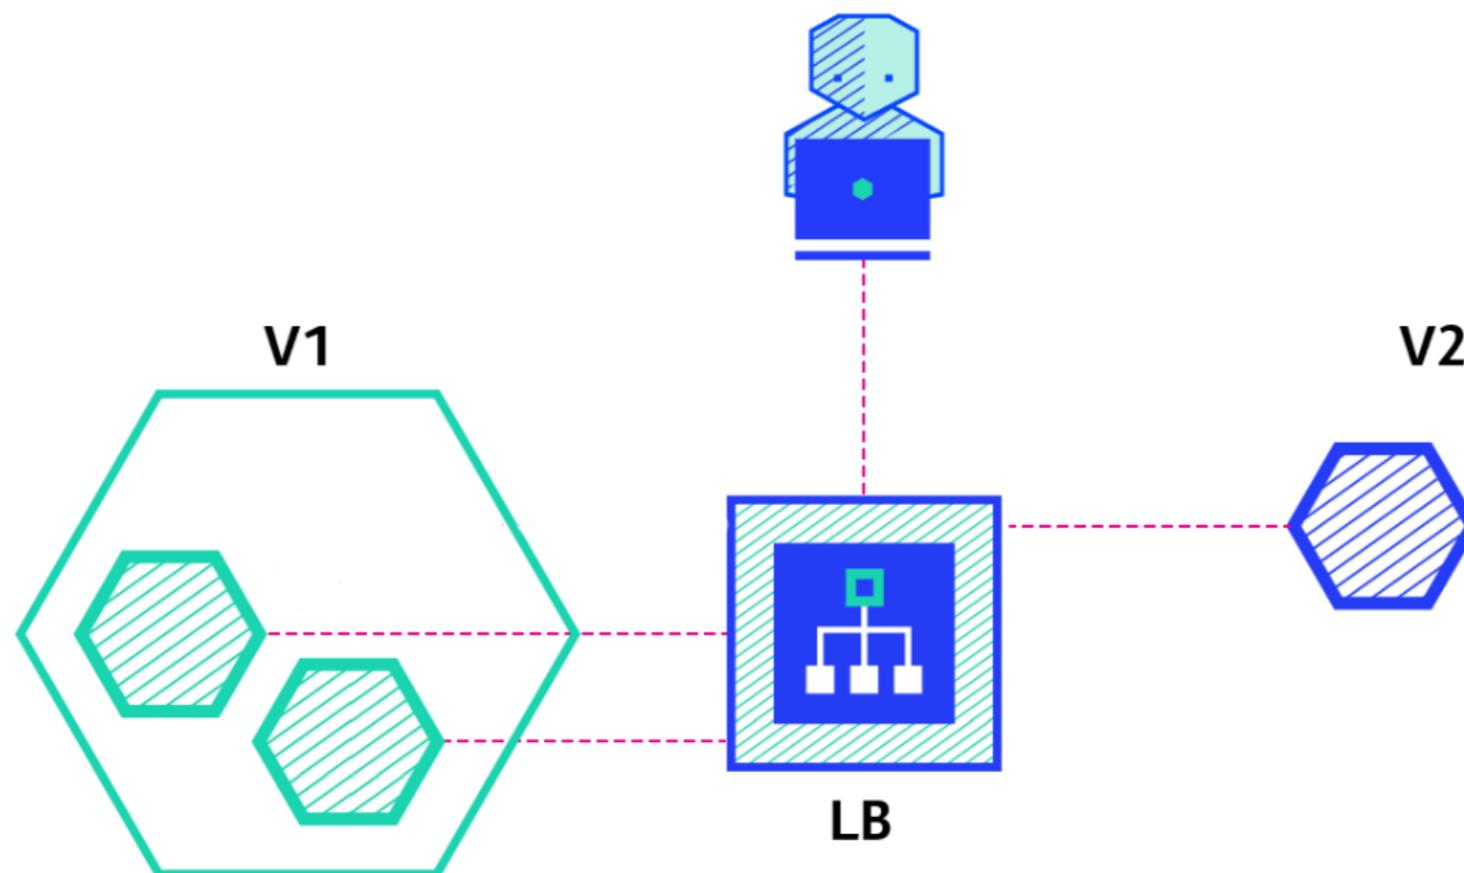
1. Recreate

Version A is terminated then version B is rollout



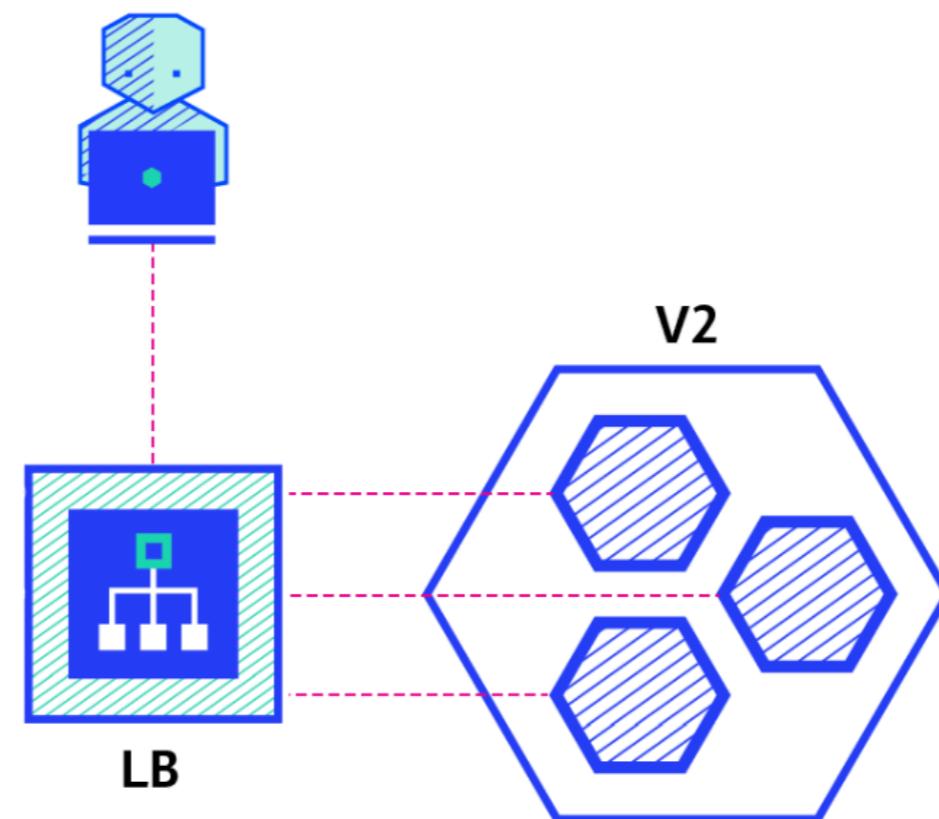
2. Ramped

Slow roll out by replace instance one-by-one



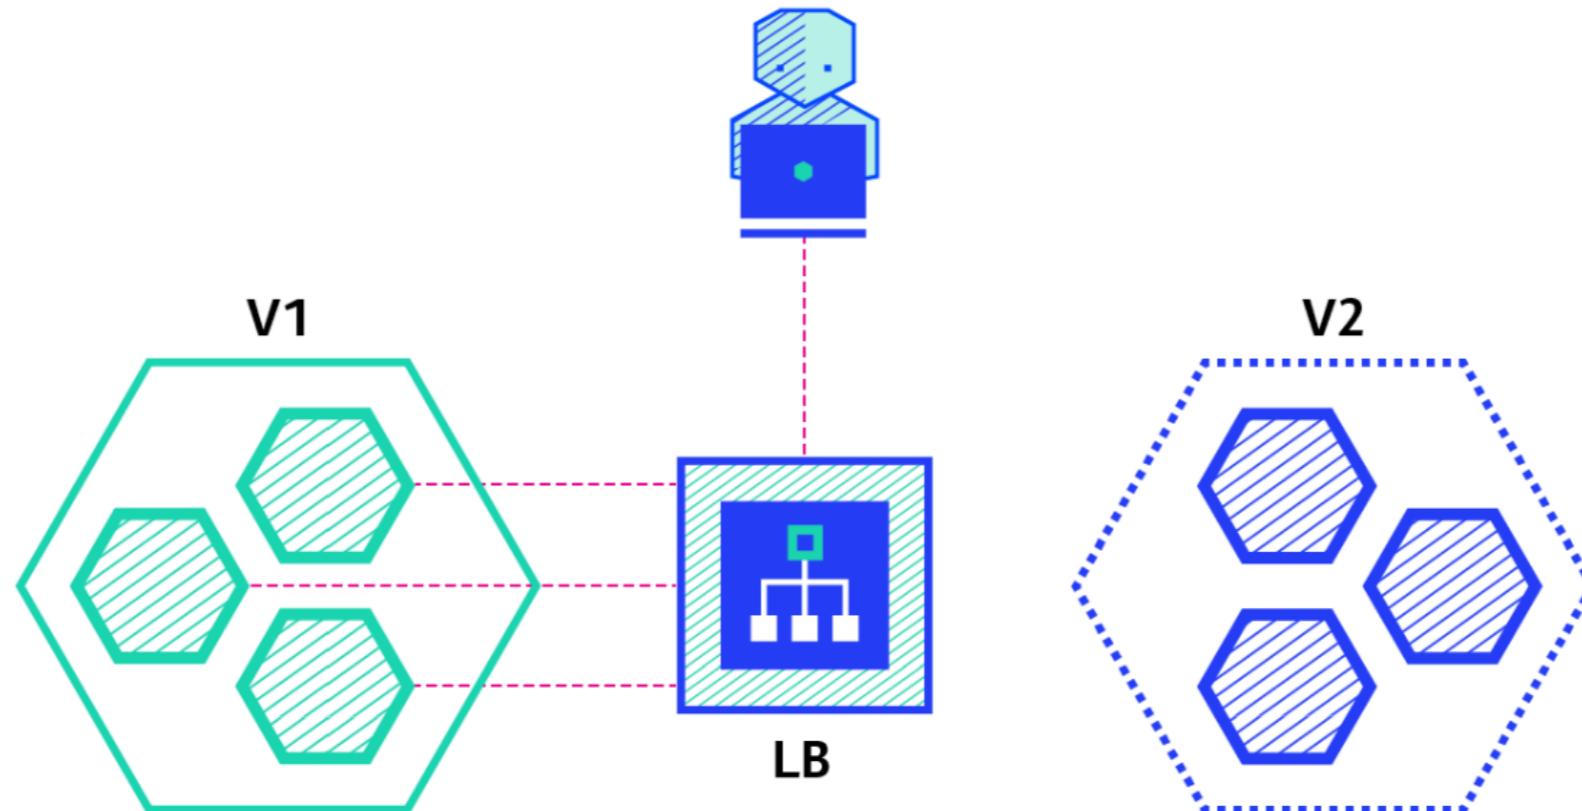
2. Ramped

Slow roll out by replace instance one-by-one



3. Blue/Green

Current version is called **Blue**
New version is called **Green**

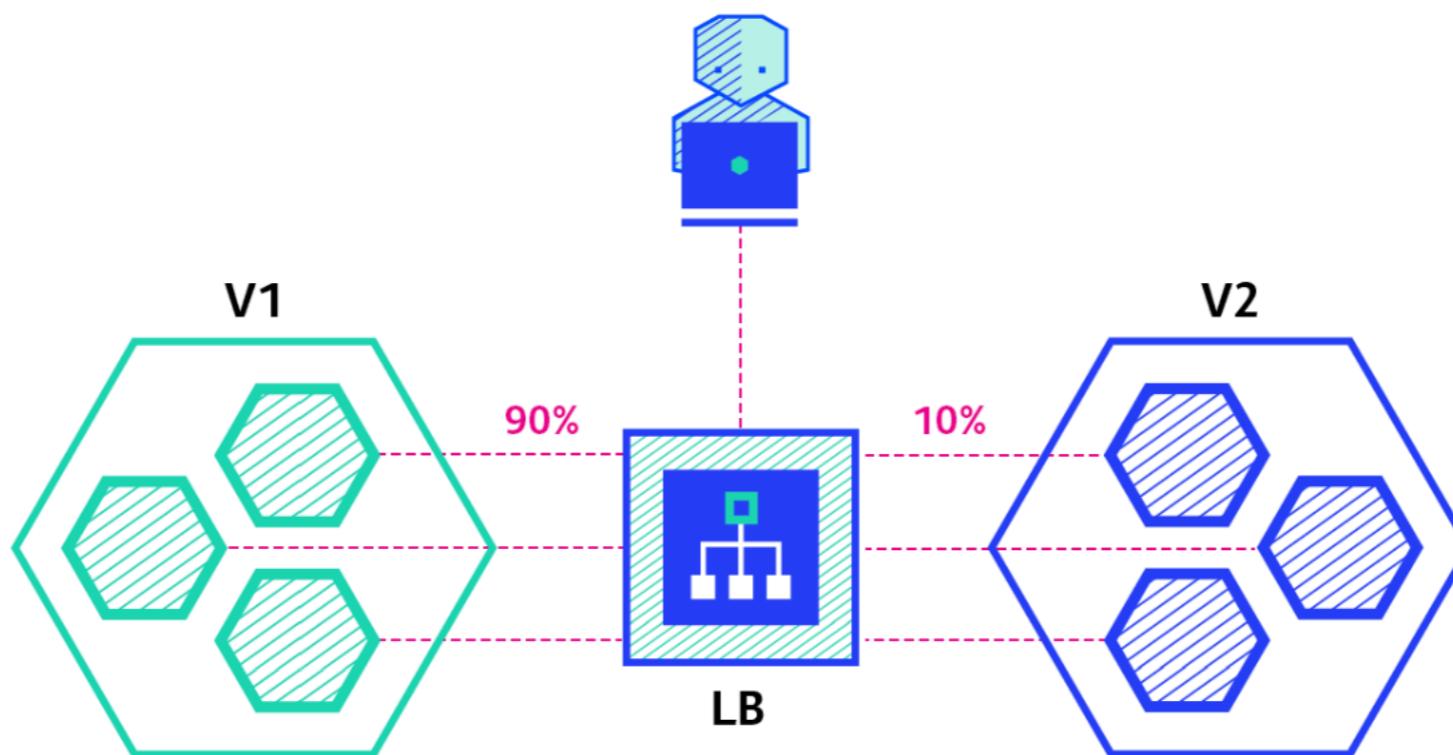


4. Canary

Shift production traffic from version A to B

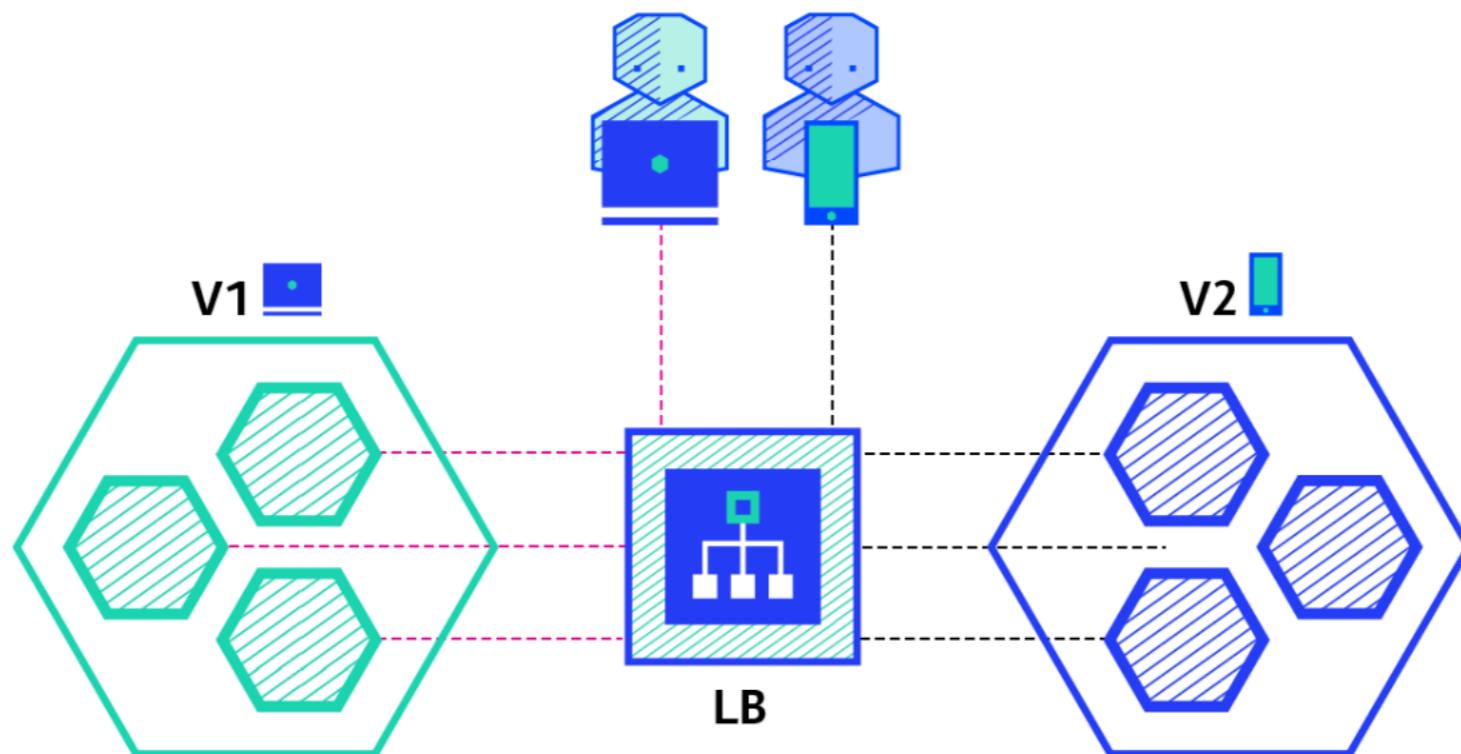
Traffic is split based-on weight

Use when tests are lacking/not reliable and less confident in system



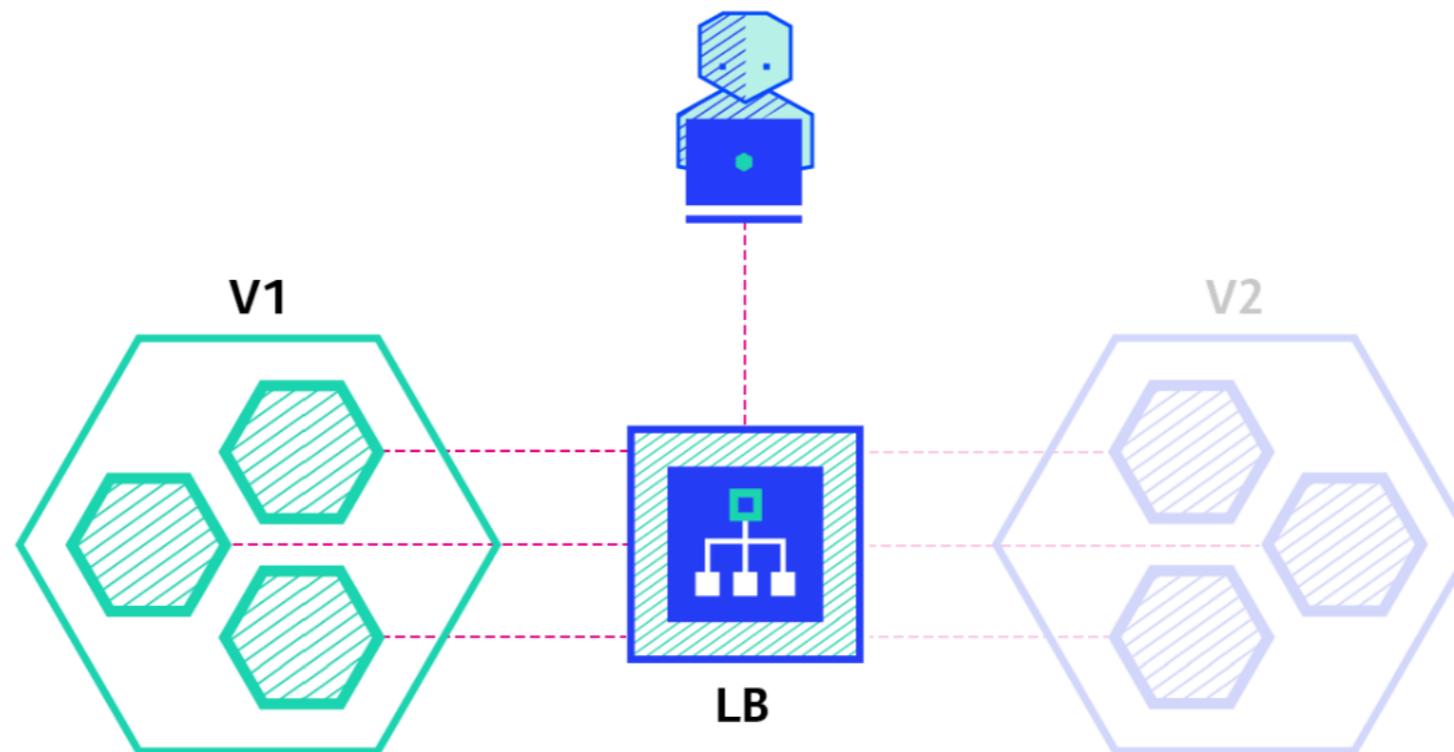
5. A/B testing

Routing the subset of users to new services under the specific condition



6. Shadow

Release version B alongside version A
Send request's A to B without production impact



DEPLOYMENT STRATEGIES

When it comes to production, a ramped or blue/green deployment is usually a good fit, but proper testing of the new platform is necessary.

Blue/green and shadow strategies have more impact on the budget as it requires double resource capacity. If the application lacks in tests or if there is little confidence about the impact/stability of the software, then a canary, a/b testing or shadow release can be used.

If your business requires testing of a new feature amongst a specific pool of users that can be filtered depending on some parameters like geolocation, language, operating system or browser features, then you may want to use the a/b testing technique.



Strategy	ZERO DOWNTIME	REAL TRAFFIC TESTING	TARGETED USERS	CLOUD COST	ROLLBACK DURATION	NEGATIVE IMPACT ON USER	COMPLEXITY OF SETUP
RECREATE version A is terminated then version B is rolled out	✗	✗	✗	■ ■ ■	■ ■ ■	■ ■ ■	□ □ □
RAMPED version B is slowly rolled out and replacing version A	✓	✗	✗	■ ■ ■	■ ■ ■	■ □ □	■ □ □
BLUE/GREEN version B is released alongside version A, then the traffic is switched to version B	✓	✗	✗	■ ■ ■	□ □ □	■ ■ □	■ ■ □
CANARY version B is released to a subset of users, then proceed to a full rollout	✓	✓	✗	■ ■ ■	■ □ □	■ □ □	■ ■ □
A/B TESTING version B is released to a subset of users under specific condition	✓	✓	✓	■ ■ ■	■ □ □	■ □ □	■ ■ ■
SHADOW version B receives real world traffic alongside version A and doesn't impact the response	✓	✓	✗	■ ■ ■	□ □ □	□ □ □	■ ■ ■

